

# The Reemergence of Datalog



A photograph of two men in a hallway. The man on the left, wearing glasses and a dark suit, is pointing his finger directly at the camera with a serious expression. The man on the right, wearing a green jacket over a blue shirt, has his mouth wide open as if shouting or screaming. The background shows a doorway and some foliage.

# Return of the Living Datalog

I like  
turtles



# Data



# Rectangles

COF_NAME	SUP_ID	PRICE	SALES	TOTAL
Pumpkin Spice	138	0.99	400,000	396,000
Dung Spice	666	2.00	null	null

# RDL

```
create table COFFEE
(COF_NAME varchar(32) NOT NULL,
SUP_ID int NOT NULL,
PRICE numeric(10,2) NOT NULL,
SALES integer NOT NULL,
TOTAL integer NOT NULL,
PRIMARY KEY (COF_NAME),
FOREIGN KEY (SUP_ID) REFERENCES SUPPLIERS (SUP_ID));
```

## **COFFEE**

**COF\_NAME** : varchar (PK)  
**SUP\_ID** : int (FK)  
**PRICE** : numeric  
**SALES** : integer  
**TOTAL** : integer

# Rectangulation

- Relationship between entities
- Sparse data
- Multi-valued attributes
- **PLace-Oriented Programming**



Java

```
public static void viewTable(Connection con) throws SQLException {  
    Statement stmt = null;  
    String query = "select COF_NAME, SUP_ID, " +  
        "PRICE, SALES, TOTAL from " +  
  
    try {  
        stmt = con.  
        ResultSet rs = stmt.executeQuery(query);  
        while (rs.next()) {  
            String cofName = rs.getString("COF_NAME");  
            int supId = rs.getInt("SUP_ID");  
            float price = rs.getFloat("PRICE");  
            int sales = rs.getInt("SALES");  
            int total = rs.getInt("TOTAL");  
            doPrint(cofName, supId, price, sales, total);  
        }  
    } catch (SQLException e) {  
        JDBCUtilities.printSQLException(e);  
    } finally {  
        if (stmt != null) { stmt.close(); }  
    }  
}
```

Where's  
the  
Data?

Java

```
public static void viewTable(Connection con) throws SQLException {
    Statement stmt = null;
    String query = "select COF_NAME, SUP_ID, " +
                   "PRICE, SALES, TOTAL from " +
                   dbName + ".COFFEE";

    try {
        stmt = con.createStatement();
        ResultSet rs = stmt.executeQuery(query);
        while (rs.next()) {
            String coffeeName = rs.getString("COF_NAME");
            int supplierID = rs.getInt("SUP_ID");
            float price = rs.getFloat("PRICE");
            int sales = rs.getInt("SALES");
            int total = rs.getInt("TOTAL");
            doPrint(coffeeName, supplierID, price, sales, total);
        }
    } catch (SQLException e) {
        JDBCUtilities.printSQLException(e);
    } finally {
        if (stmt != null) { stmt.close(); }
    }
}
```

# Java

```
public static void viewTable(Connection con) throws SQLException {  
  
    ResultSet rs = stmt.executeQuery(query);  
    while (rs.next()) {  
        String coffeeName = rs.getString("COF_NAME");  
        int supplierID = rs.getInt("SUP_ID");  
        float price = rs.getFloat("PRICE");  
        int sales = rs.getInt("SALES");  
        int total = rs.getInt("TOTAL");  
        doPrint(coffeeName, supplierID, price, sales, total);  
    }  
}
```

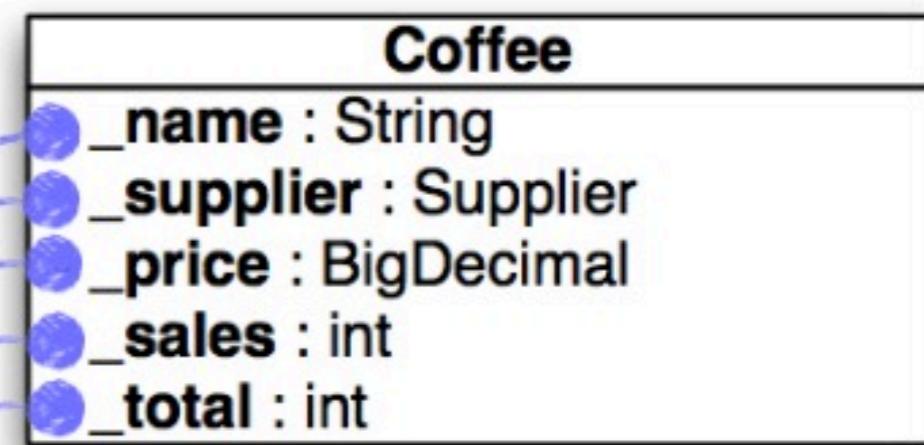
# Java



# RDL

```
public static void viewTable(Connection con) throws SOLEException {
    Statement stmt = null;
    String query = "select COF_NAME, SUP_ID, " +
                   "PRICE, SALES, TOTAL from " +
                   dbName + ".COFFEES";

    try {
        stmt = con.createStatement();
        ResultSet rs = stmt.executeQuery(query);
        while (rs.next()) {
            String coffeeName = rs.getString("COF_NAME");
            int supplierID = rs.getInt("SUP_ID");
            float price = rs.getFloat("PRICE");
            int sales = rs.getInt("SALES");
            int total = rs.getInt("TOTAL");
            doPrint(coffeeName, supplierID, price, sales, total);
        }
    } catch (SQLException e) {
        JDBCUtilities.printSQLException(e);
    } finally {
        if (stmt != null) { stmt.close(); }
    }
}
```





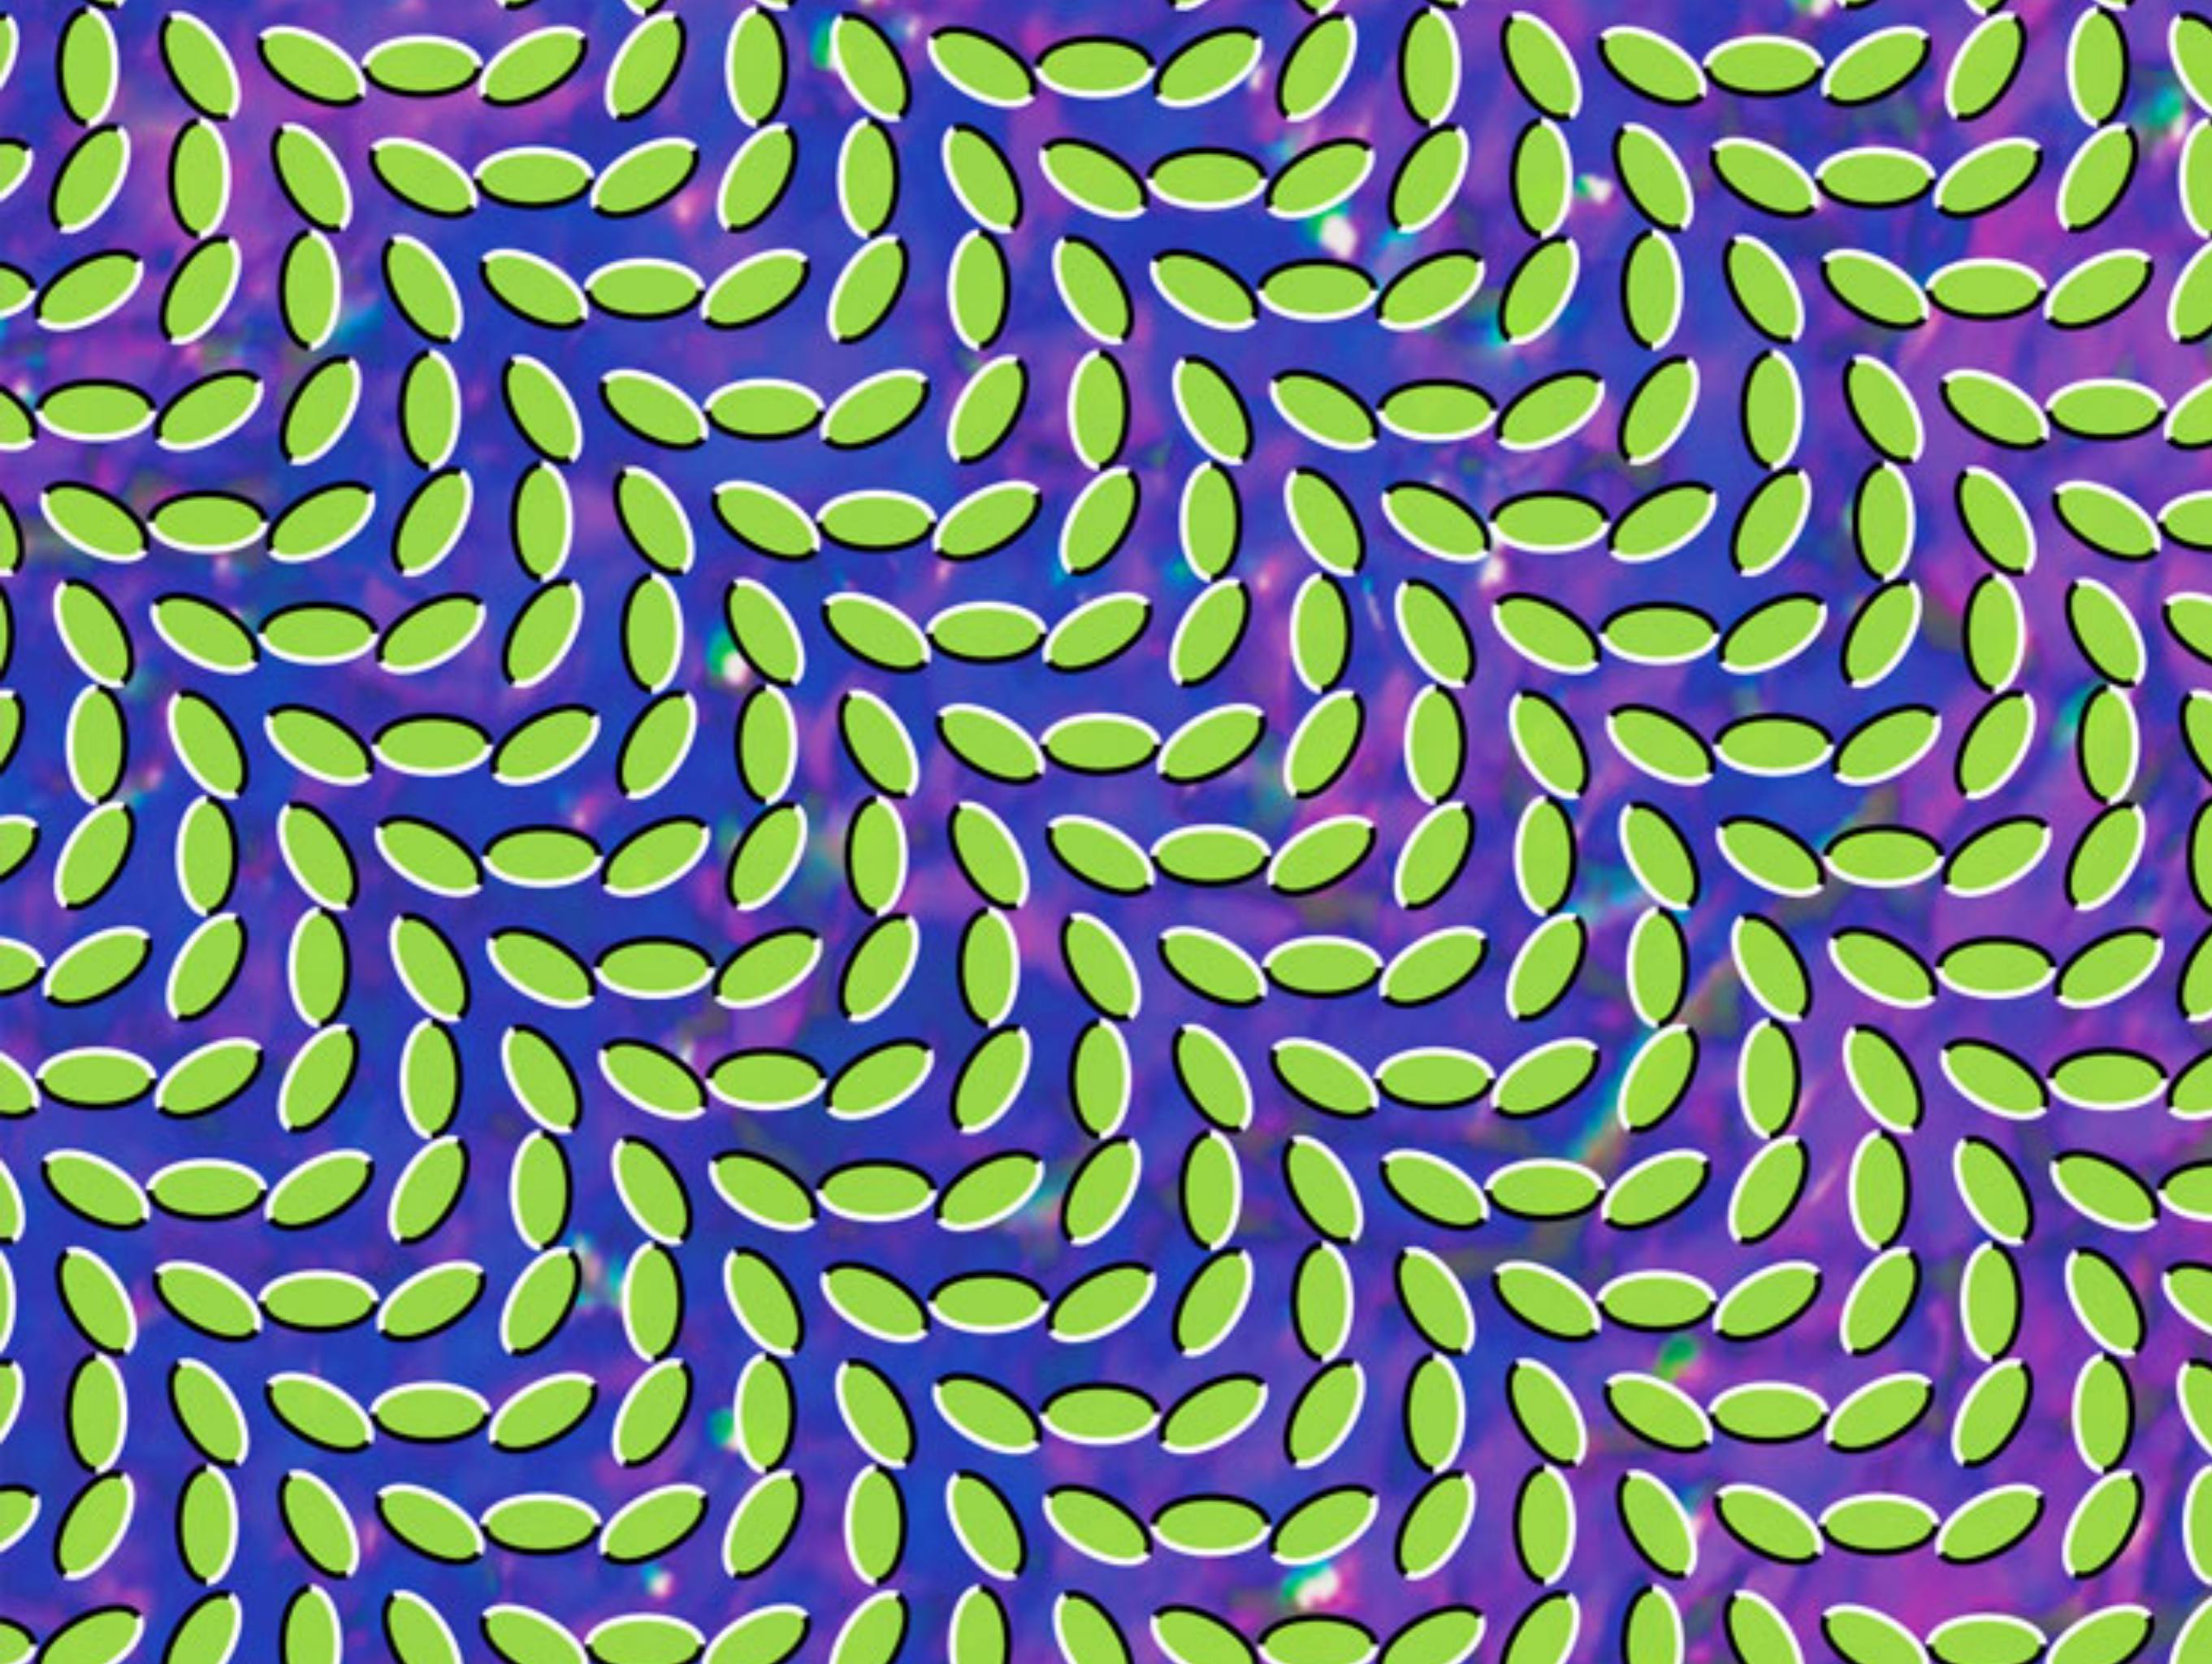
A black and white portrait of a man with short, light-colored hair and glasses, smiling broadly. He is wearing a dark jacket over a light-colored shirt. A large, dark blue speech bubble originates from his mouth, containing the text "ORMG!" in a bold, white, sans-serif font.

**ORMG!**

Code as Code.

---

Data as Data.



Cookies

User information

Protocols

Lisp

Schemas

Events

Chess moves

TX

...

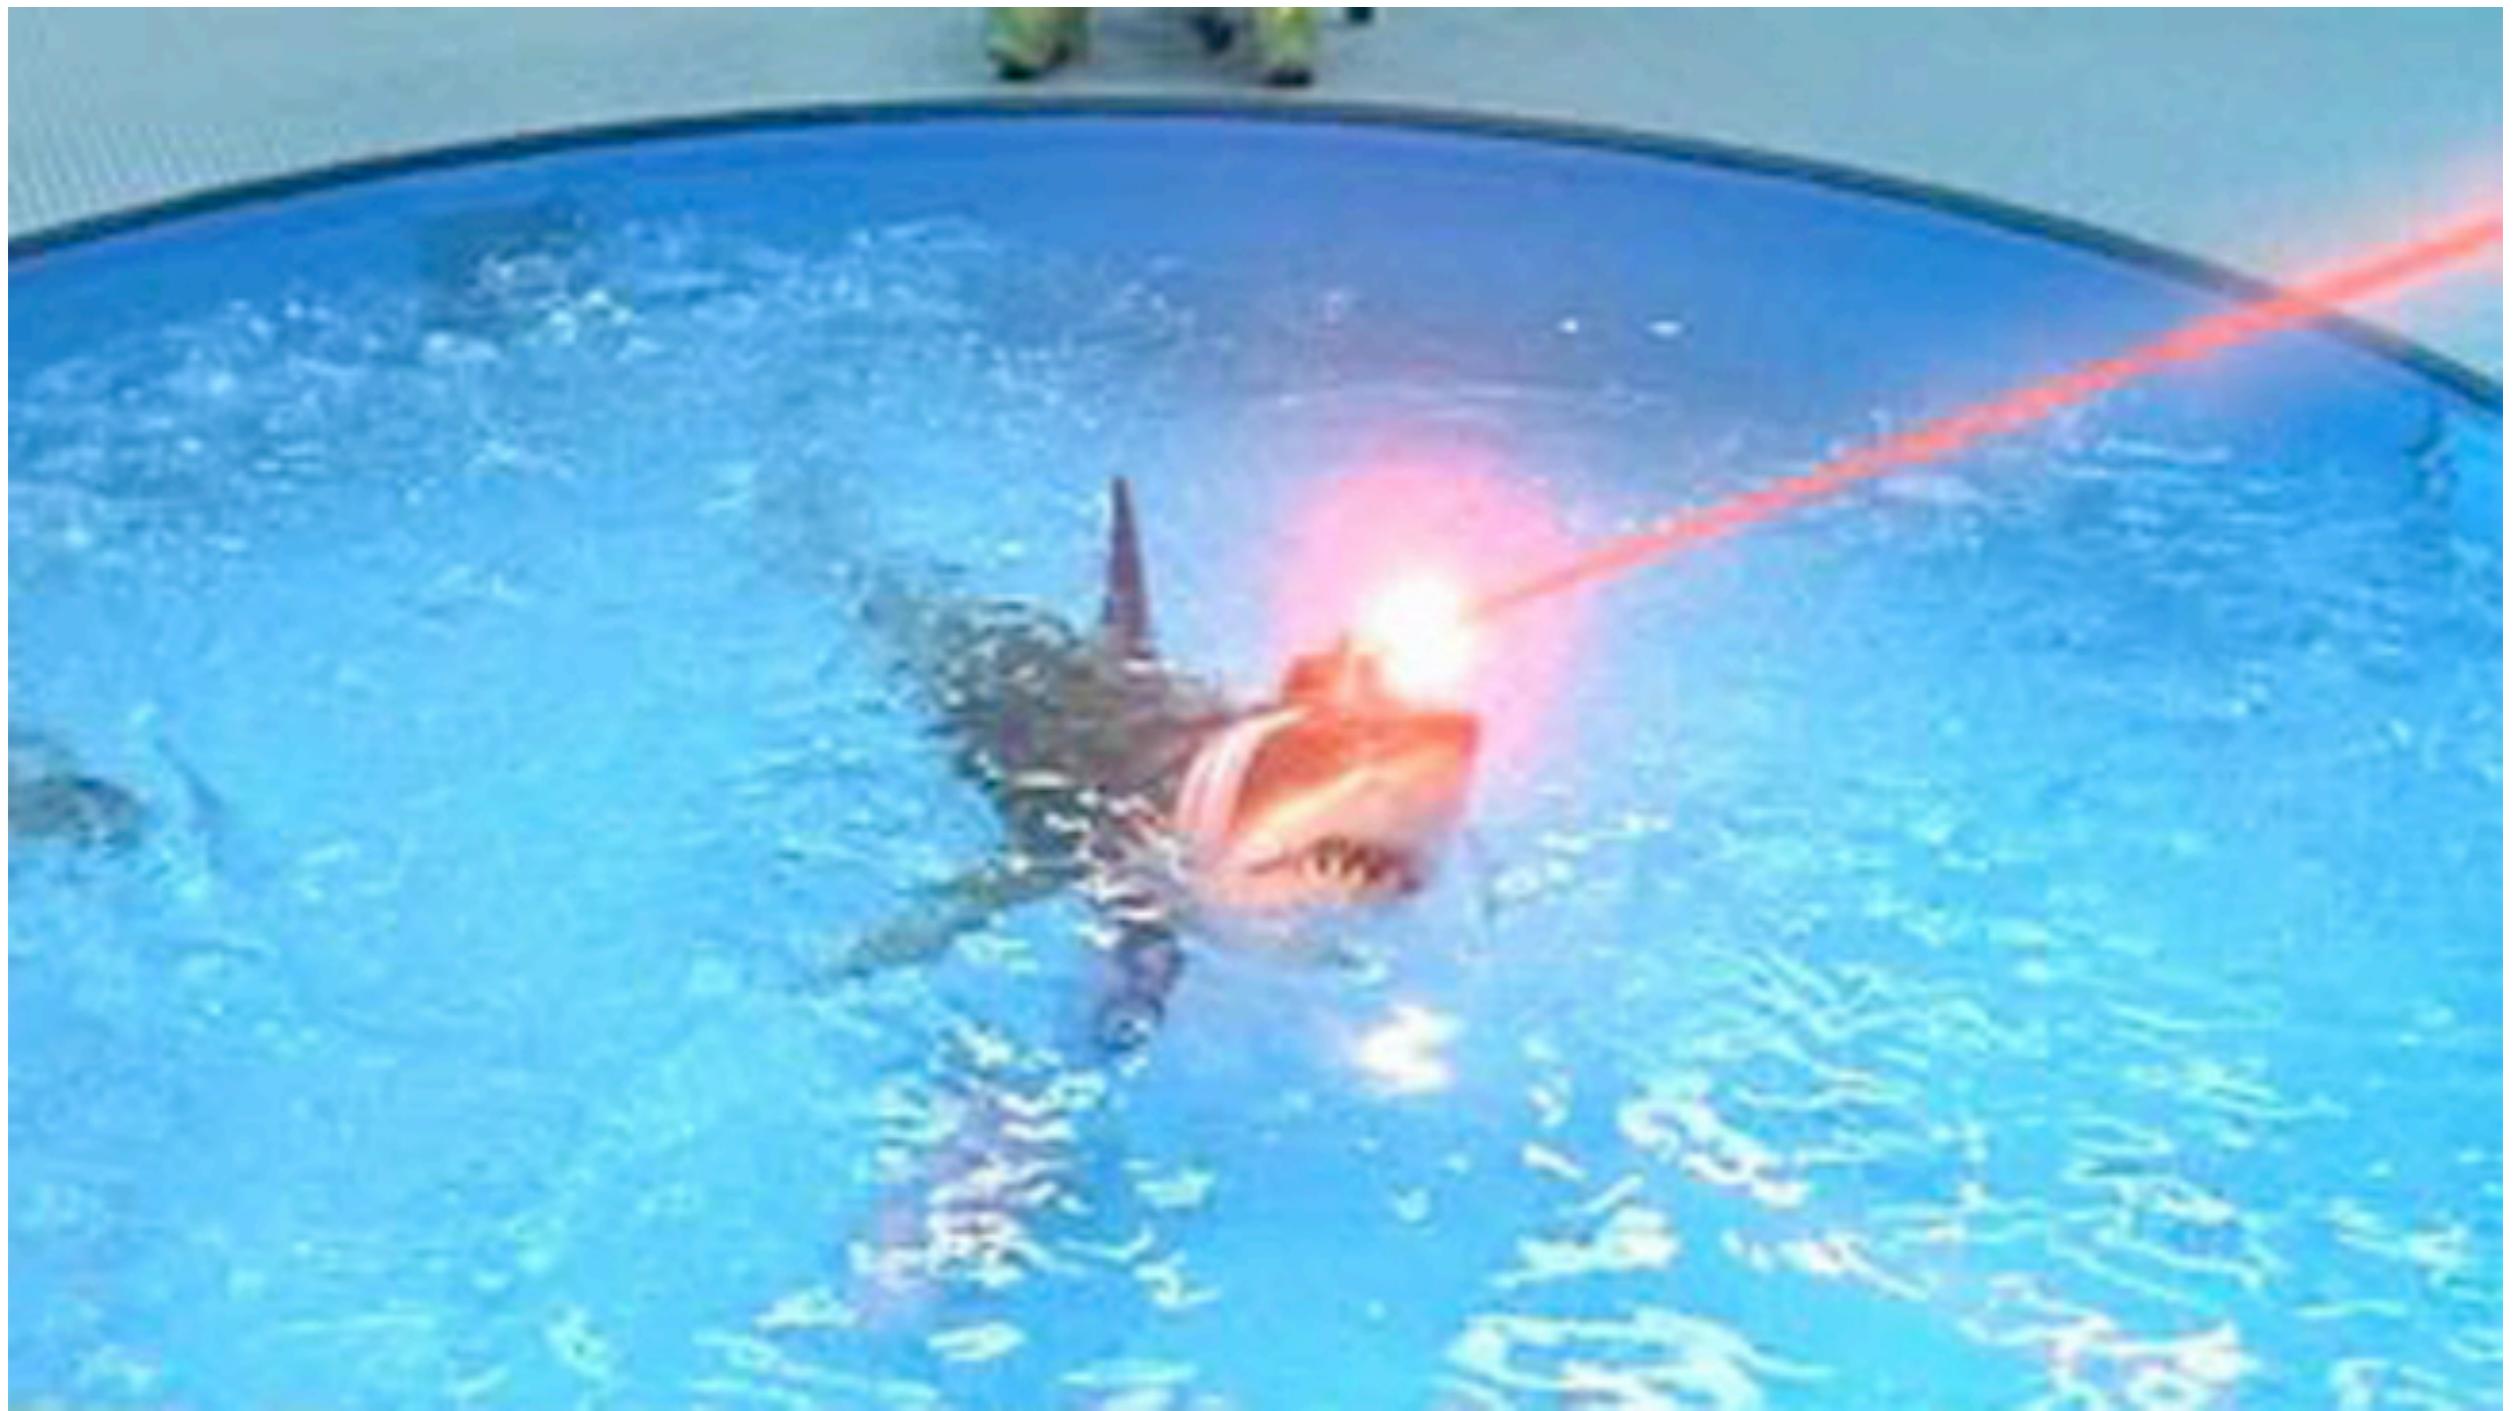
as Data.

Code as Code.

---

Data as Data.

# Unification



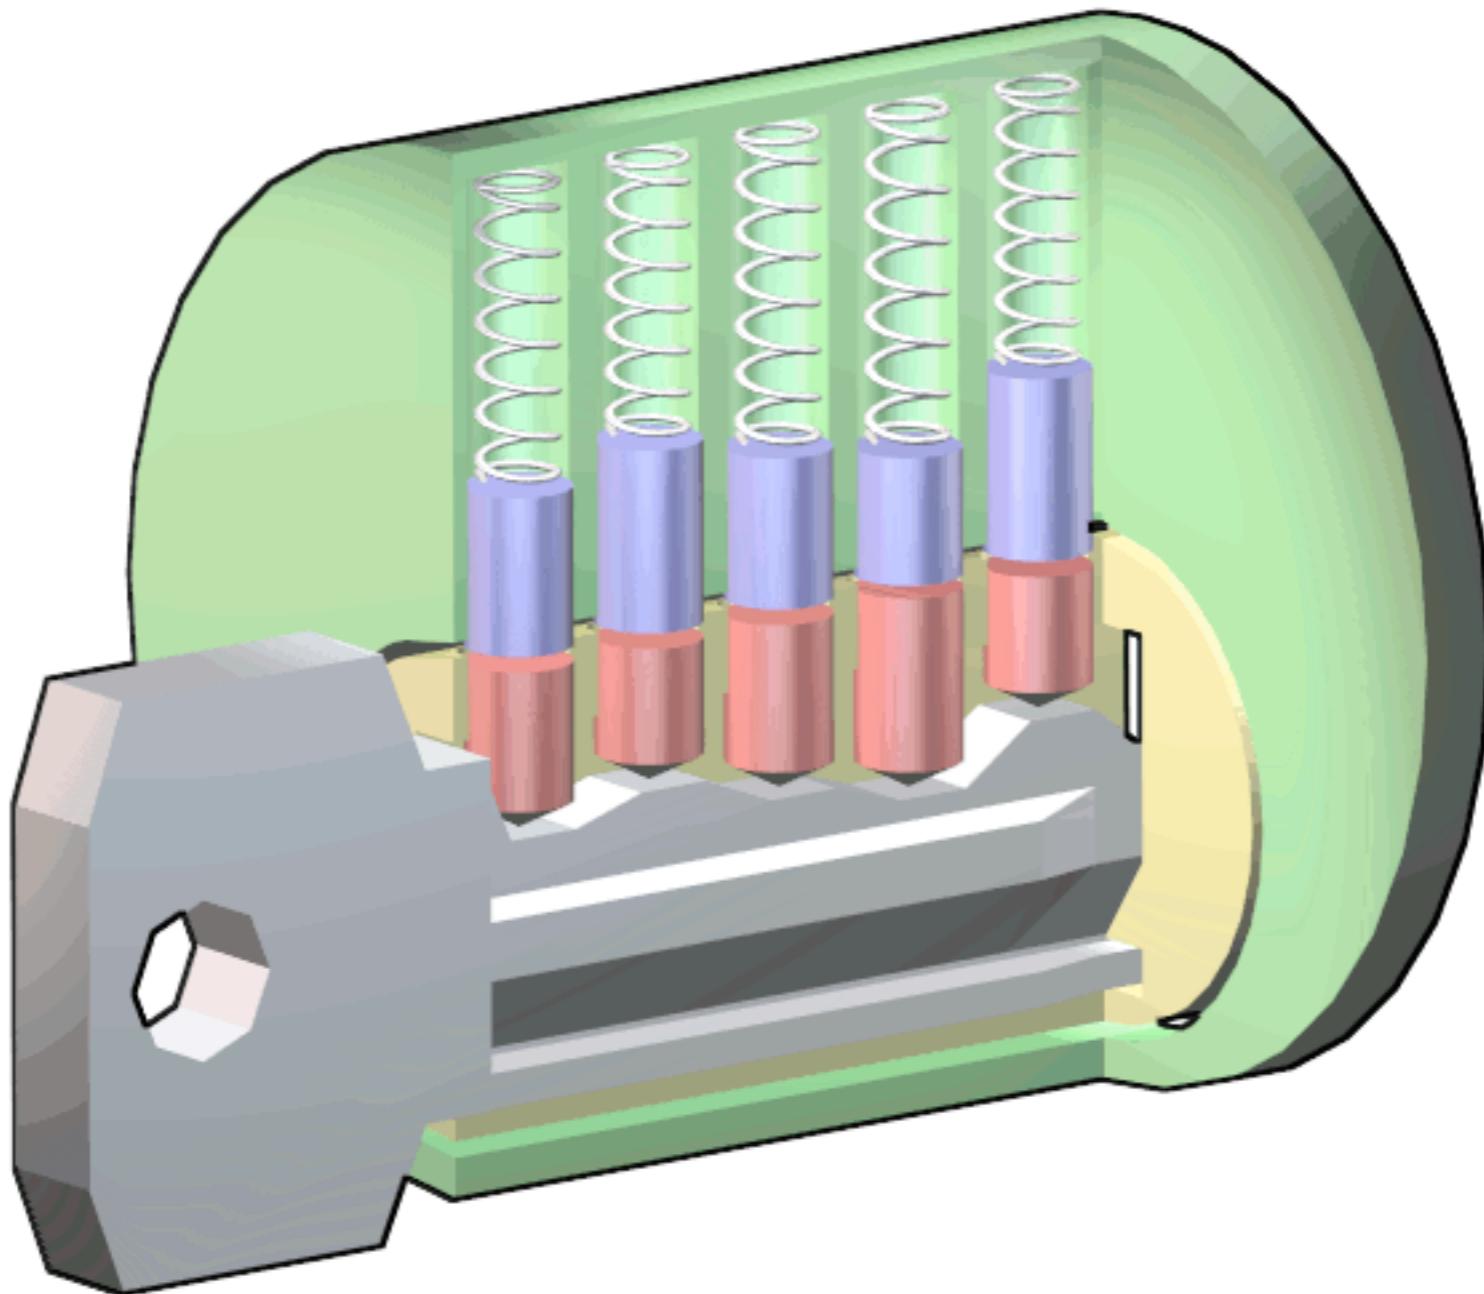
# Your data

```
{:first "Dario"  
:last  "Argento"  
:genre :giallo}
```

# *punching holes in* Your data

```
{:first ?  
:last ?  
:genre :giallo}
```

# *fitting the holes in Your data*



# Variables

```
{:first '?first  
:last  '?last  
:genre :giallo}
```

# Deriving bindings

```
(unify
  { :first '?first
    :last  '?last
    :genre :giallo}

  { :first "Dario"
    :last  "Argento"
    :genre :giallo})

;=> {?first "Dario", ?last "Argento"}
```

# Substitution

```
(subst
  ['?first '?last]          ;; TEMPLATE
  {'?first "Dario",
   '?last "Argento"})      ;; BINDINGS
;=> ["Dario" "Argento"]
```

# Leaving variables

```
(subst '[1 2 ?x ?y]  
' {?x [3 4 ?y 6]})  
;=> [1 2 [3 4 ?y 6] ?y]
```

# Related variables

```
(subst '[1 2 ?x ?y]  
' {?x [3 4 ?y 6]  
  ?y :BUB})  
;=> [1 2 [3 4 :BUB 6] :BUB]
```

# MGU

```
(unifier '[ (?a * ?x | 2) + (?b * ?x) + ?c]
          '[ ?z + (4 * 5) + 3])
;=> [ (?a * 5 | 2) + (4 * 5) + 3]
```



# Prolog

# Prolog data

type	example
atom	bill
atom	'hi there'
int	42
float	3.14
variable	X
variable	Name
variable	[L R]
functor	first(L).
list	[1,2]
operator	+ - / *

# Prolog facts

```
spawned(bill, carl).
spawned(bill, dilbert).
spawned(dilbert, phil).

aged(bill,81).
aged(carl,33).
aged(dilbert, 40).
aged(phi1,10).
```

# Prolog rules

```
descendant(X,Y) :- spawned(X,Y).  
descendant(X,Y) :- descendant(X,Z), spawned(Z,Y).
```

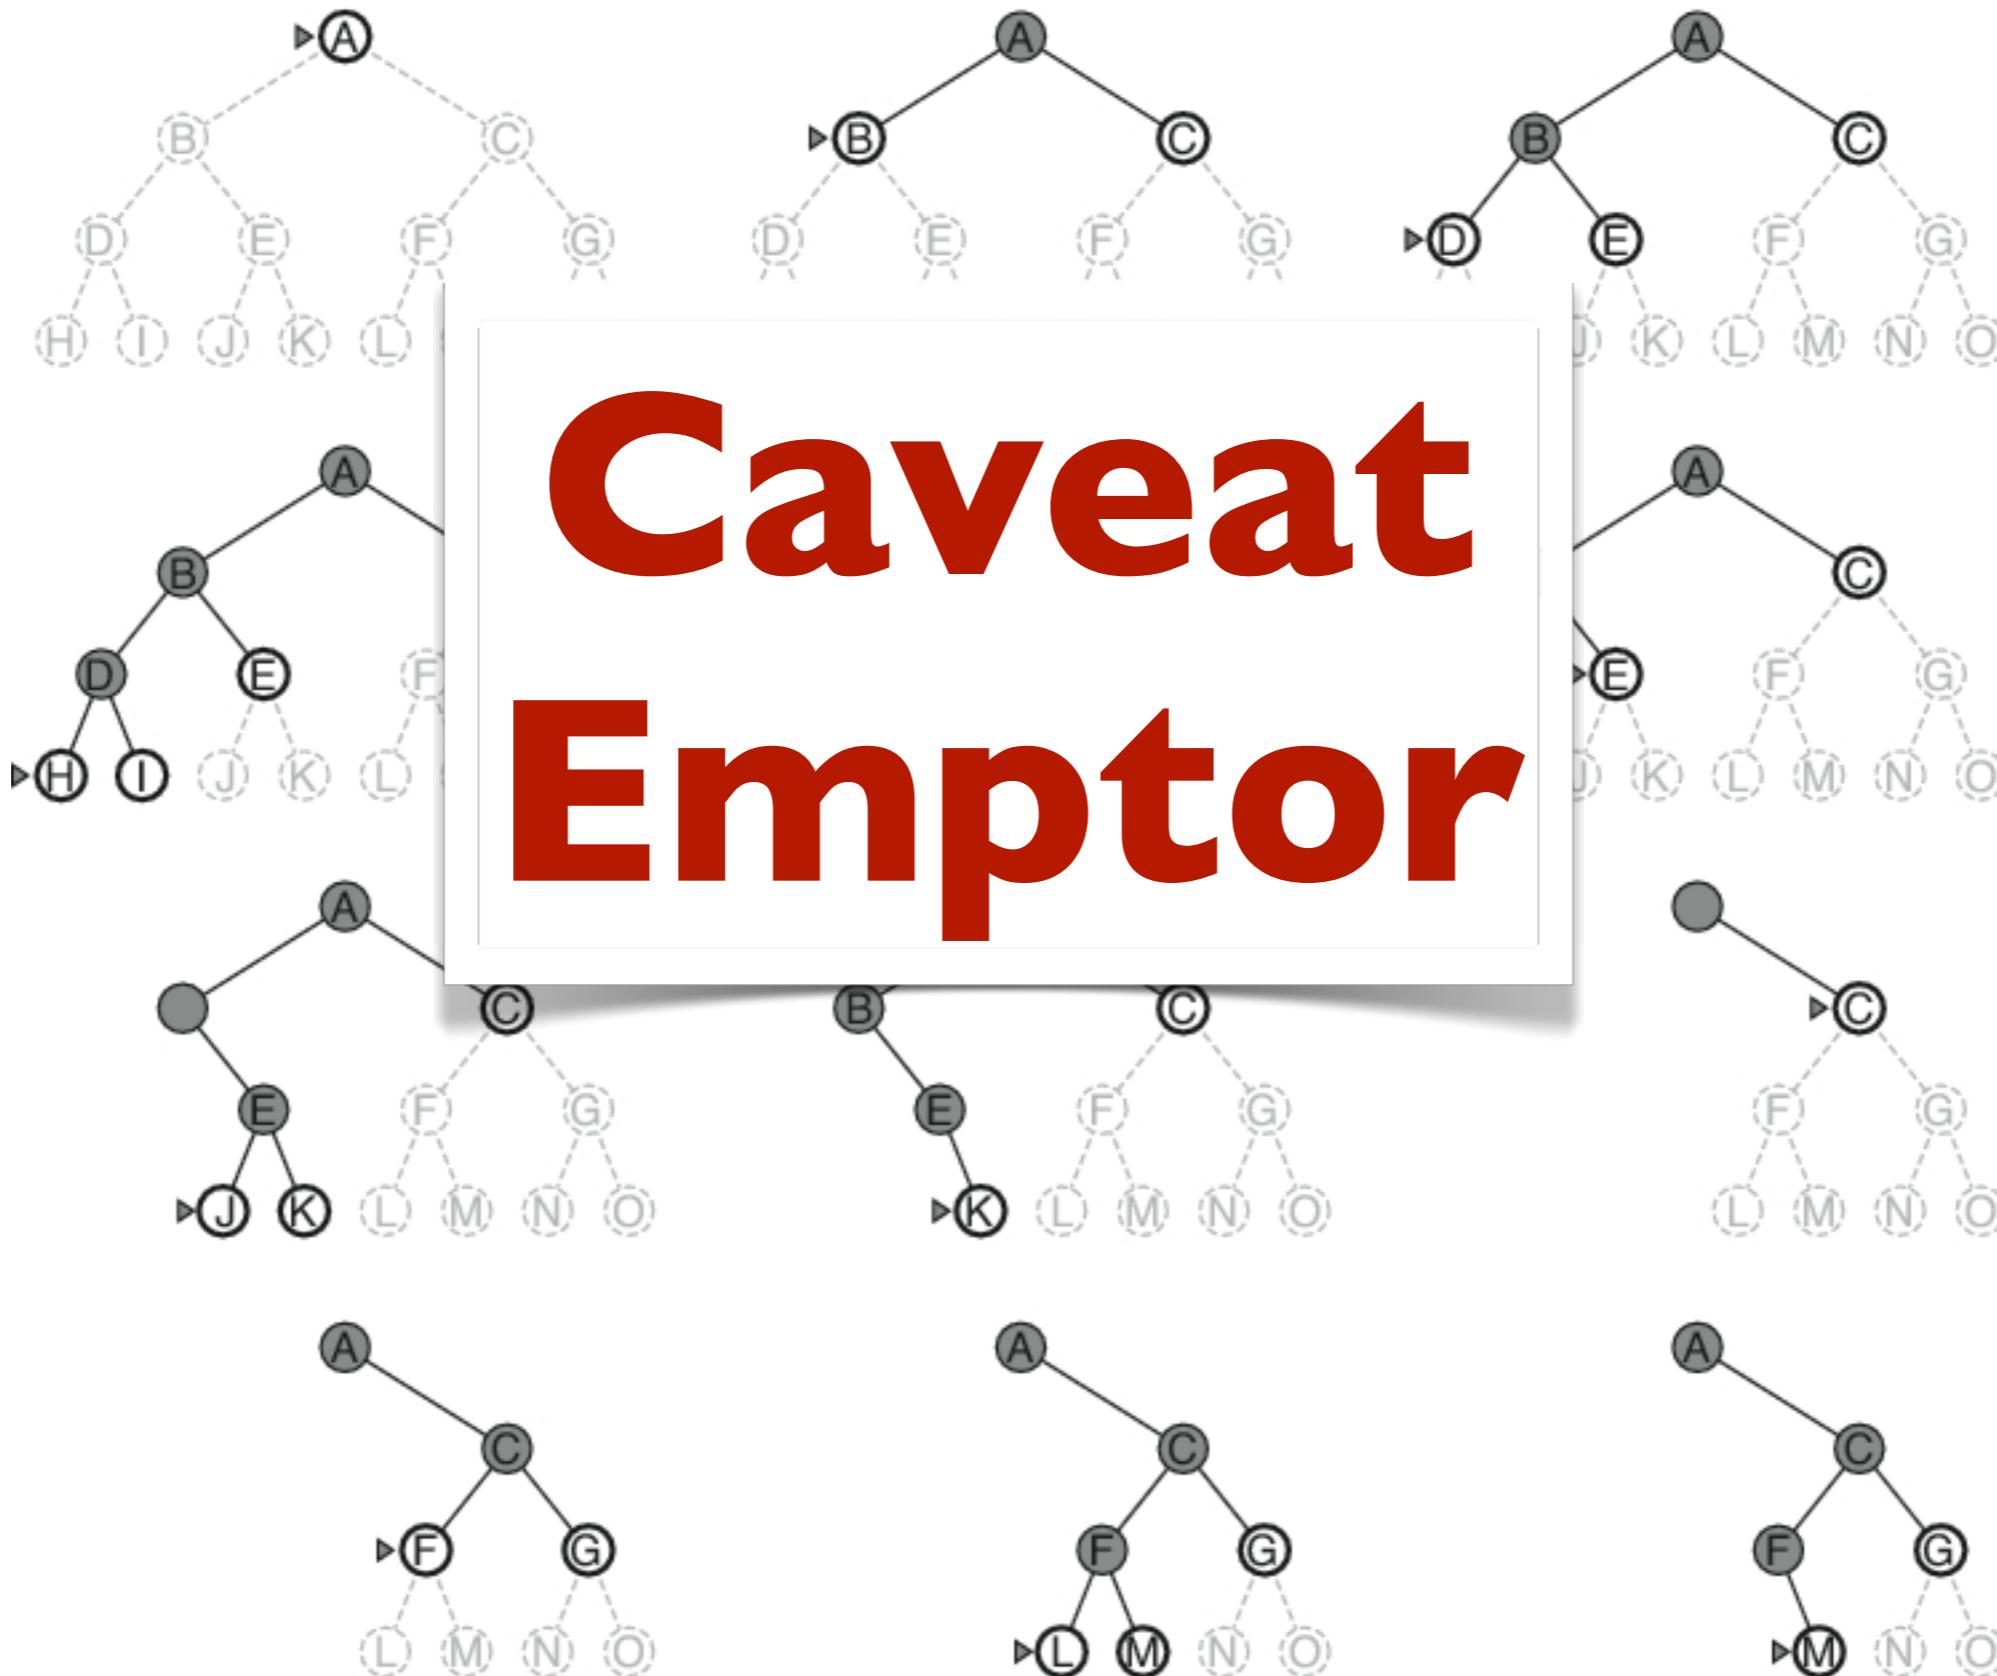
# Prolog query

```
descendant(bill,Kid),  
aged(Kid,Age),  
Age < 13.
```

```
% Kid = phil,  
% Age = 10 .
```

# Data as Code!

**Caveat  
Emptor**



# Caveats

- Clause-order dependence
- Non-termination
- Infection of imperative

# Clause order dependence

```
descendant(X,Y) :- spawned(X,Y).
descendant(X,Y) :- descendant(X,Z), spawned(Z,Y).

descendant(bill,D).

% D = carl ;
% D = dilbert ;
% D = phil ;
% runs forever!!!
```

# Clause order dependence

```
descendant(X,Y) :- spawned(X,Y).
descendant(X,Y) :- spawned(Z,Y), descendant(X,Z).

descendant(bill,D).

% D = carl ;
% D = dilbert ;
% D = phil ;
% false.
```

# Non-termination

```
married(A,B) :- married(B,A).
```

# Non-termination

```
married(edard,catelyn).  
% never stops
```

# Fixed?

```
are_married(A,B) :- married(A,B).  
are_married(A,B) :- married(B,A).  
married(edard, catelyn).  
  
?- are_married(edard,catelyn).  
true.  
  
?- are_married(catelyn,edard).  
true.
```

# Imperative infection

- !
- fail

# Cut

```
statement :- condition, !, then.  
statement :- else.
```

```
zzz(T,A) :- T > 2100, !, A = sleep.  
zzz(_,A) :- A = play_chess.
```

```
?- zzz(2000, A).  
A = play_chess.
```

```
?- zzz(2300, A).  
A = sleep.
```

# fail

```
pparts(L) :- append(A, B, L),
            write(A), write(' '), write(B), nl,
            fail.

?- pparts([1,2]).  
[ ] [1,2]  
[1] [2]  
[1,2] []  
false.
```

# Prolog is prelude

- Powerful
- Often beautiful
- Not as declarative as we'd like



# Datalog

# Datalog is...

- A query language
- Not Turing complete
- Explicit
- Simple
  - to use
  - ... and implement
  - ... kinda

# History

- 1977: Gallaire and Minker's Symposium on Logic Data Bases
- 1980s: Nail, LDL, Coral
- 1995: Stonebraker and Hellerstein declare "no practical applications ..."
- **The dark years...**

# History

- 2002: Binder, a logic-based security language by DeTreville
- 2000s: Declarative networking, bddbddb, Orchestra CDSS, Doop, SecureBlox, Dedalus, more
- 2010: The Declarative Imperative by Hellerstein!
- Today: Bloom, Cascalog, Datomic, LogicBlox, more

# Datalog is also

- Declarative logic programming with termination
- Recursive queries
- Implicit joins

# EAV

Entity	Attribute	value
999-99-9999	:person/age	42

- “Entities” (objects?) - a grouping of tuples
- Make that efficient

# Query elements

element	specifies	details
:where	constraints	patterns, predicates, functions
:find	return values	variable bindings
:in	data sources	databases, collections, objects

# Patterns

pattern	binds type	input	?a binding
?a	scalar	42	42
[ ?a ?b ]	tuple	[1 2]	1
[ ?a ... ]	collection	[1 2]	1, 2
[ [ ?a ?b ?c ] ]	relation	[john :likes :pizza] [jane :likes :pasta]	john, jane

# Simple query

- Find all language entities with a website entry

```
[ :find ?language
:where [?language :lang/website]]  
;=> [[1234567890] [1234567891] ...]
```

# Simple query

- Find all language entities with a website entry

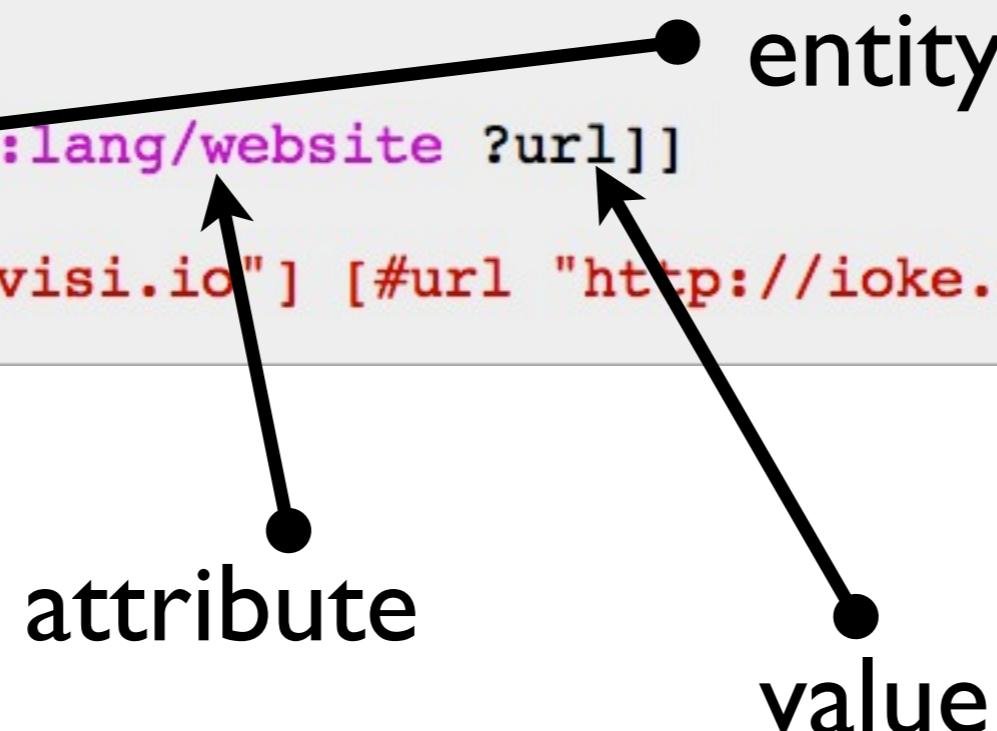
```
[ :find ?language
:where [?language :lang/website]
;=> [[1234567890] [1234567891] ...]
```

entity  
attribute

# Binding query

- Find all language URLs with a website entry

```
[ :find ?url  
:where [?language :lang/website ?url]]  
;=> [[#url "http://visi.io"] [#url "http://ioke.org"] ...]
```



# Join

```
[ :find ?name ?url  
:where [ ?language :lang/website ?url]  
       [ ?language :lang/name ?name] ]  
  
;=> [[ "Visi" #url "http://visi.io"]  
;      ["Ioke" #url "http://ioke.org"]  
;      ...]
```

- Repeating ?language indicates a join

# Rules

```
[[ (lang-anchor ?name ?url) ←  
  [?language :lang/website ?url]  
  [?language :lang/name ?name] ]]
```

```
[ :find ?name ?url  
  :where (lang-anchor ?name ?url)]
```

```
;=> [ ["Visi" #url "http://visi.io"]  
;      ["Ioke" #url "http://ioke.org"]  
;      ... ]
```

head  
body

- All variables in the head, **must** appear in the body

# Recursive rules

```
' [[ (direct-influence ?old ?new)
  [ ?old :influenced ?new] ]

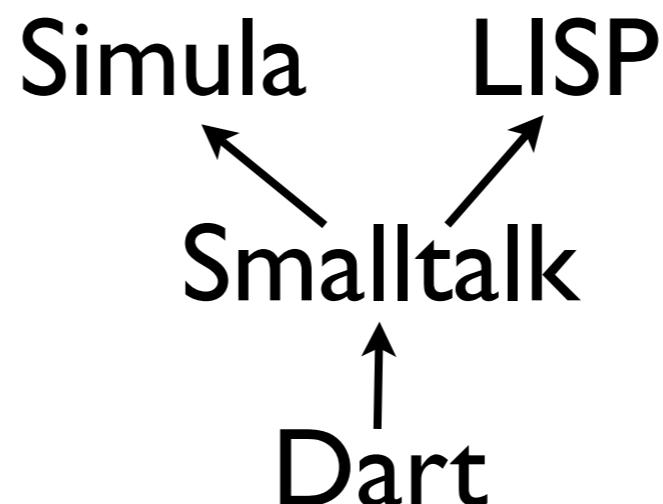
[ (remote-influence ?old ?new)
  (direct-influence ?old ?new) ]

[ (remote-influence ?old ?new)
  (direct-influence ?old ?intermediate)
  (remote-influence ?intermediate ?new) ]]
```

# Recursive rules

```
(q '[ :find ?influence
  :in $ %
  :where
  [ ?new :lang "Dart"]
  (remote-influence ?old ?new)
  [ ?old :lang ?influence])
  influences-db
  influces-rules)

;;=> #[["LISP"] ["Smalltalk"] ["Simula"]]
```





# Datomic

# Datalog plus

- No need for a database
- Time travel

# Where's the DB?

```
public static void viewTable(Connection con) throws SQLException {  
    Statement stmt = null;  
    String query = "select COF_NAME, SUP_ID, " +  
        "PRICE, SALES, TOTAL from " +  
        dbName + ".COFFEE";  
  
    try {  
        stmt = con.createStatement();  
        ResultSet rs = stmt.executeQuery(query);  
        while (rs.next()) {  
            String coffeeName = rs.getString("COF_NAME");  
            int supplierID = rs.getInt("SUP_ID");  
            float price = rs.getFloat("PRICE");  
            int sales = rs.getInt("SALES");  
            int total = rs.getInt("TOTAL");  
            doPrint(coffeeName, supplierID, price, sales, total);  
        }  
    } catch (SQLException e) {  
        JDBCUtilities.printSQLException(e);  
    } finally {  
        if (stmt != null) { stmt.close(); }  
    }  
}
```

# Where's the DB?

```
public static void viewTable(Connection con) throws SQLException {  
  
    stmt = con.createStatement();  
  
}  
}
```

# Where's the DB?

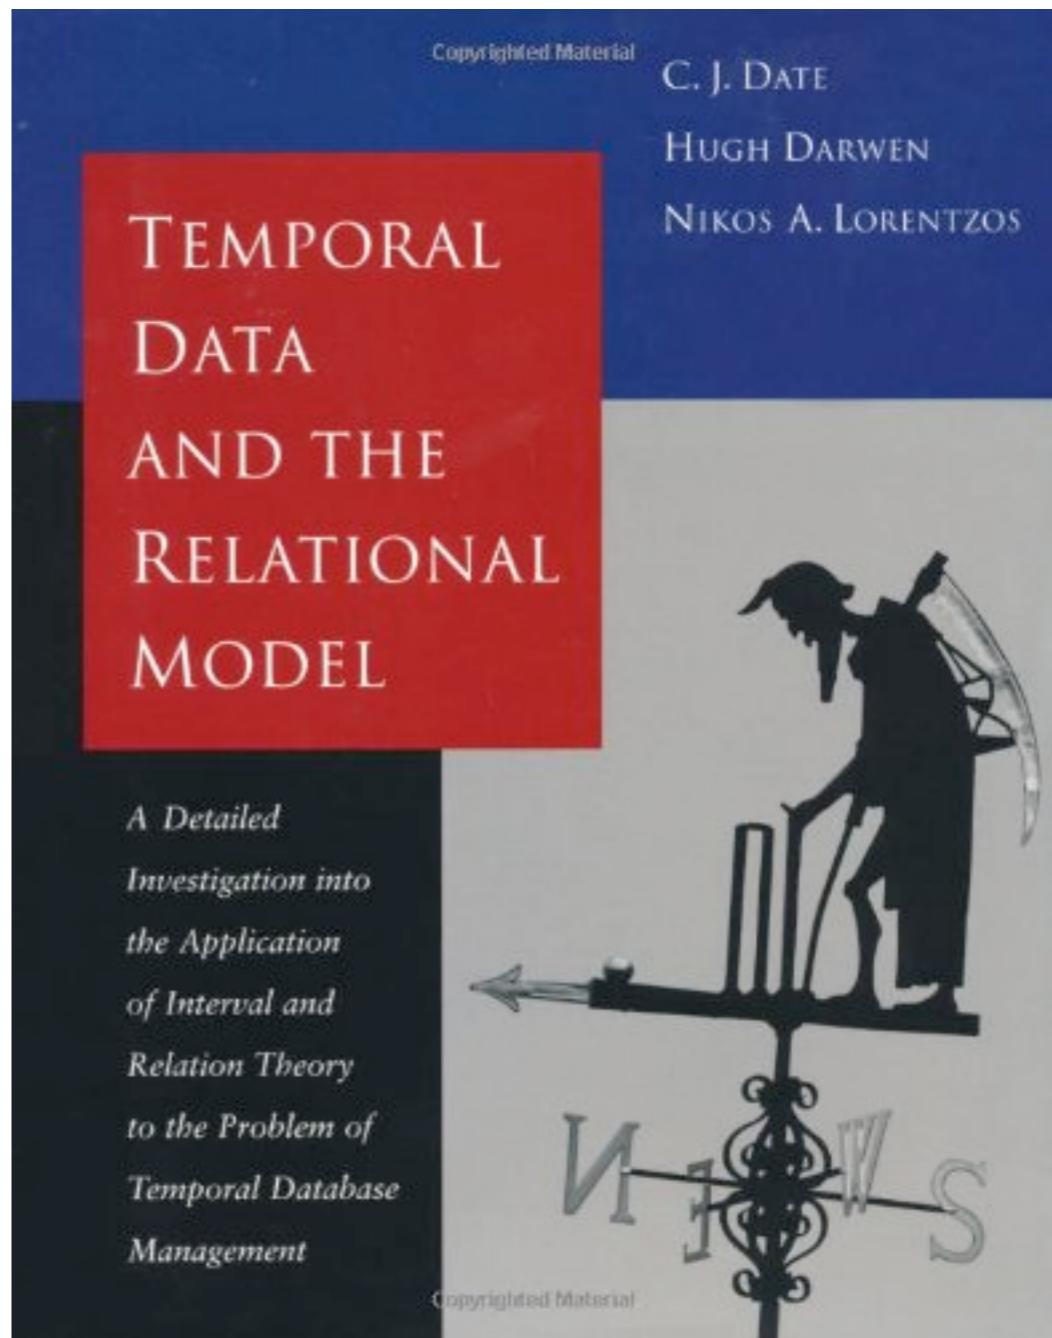
```
(datomic/q '[:find ?name ?url  
           :where (lang-anchor ?name ?url)]  
           the-database  
           the-rules)  
;; <- HINT: RIGHT HERE!
```

# Where's the DB?

```
(datomic/q '[:find ?name ?url
               :where (lang-anchor ?name ?url)]
               [[-100 :language/name "Visi"]
                [-200 :language/name "Ioke"]
                [-3   :language/name "Frink"]
                [-4   :language/name "Roy"]
                [-100 :language/url #url "http://visi.io"]
                [-200 :language/url #url "http://ioke.org"]])
the-rules)
```

# now

- How do you keep a notion of time in a relational database?
- Ever write now()?



# Always?

```
[john :favorite/food :pizza]
```

# Time

```
[john :favorite/food :pizza <last week>]    ;; ASSERTION  
[john :favorite/food :pho <now>]                ;; ANOTHER ASSERTION
```

- Total ordering of transactions
- Every datom retains a reference to its enclosing transaction
- Transactions are first-class entities, can have their own attributes

# Time

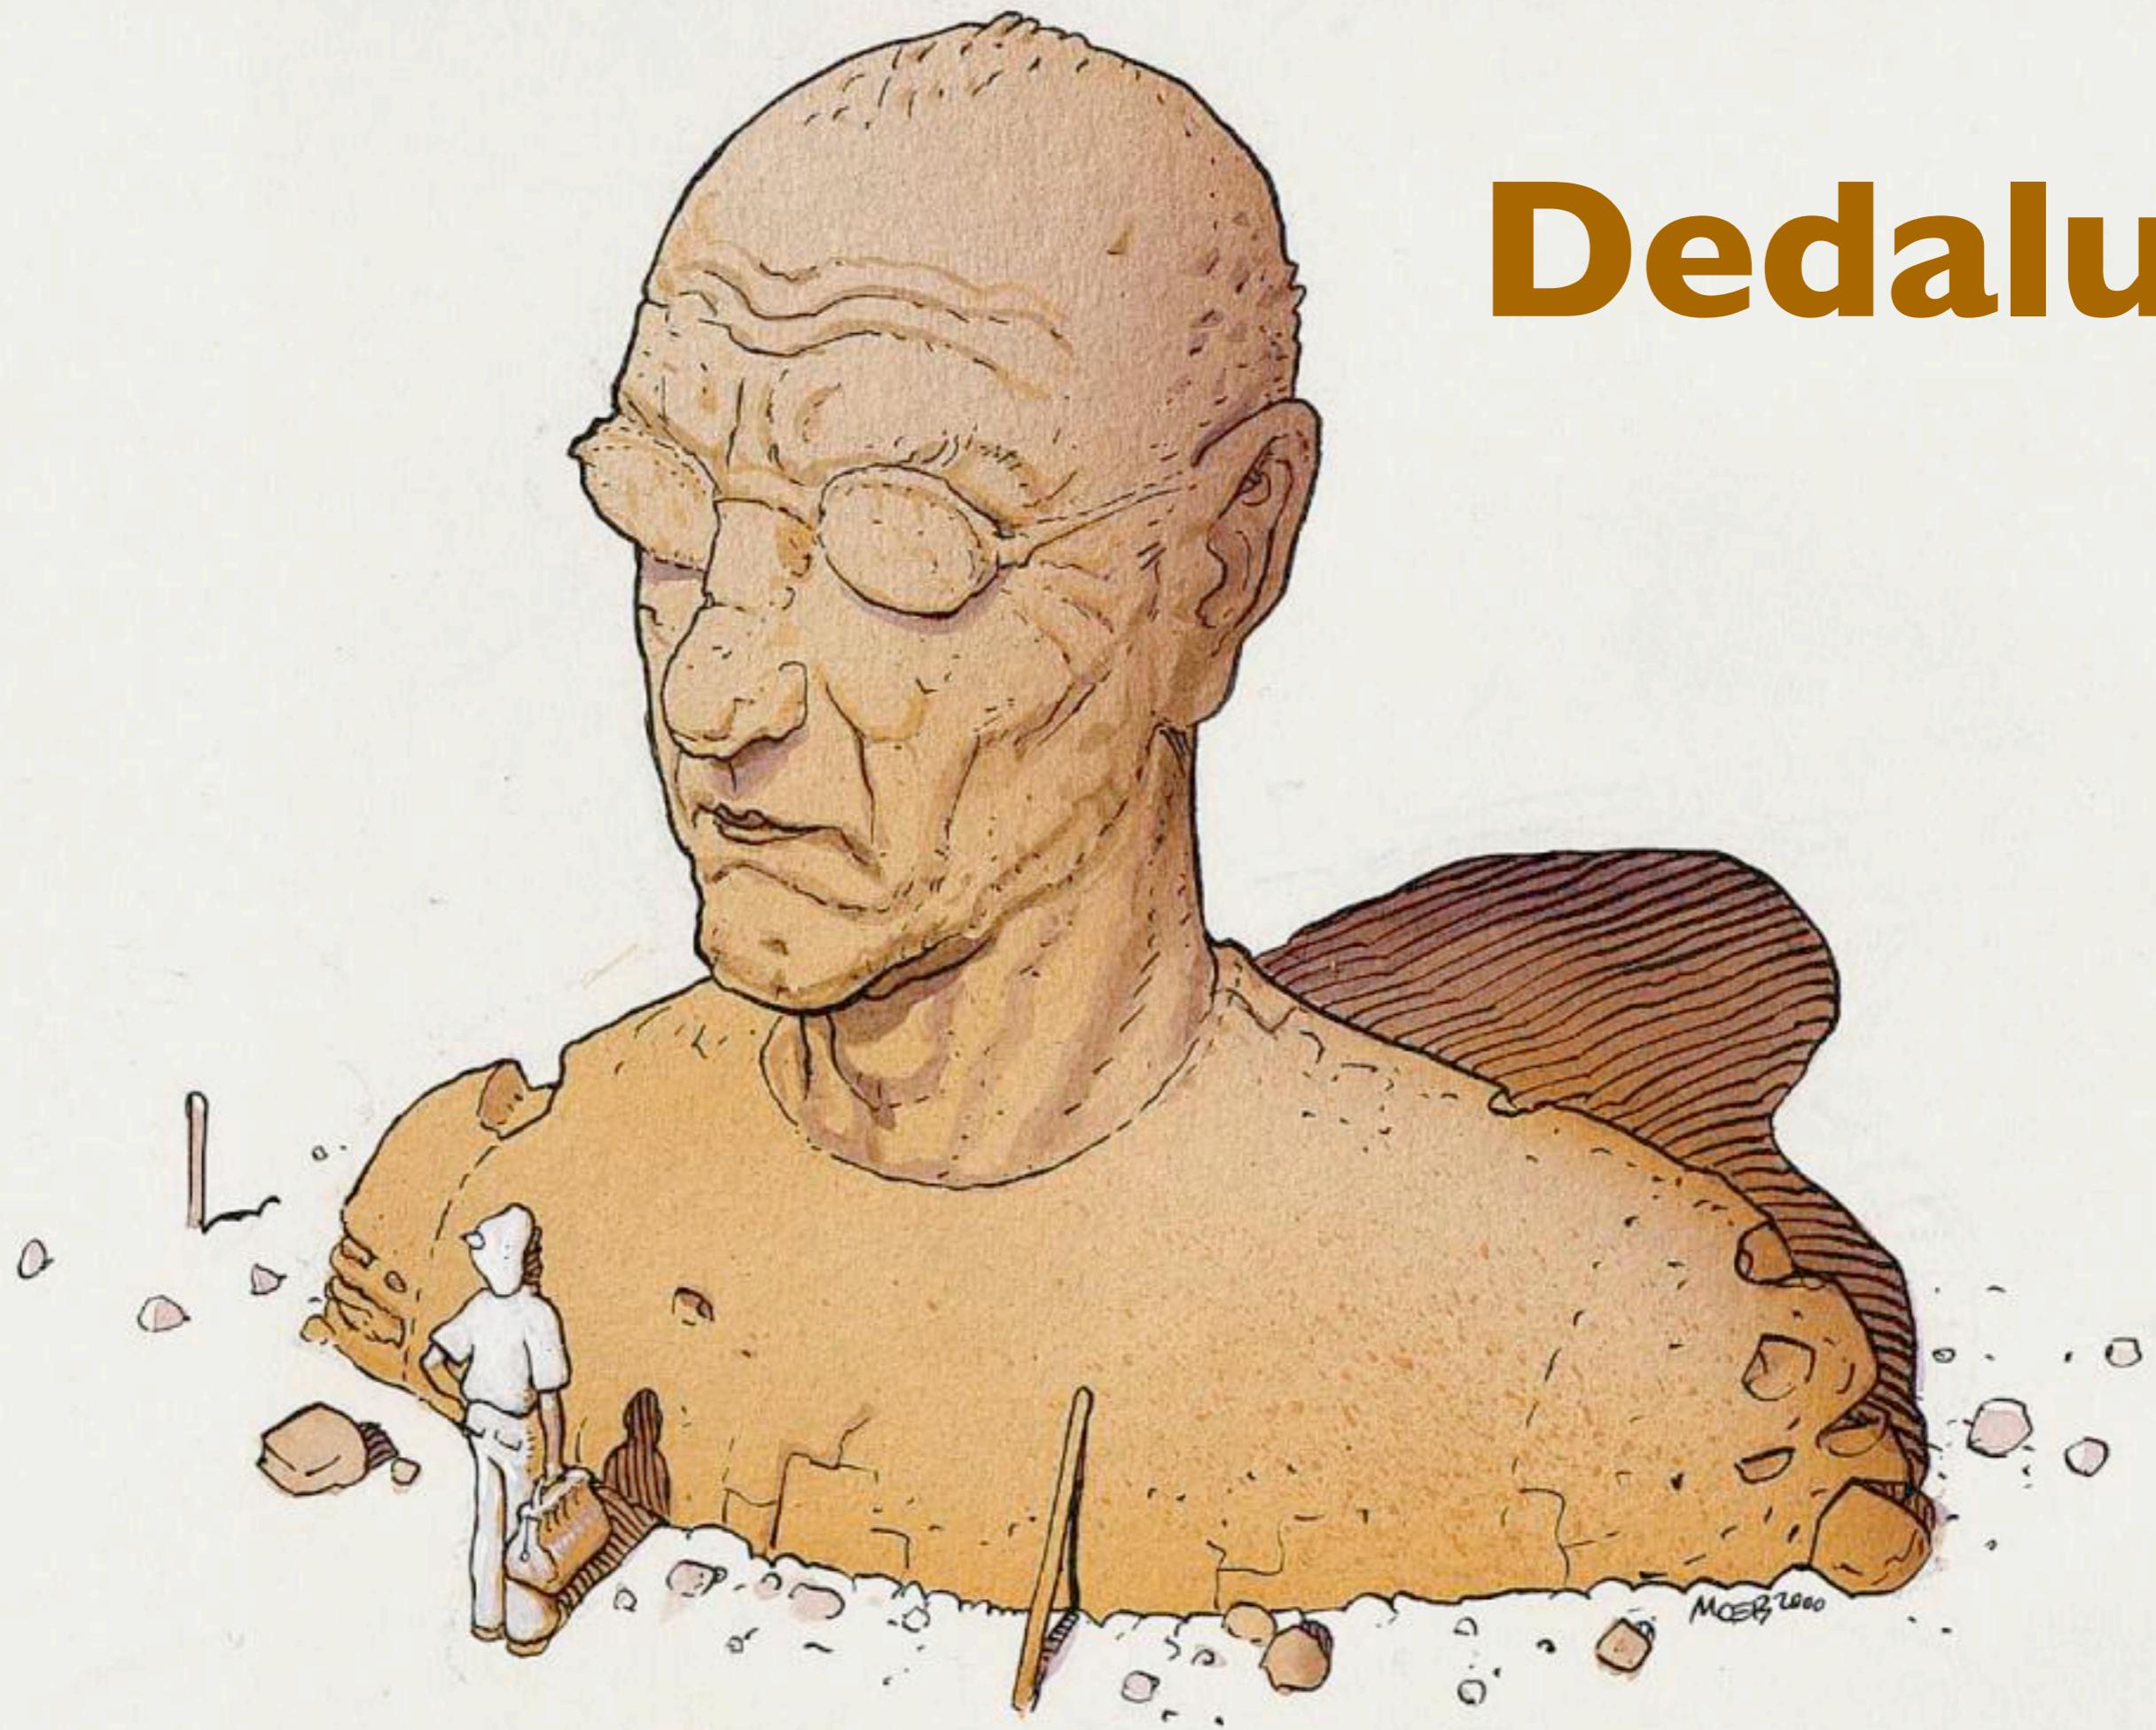
```
(q '[:find ?food
      :where [?person :person/name "John"]
              [?person :favorite/food ?food]
              (d/since database #inst "2012-01-01T00:00:00.000-00:00"))
;=> [[:pho]])
```

- You can obtain the value of the db as-of, or since, a point in time, or both
  - without parameterizing your logic with a time argument
- You can also get the entire history of an entity!

# Datomic is also

- Fully navigable lazy entity maps
- Query across databases
- Optimistic and pessimistic concurrency
- “Upsertting”
- <http://datomic.com>

# Dedalus



# Datalog plus

- Time
- State via rules

# Time

```
loves(john, pizza, <last week>).
```

- Tick model
- Time is an element of the tuple

# Deductive time

```
loves(john, pizza) :- order(john, Food).
```

- “Right now”
- All terms have the same time...

```
loves(john, pizza, <now>) :- order(john, Food, <now>),  
<now> = <now>.
```

# Inductive time

```
loves(john, pizza)@next :- order(john, Food).
```

- “Some other time”
- Next time tick

```
loves(john, pizza, T) :- order(john, Food, S).  
successor(T, S).
```

# Async time

```
loves(john, pizza)@async :- order(john, Food).
```

- Unreliable network

# State

```
loves_pos(john, pizza).
```

- At time tick 0
- Facts

```
loves(john, pizza)@0;
```

# Update

```
loves_pos(john, pho).
```

- At time tick 100
- Facts

```
loves(john, pizza)@0;  
loves(john, pho)@100;
```

# Update

```
loves_neg(john, pizza).
```

- At time tick 300
- Facts

```
loves(john, pizza)@0;  
loves(john, pho)@100;  
loves_neg(john, pizza)@300;
```

# Mutable persistence rule

```
loves_pos(Person, Food) :- loves(Person, Food).  
loves_pos(Person, Food)@next :-  
    loves_pos(Person, Food),  
    ~loves_neg(Person, Food).
```

- Update

```
loves(john, pizza)@<last year>;  
loves(john, pho)@<now>;  
loves_neg(john, pizza)@<now>;
```

A photograph of a forest scene. The foreground is dominated by the dark trunks and branches of several tall trees. Sunlight filters through the dense canopy of leaves above, creating bright highlights and deep shadows. The overall atmosphere is natural and serene.

# Cascalog

# Datalog plus

- Map/reduce processing
  - Order independence a win

# Cascalog

```
(def people [[ "ben" 35 ]
             [ "jerry" 41 ]])

( ??- (<- [?name ?age]
            (people ?name ?age)
            (< ?age 40)))

;=> (([ "ben" 35]))
```

# Three stages

- Pre-aggregation
- Aggregation
- Post-aggregation



# Pre-aggregation

- Joins aggregator functions
- Applies bindable functions and filters
- Dataflow-esque



# Aggregation

- Partition result tuples along logic variables
- Execute aggregators for each logic var



# Post-aggregation

- Execute the remaining filters and functions on dependent aggregator output



# Example

```
(<- [ ?word]
  ;; pre
  (sentence ?sentence)
  (split ?sentence :> ?word)
  ;; agg
  (c/count ?count)
  ;; post
  (> ?count 5))
```



# Pre-aggregation

```
(<- [ ?word ]
;; pre
(sentence ?sentence)
(split ?sentence :> ?word)
;; agg
(c/count ?count)
;; post
(> ?count 5))
```



# Aggregation

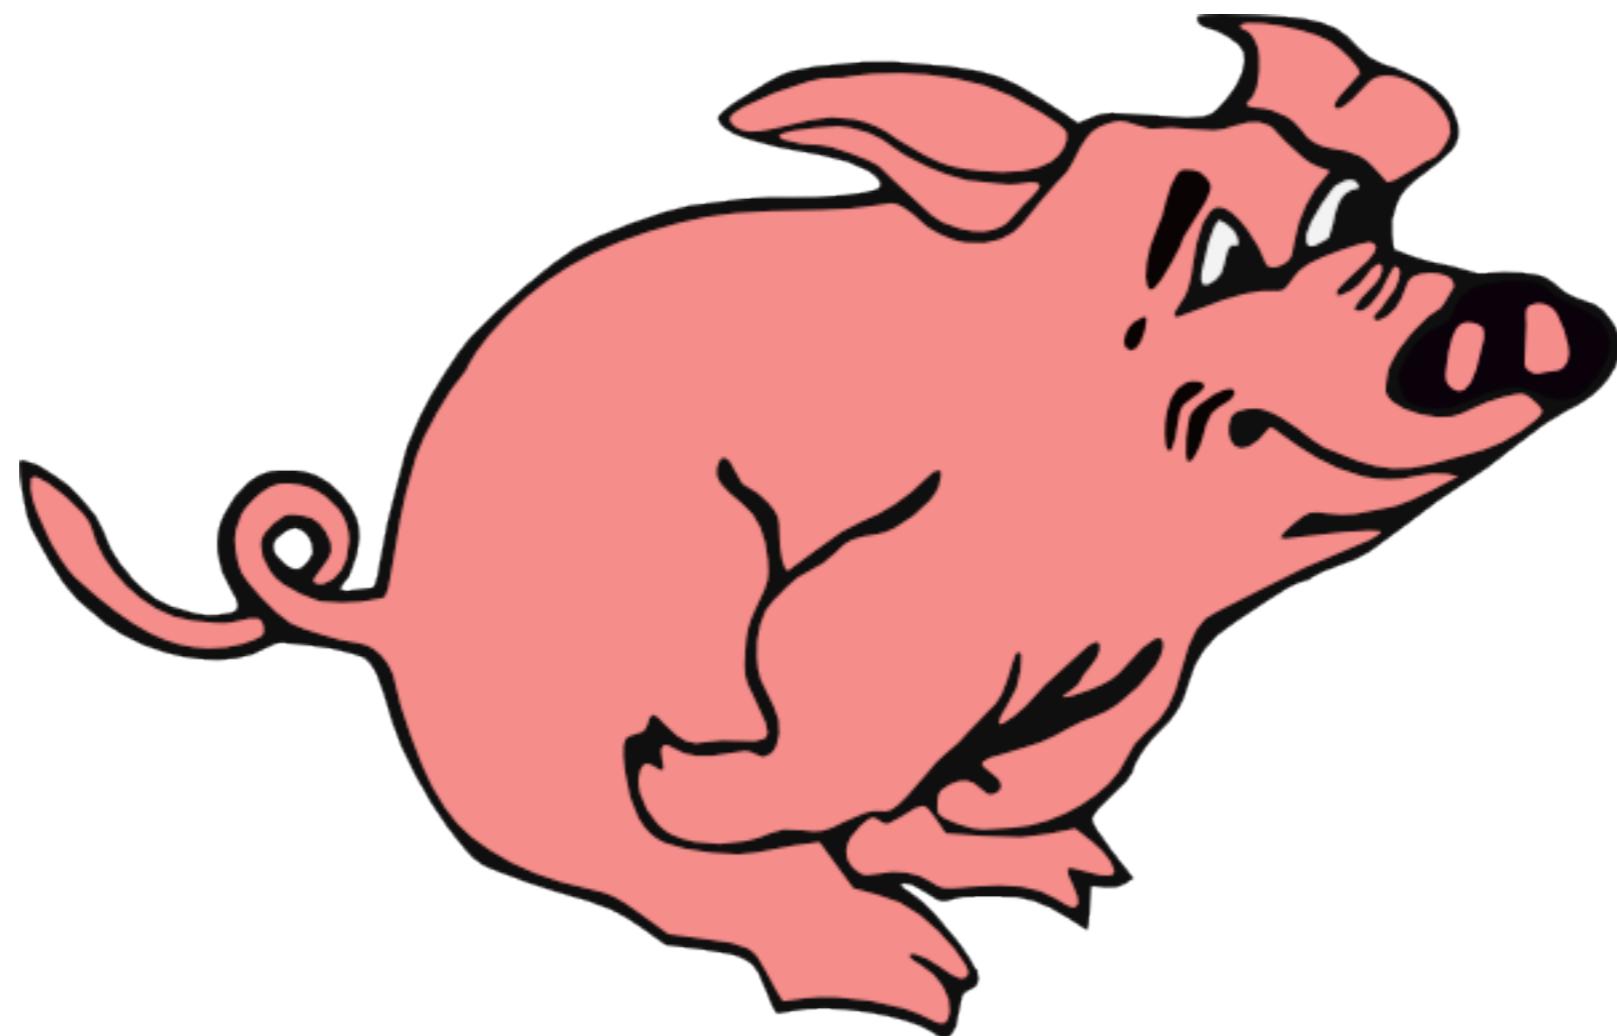
```
(<- [ ?word]
    ;; pre
    (sentence ?sentence)
    (split ?sentence :> ?word)
    ;; agg
    (c/count ?count)
    ;; post
    (> ?count 5))
```



# Post-aggregation

```
(<- [ ?word ]
    ;; pre
    (sentence ?sentence)
    (split ?sentence :> ?word)
    ;; agg
    (c/count ?count)
    ;; post
    (> ?count 5))
```





**Bacwn**

# Datalog plus

- Negation



# Delicious!

# Bacwn facts

```
(def mst3k-db
  (-> mst3k-schema
    (facts {:character/db.id 0 :character/name "Joel" :character/human? true}
           {:character/db.id 1 :character/name "Crow" :character/human? false}
           {:character/db.id 2 :character/name "TV's Frank" :character/human? true}
           {:location/db.id 0 :location/character 0 :location/name "SoL"}
           {:location/db.id 0 :location/character 1 :location/name "SoL"}
           {:location/db.id 1 :location/character 2 :location/name "Gizmonics"}))))
```

# Bacwn rules!

```
(def locate-rule
  (rules-set
    (<- (:stationed-at :location/name ?loc-name :character/name ?char-name)
         (:location :name ?loc-name :character ?char)
         (:character :character/db.id ?char :name ?char-name))))
```

# Bacwn query

```
(q (?- :stationed-at :location/name '??loc :character/name ?char-name)
mst3k-db
locate-rule
{'??loc "SoL"')

;;=> ({:location/name "SoL", :character/name "Crow"}
;;      {:location/name "SoL", :character/name "Joel"})
```

# Bacwn negation

```
(def non-human-locate-rule
  (rules-set
    (<- (:stationed-at :location/name ?loc-name :character/name ?char-name)
         (:location :name ?loc-name :character ?char)
         (:character :character/db.id ?char :name ?char-name)
         (not! :character :character/db.id ?char :human? true))) ;; ADDED NEGATION
```

# Bacwn negation

```
(q (?- :stationed-at :location/name '??loc :character/name ?char-name)
    mst3k-db
    non-human-locate-rule
    {'??loc "SoL"})

;;=> ({:location/name "SoL", :character/name "Crow"})
```

- Find all non-humans at a location

# Open questions

- Query plans
- Optimizations

# Query plans

- The holy grail of DBs is that they **do the right thing**
- Query plans
- No runtime guarantees

# Gaming the query

- /\*+ Hinting \*/
- Prolog - Ordering for termination
- Datalog - Ordering for speed

# Gaming the query

- `/*+ Hinting */`
- Prolog - Ordering for termination
- Datalog - Ordering for speed

# Wut

```
[ :find ?desc
:where
[?root :num 100]
(descendant ?root ?desc)
[?desc :num 4]]
```

- Find all descendants of root number #100 with value = 4
- SLLLLL000000WWWWWW

# Wat

```
[ :find ?desc
:where
[?root :num 100]
(descendant ?root ?desc)
[?desc :num ?n] ; INDIRECTION
[(= ?n 4)]] ; CHECK 4-NESS
```

- Find all descendants of root number #100 with value = 4
- FAST!!

# Wat

```
[ :find ?desc
:where
[?root :num 100]
[?desc :num 4]
(descendant ?root ?desc)]
```

- Find all descendants of root number #100 with value = 4
- Most-bound

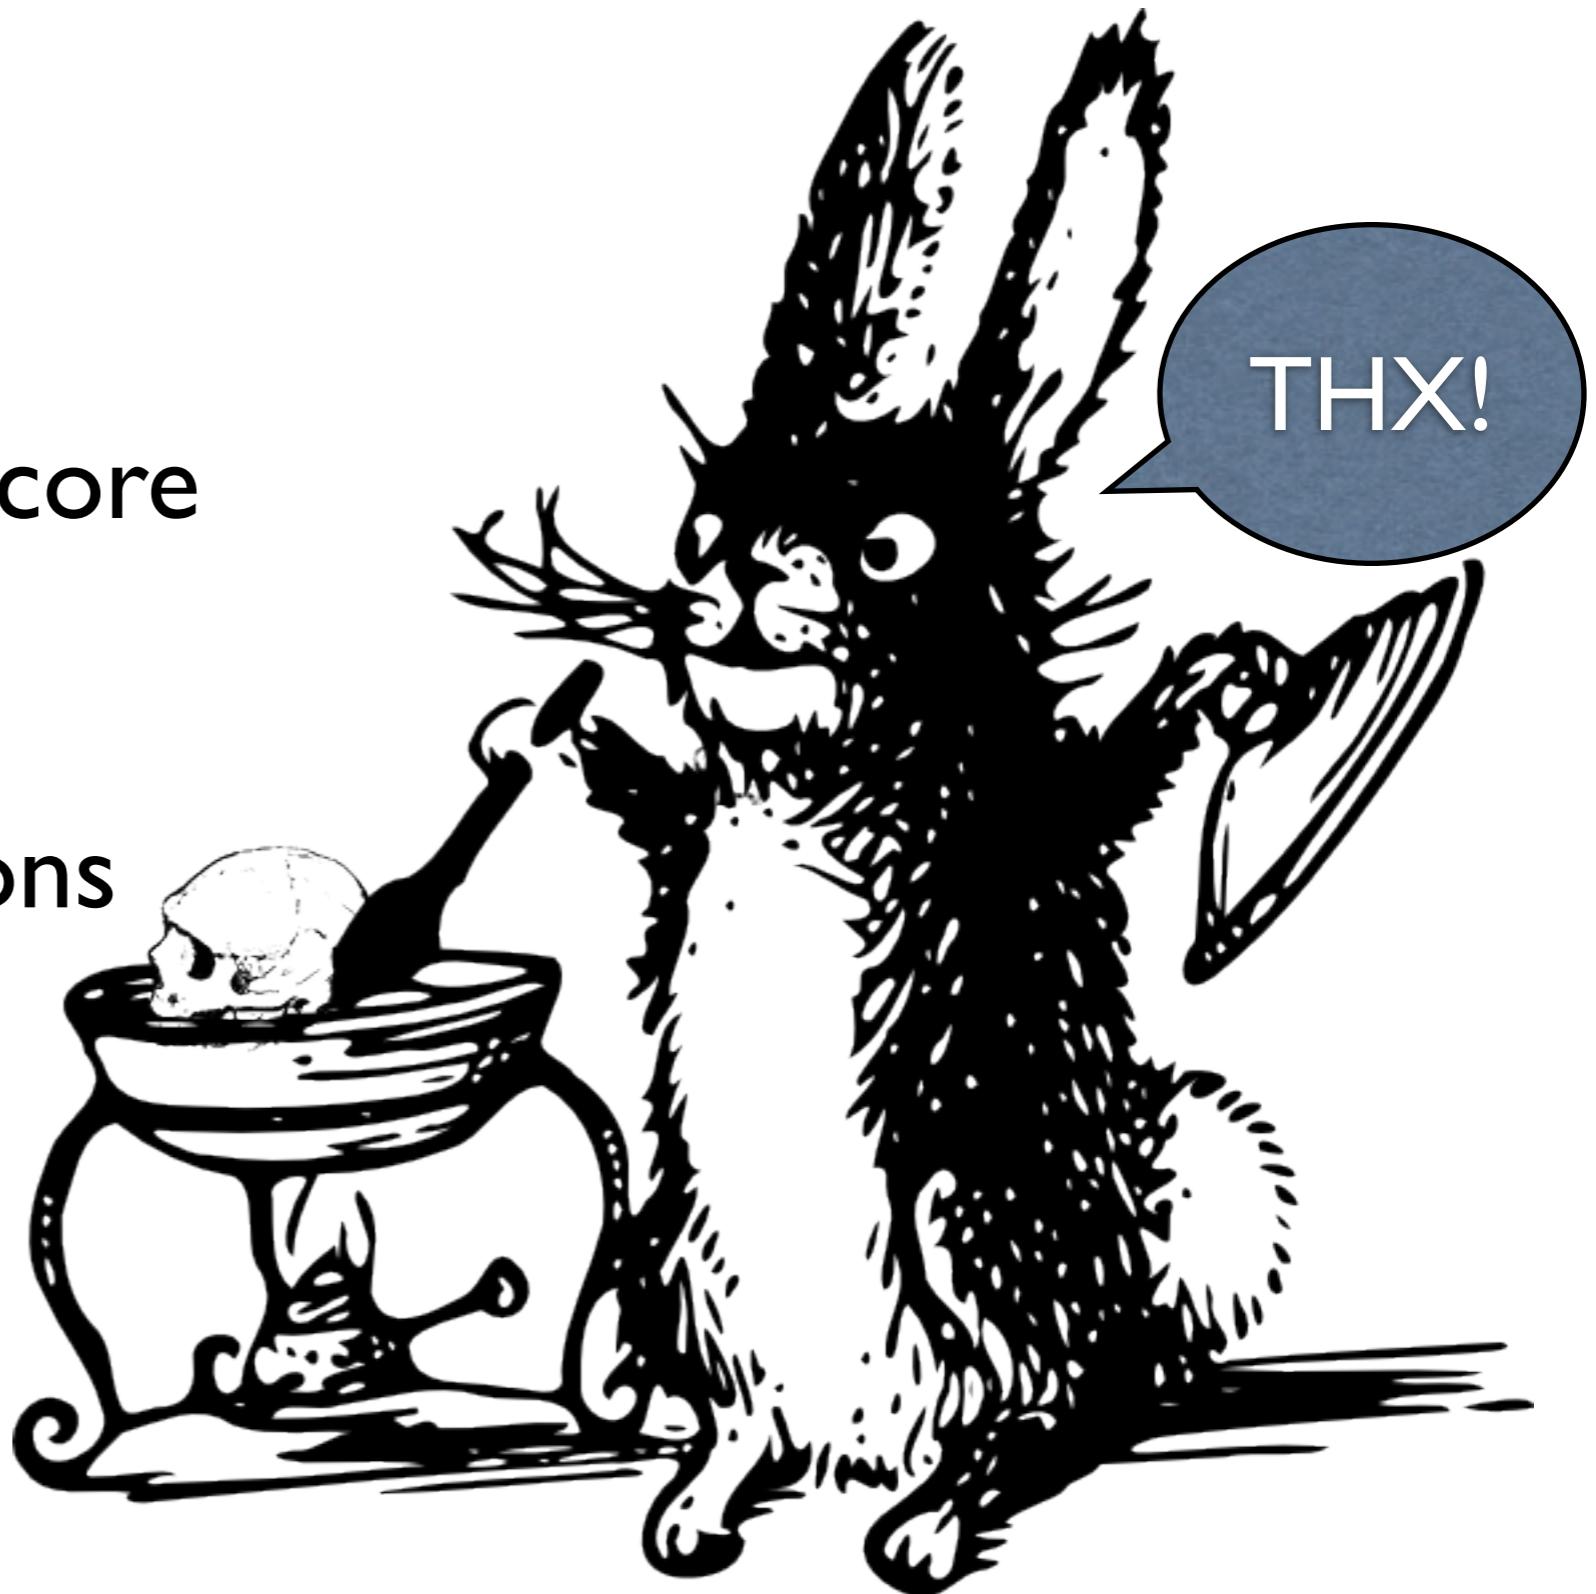
**Order doesn't  
matter**

**Except when  
it does**

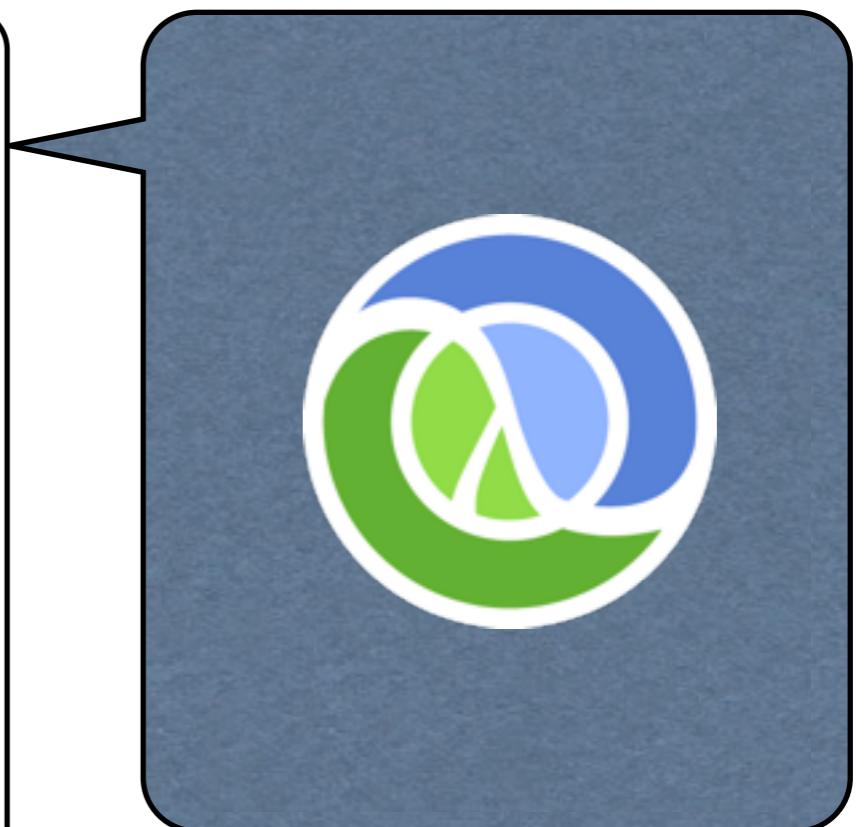
# Pluggable Optimizers

- Order agnostic is a win
- Some orders are better than others
- Plug in your own optimizer
  - That knows your data
  - Will not affect other Datalog engine optimization techniques

- Rich Hickey, Stu Halloway, Clojure/core
- Clojure/dev
- Manning Publications
- The fam
- You



**double-secret slides**



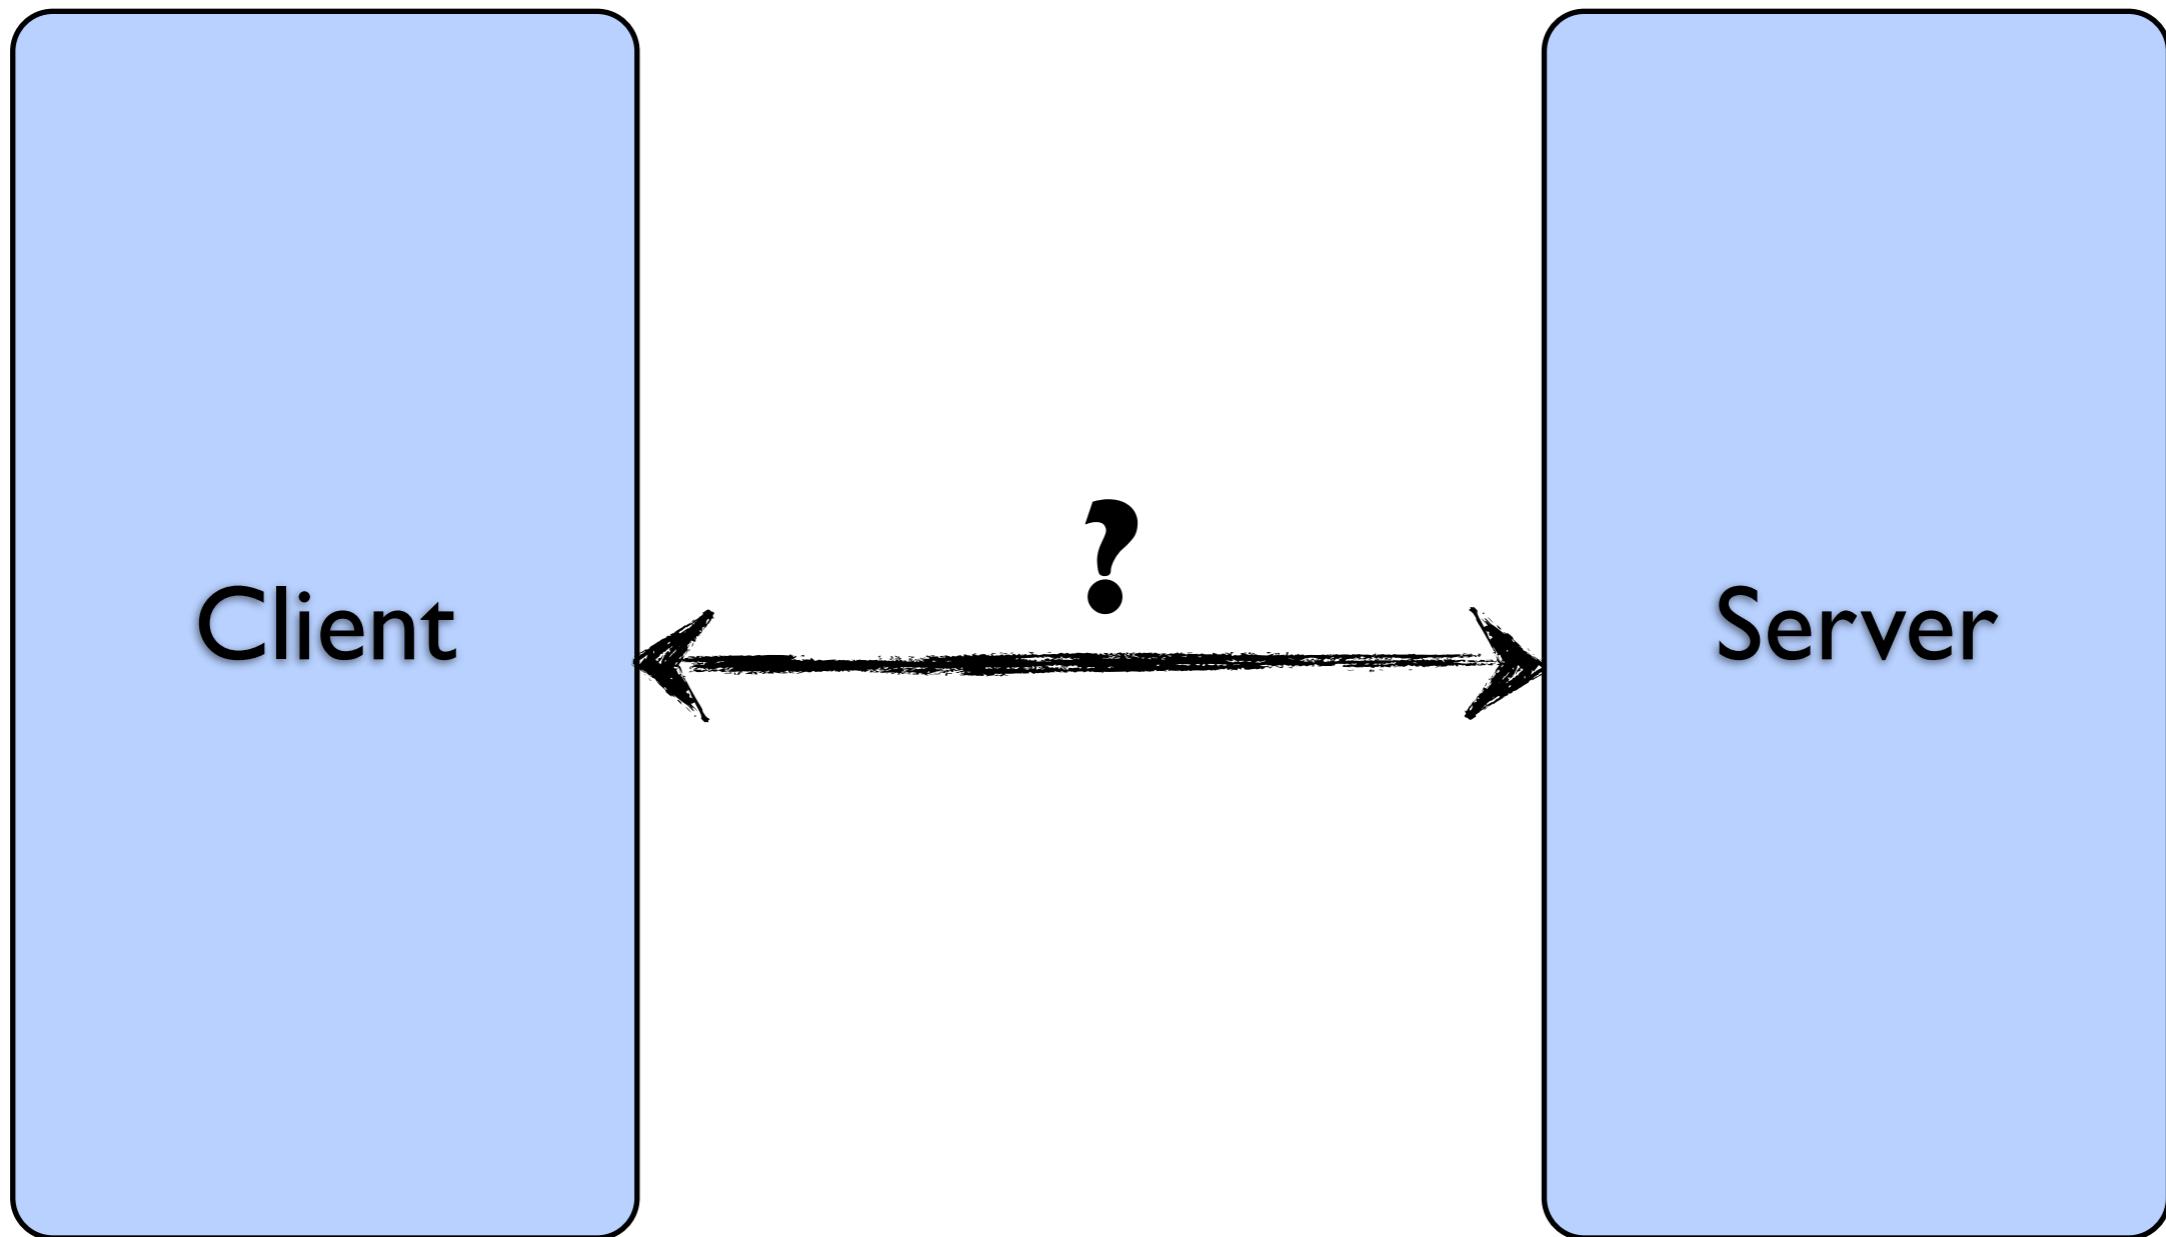
Data as Data.

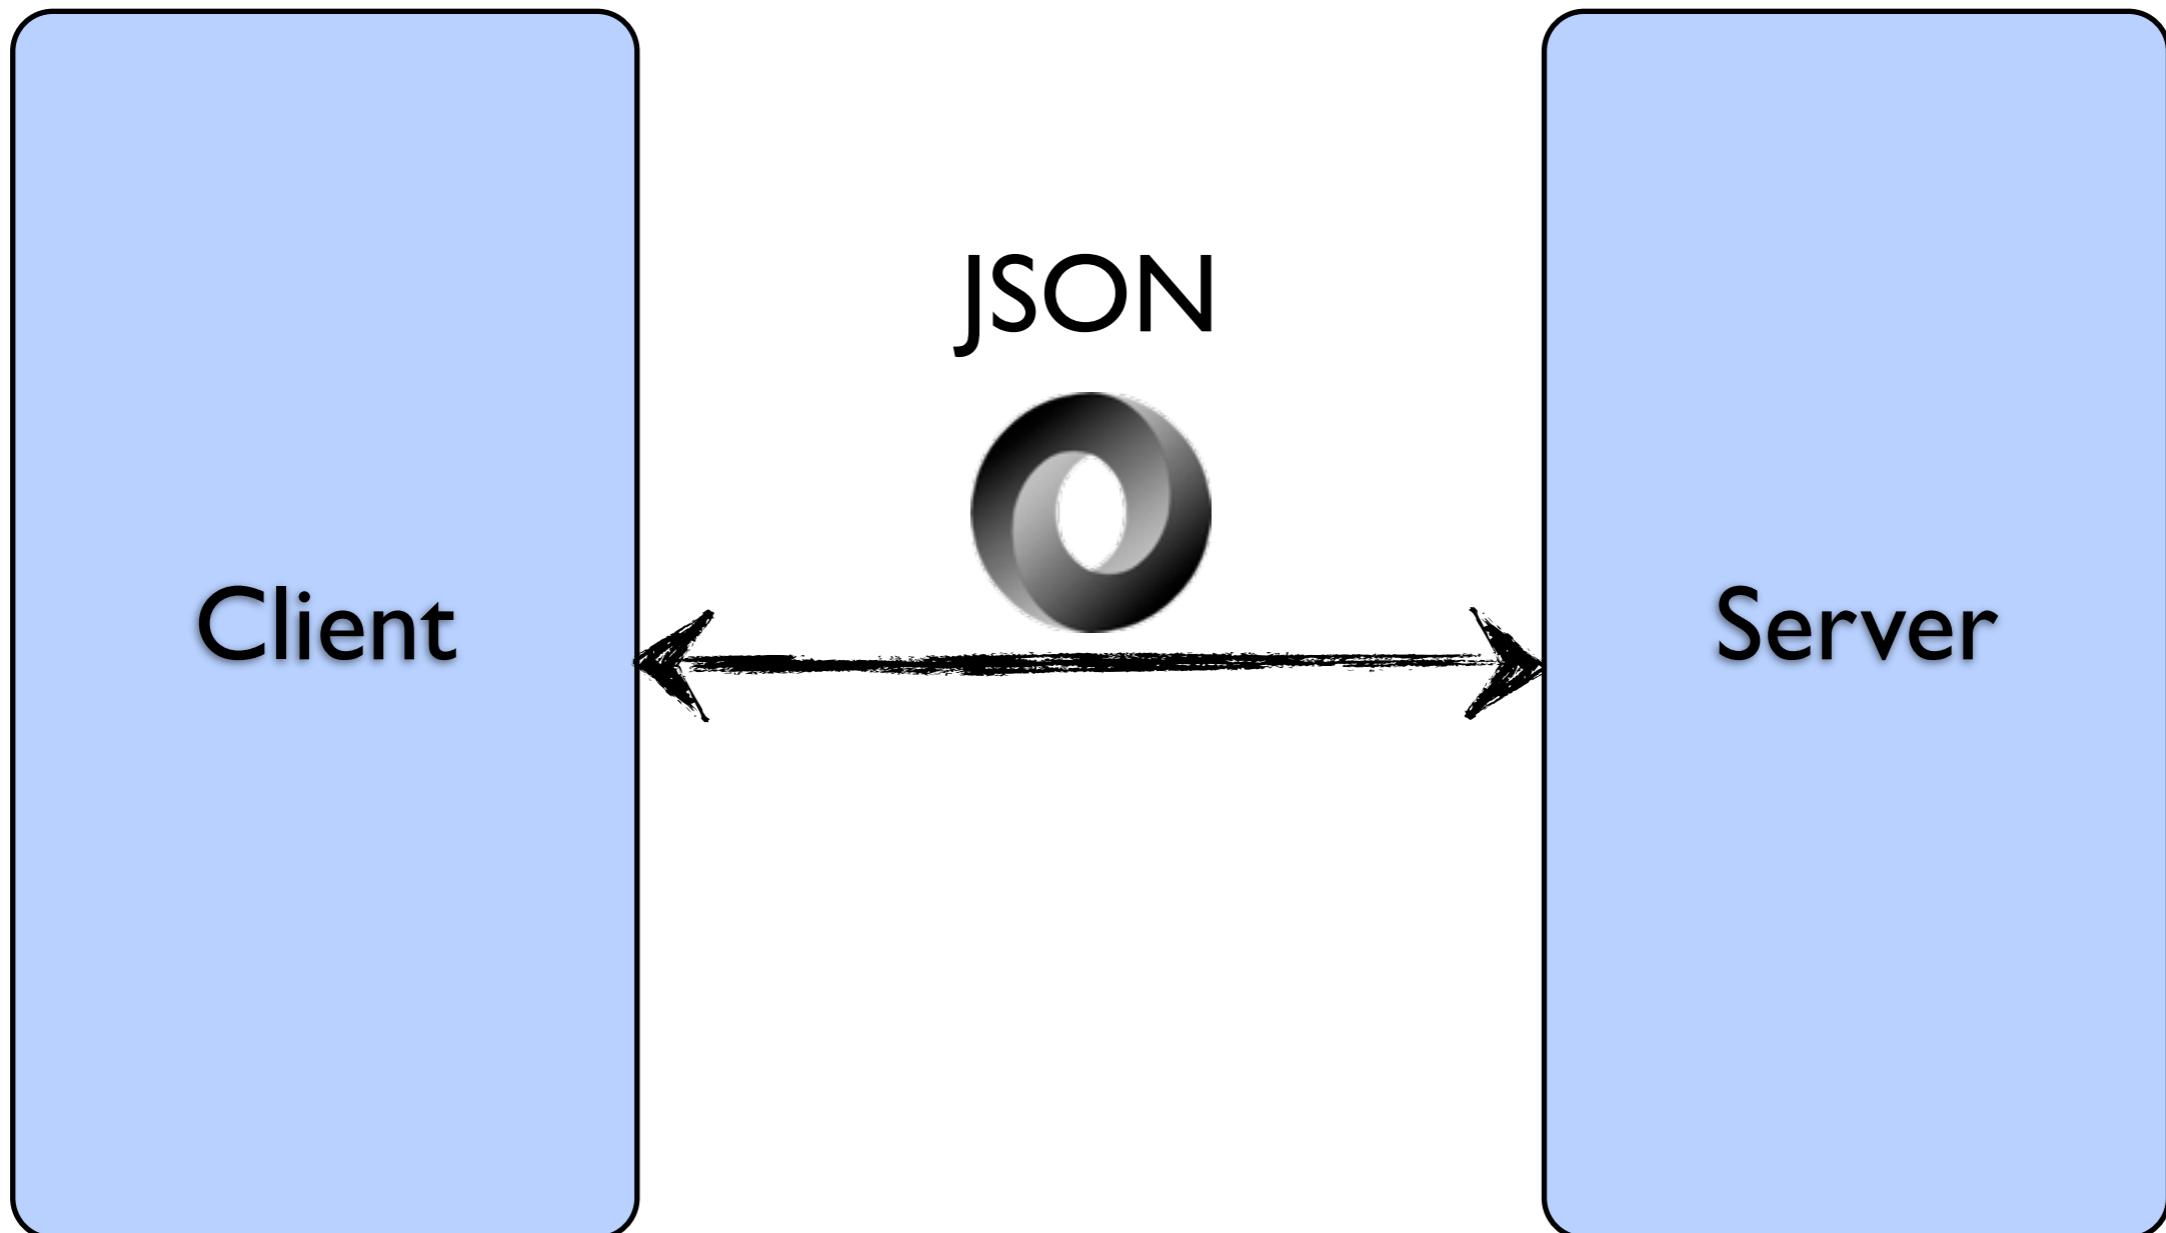
# Scalar

type	in	example
string	–	"hi"
int	–	42
float	–	3.14159
char	Scheme	#\Z
boolean	Scheme	#t
nullity	Scala	null
nothingness	JavaScript	undefined
ratio	Clojure	22/7
symbol	Ruby	:hi
keyword	Clojure	:hi

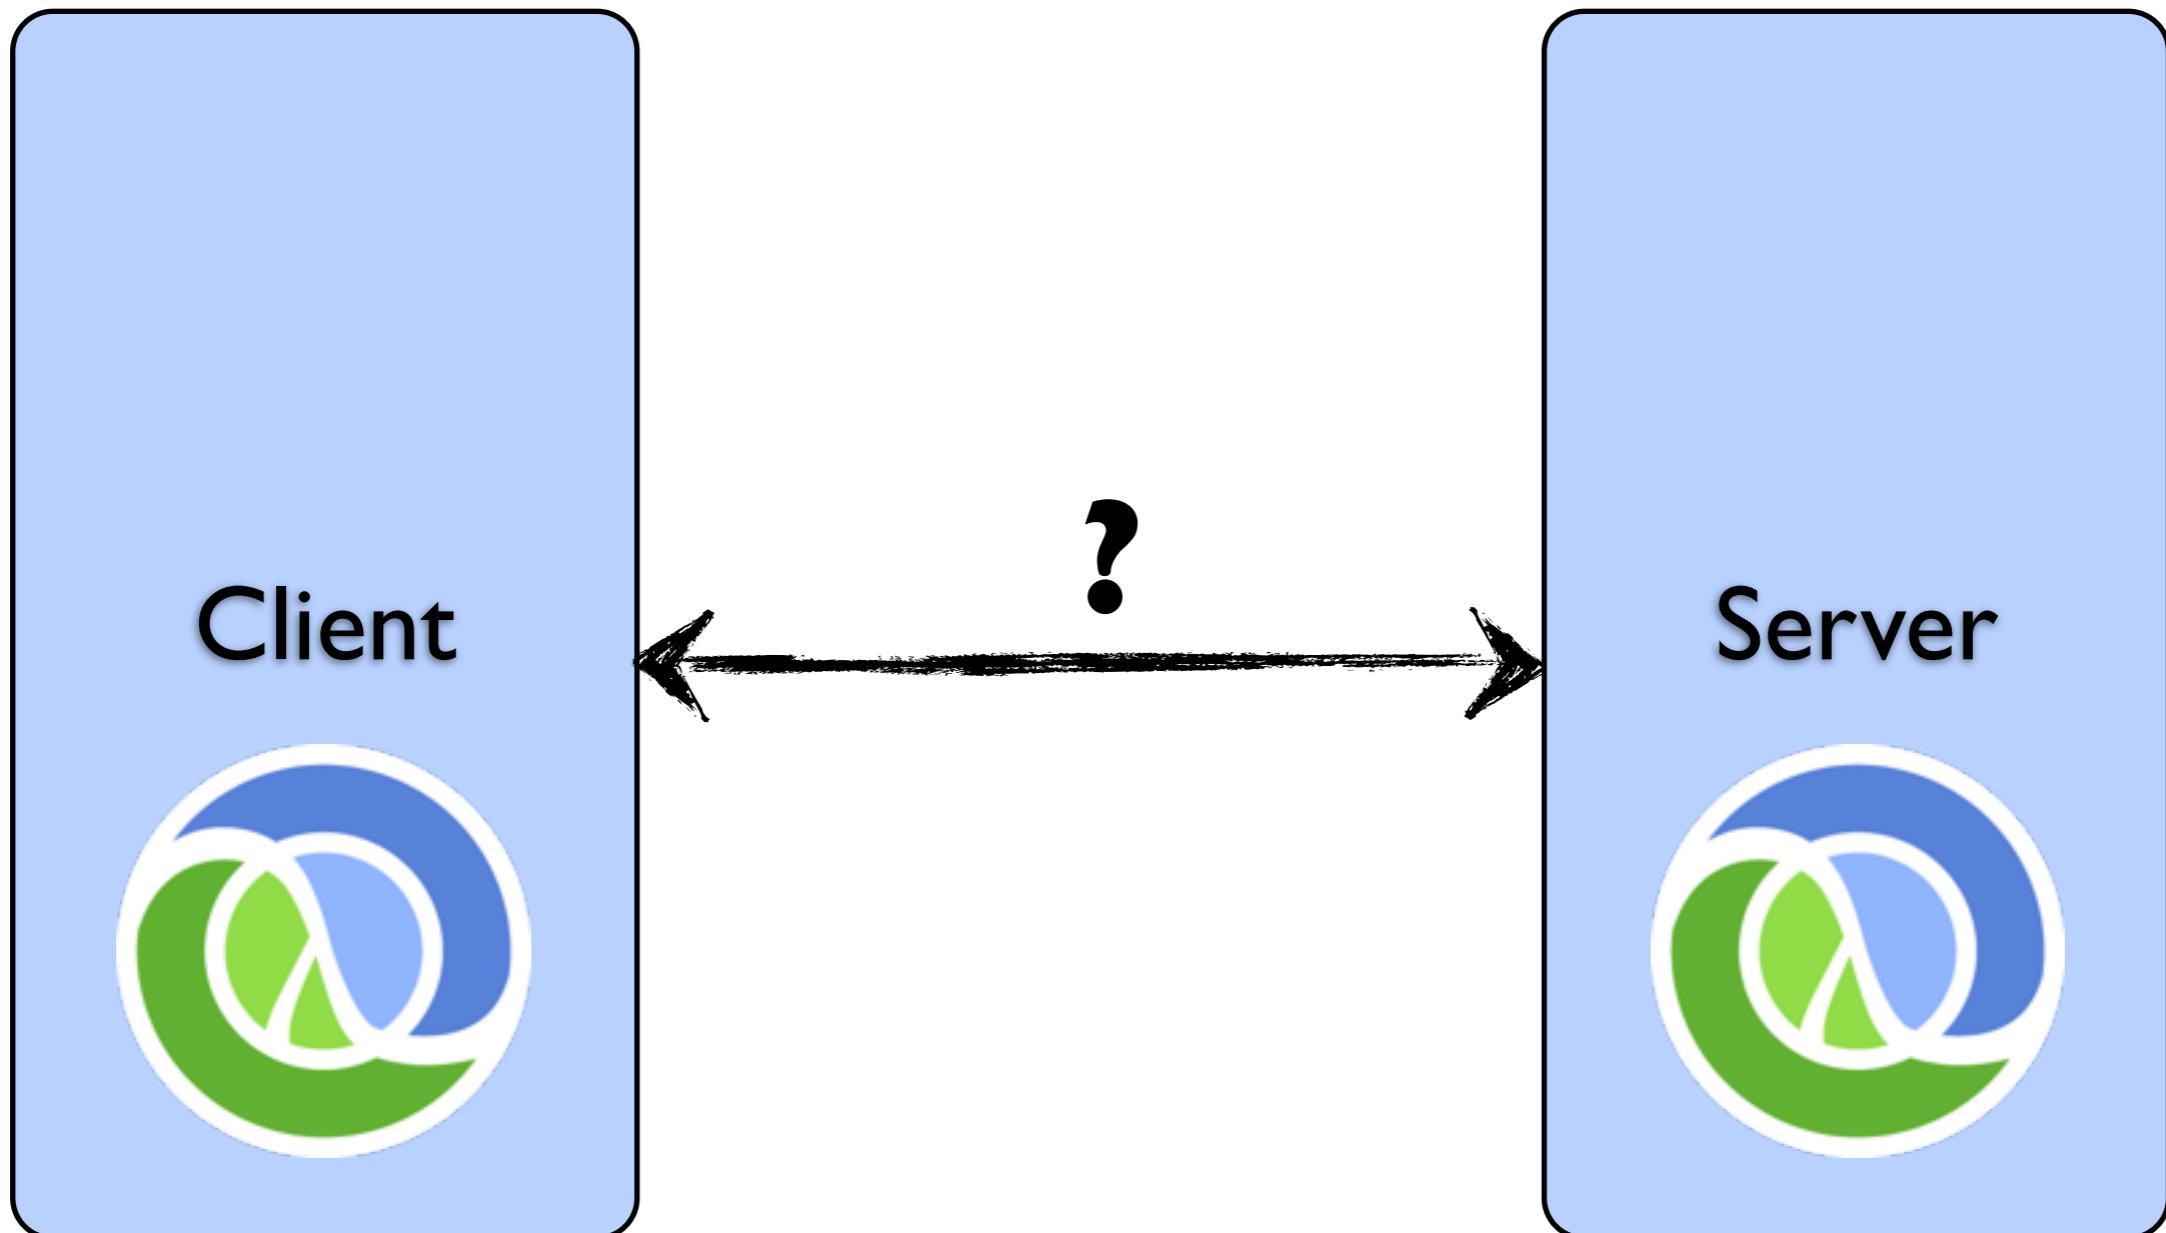
# Scalar

type	in	example
list	Haskell	[1,2,3,4]
tuple	Erlang	{1,2,3,4}
vector	Clojure	[1 2 3 4]
dict	Python	{'a': 1, 'b': 2}
set	Clojure	#{}{1 2 3}
array	Pike	({ 1, 2, 3 })
table	J	22\$1234
queue	Clojure	#queue [1 2 3]

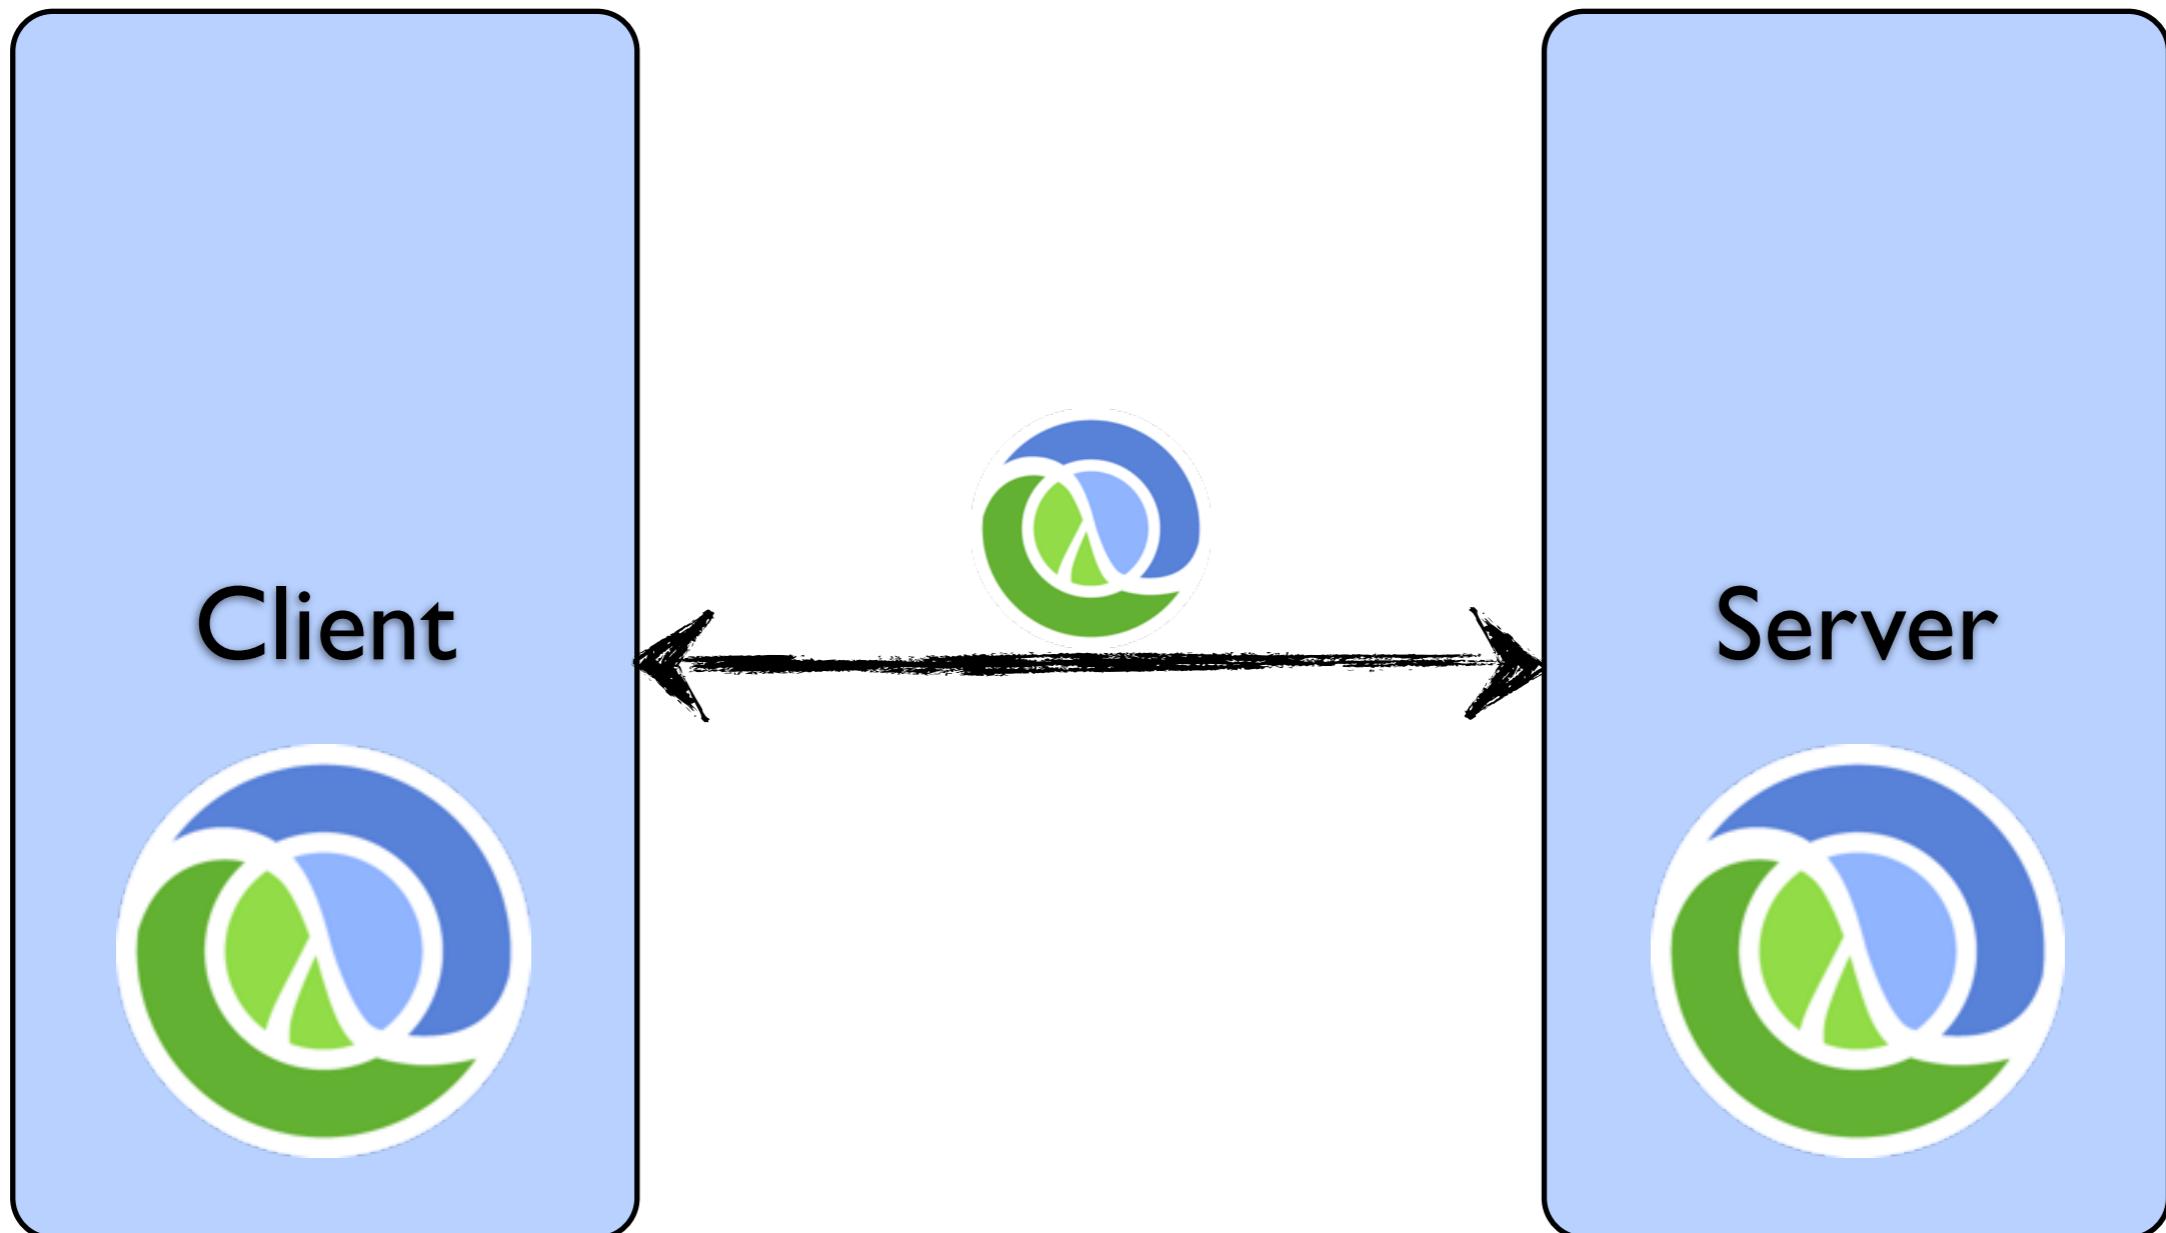




# **JSON as Data**



# Clojure as Code



# Clojure as Data

# Data

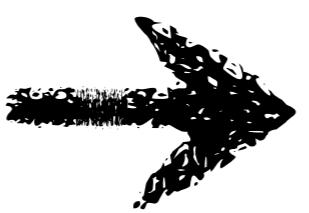
# as Code.

**Engine**

**Specification**

**Engine**

**Specification**



**Engine**

**Specification**



**Engine**

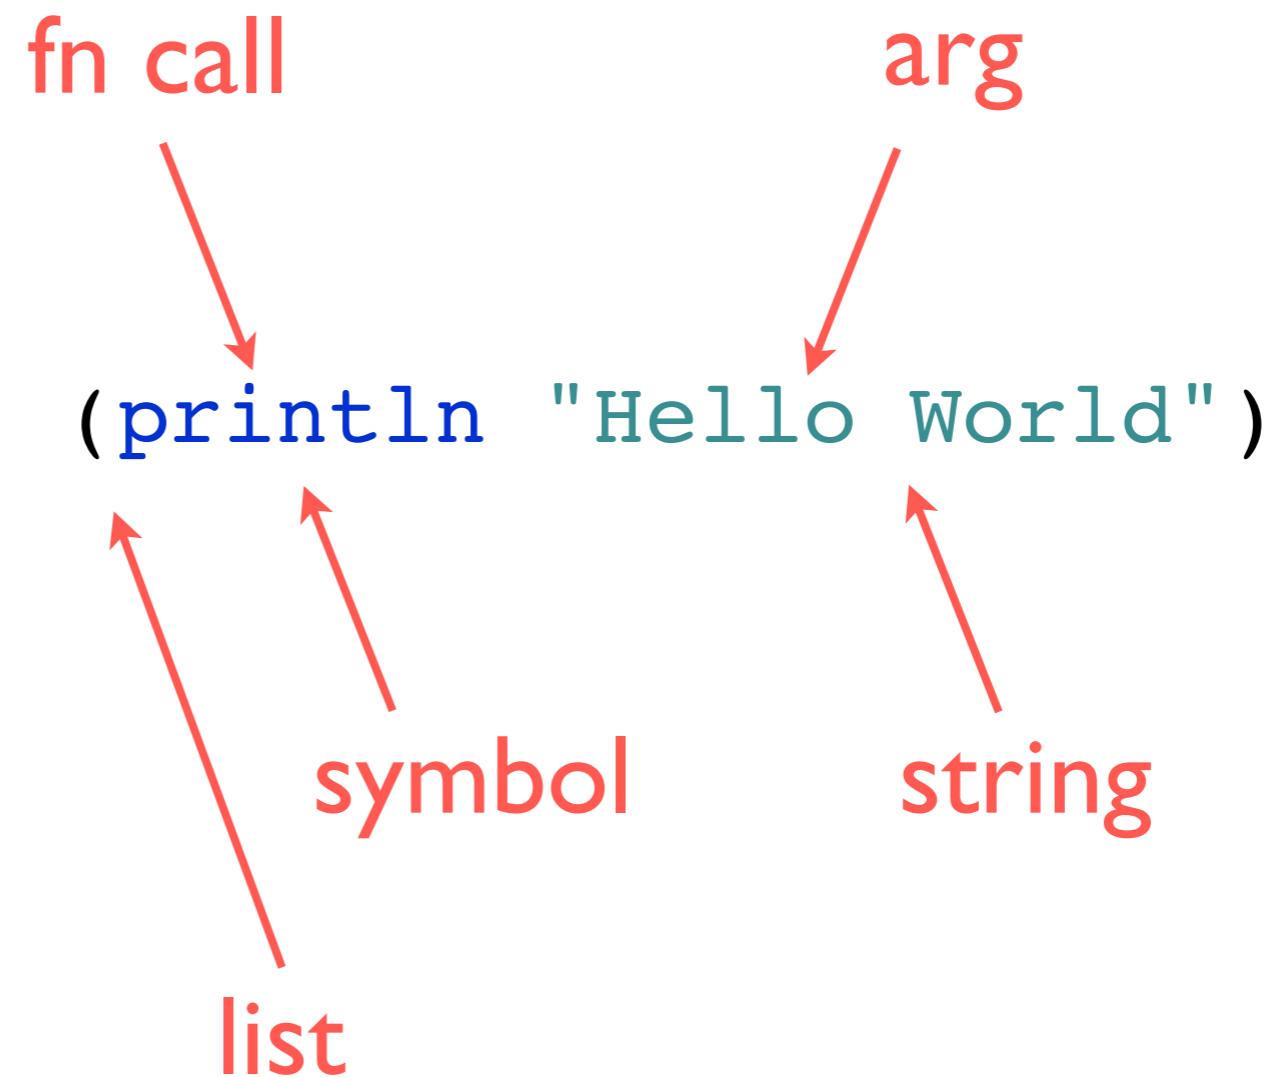
I'm really  
doing  
something  
cool  
now



Code  
as Data.

# Function call

semantics:



structure:

# Function definition

```
(defn greet
  "Returns a friendly greeting"
  [your-name]
  (str "Hello, " your-name))
```

define a fn      fn name      docstring  
arguments      fn body

# It's all data

```
(defn greet
  "Returns a friendly greeting"
  [your-name]
  (str "Hello, " your-name))
```

symbol                    symbol                    string

vector

list

The diagram illustrates the structure of a Clojure function definition. It points to several parts of the code with red arrows:

- A red arrow labeled "symbol" points to the opening parenthesis of the defn form: "(defn".
- A red arrow labeled "symbol" points to the closing parenthesis of the docstring: ")".
- A red arrow labeled "string" points to the text within the docstring: "Returns a friendly greeting".
- A red arrow labeled "vector" points to the opening parenthesis of the argument list: "[".
- A red arrow labeled "list" points to the closing parenthesis of the argument list: "]".