

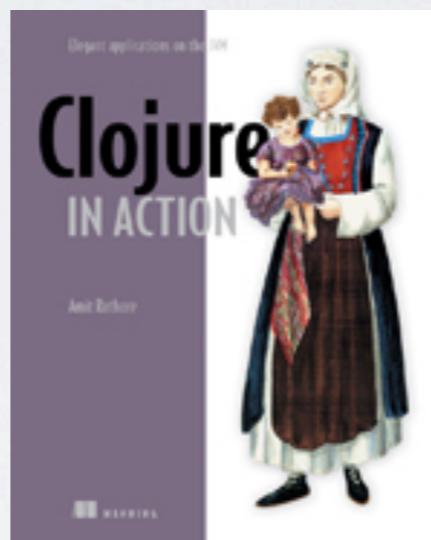
clojure
+ datomic
+ storm

= **zolodeck**



Amit Rathore

Startup Guy



amit@zololabs.com



software applications today



<insert buzzword>



big data



NoSQL



machine-learning



real-time



performant



availability



cloud



no ops



web-scale



choices



architecture



stack



language



database



productivity



startup world



delivered yesterday



with 2 engineers



(usually)



must provide all the cool features



must provide all the cool features
(that you think you need)



cheaper to operate than what customers pay



“pivot” on a dime



Right?



design



design to scale



performance vs scalability



typical path



simple beginnings



a web-server,
a database-server



Low complexity, no overhead
Quick development, agility, lots of features
No redundancy, low operational cost

non-trivial number of users



more web-servers



zolodeck

Add some for redundancy, for performance
Optimize the DB,
add redundancy at the DB level
Still fairly simple from the app standpoint

more web-servers
scale up the database



zolodeck

Add some for redundancy, for performance
Optimize the DB,
add redundancy at the DB level
Still fairly simple from the app standpoint

more growth



caching content



reverse proxy, fancy load balancers
may start needing some re-coding of the app

caching content
more web-servers



reverse proxy, fancy load balancers
may start needing some re-coding of the app

caching content
more web-servers
single database-server not enough



reverse proxy, fancy load balancers
may start needing some re-coding of the app

caching content
more web-servers

single database-server not enough



zolodeck

reverse proxy, fancy load balancers
may start needing some re-coding of the app
incidental complexity

caching content
more web-servers

single database-server not enough

may need some re-coding



reverse proxy, fancy load balancers
may start needing some re-coding of the app

database replication



- multiple DB servers can handle high read volumes (load is spread across)
- writes: not so much – all the replicas must write the same data

database replication

reads



zolodeck

- multiple DB servers can handle high read volumes (load is spread across)
- writes: not so much – all the replicas must write the same data

database replication

reads
writes?



- multiple DB servers can handle high read volumes (load is spread across)
- writes: not so much – all the replicas must write the same data

even more growth



memcached



zolodeck

- replication starts taking too long
- certain features get their own DB
- significant re-architecture

memcached CDN



zolodeck

- replication starts taking too long
- certain features get their own DB
- significant re-architecture

memcached
CDN
database partitioning



zolodeck

- replication starts taking too long
- certain features get their own DB
- significant re-architecture

memcached,
CDN,
database partitioning

significant re-architecture



zolodeck

- replication starts taking too long
- certain features get their own DB
- significant re-architecture

even more growth



zolodeck

- hasn't someone done this before?
- why wasn't this designed to scale?
- partition on features – geo, last-names, user-id,
- create user-clusters, all available features on every cluster,
- hashing or master DB to locate cluster

even more growth

Aaaaaa!!



zolodeck

- hasn't someone done this before?
- why wasn't this designed to scale?
- partition on features – geo, last-names, user-id,
- create user-clusters, all available features on every cluster,
- hashing or master DB to locate cluster

even more growth

Aaaaaa!!

Aaaaaaaaaa!!



zolodeck

- hasn't someone done this before?
- why wasn't this designed to scale?
- partition on features – geo, last-names, user-id,
- create user-clusters, all available features on every cluster,
- hashing or master DB to locate cluster

even more growth

Aaaaaa!!

Aaaaaaaaaa!!

rethink everything



zolodeck

- hasn't someone done this before?
- why wasn't this designed to scale?
- partition on features – geo, last-names, user-id,
- create user-clusters, all available features on every cluster,
- hashing or master DB to locate cluster

new features?



- everyone only working on infrastructure

finally...



stability through scalability



zolodeck

- OK performance
- new features can begin
- somewhat manageable

ready for more growth



ready for more growth?



here's the point



scalability vs agility



a real tradeoff



typical path for a reason



typical because it is logical



incidental complexity



a proposed solution

data
processing



hadoop anyone?

simple



simple

and easy



zolodeck



zolodeck

improve life through artificial intelligence



networking is good
(hello Strange Loop)



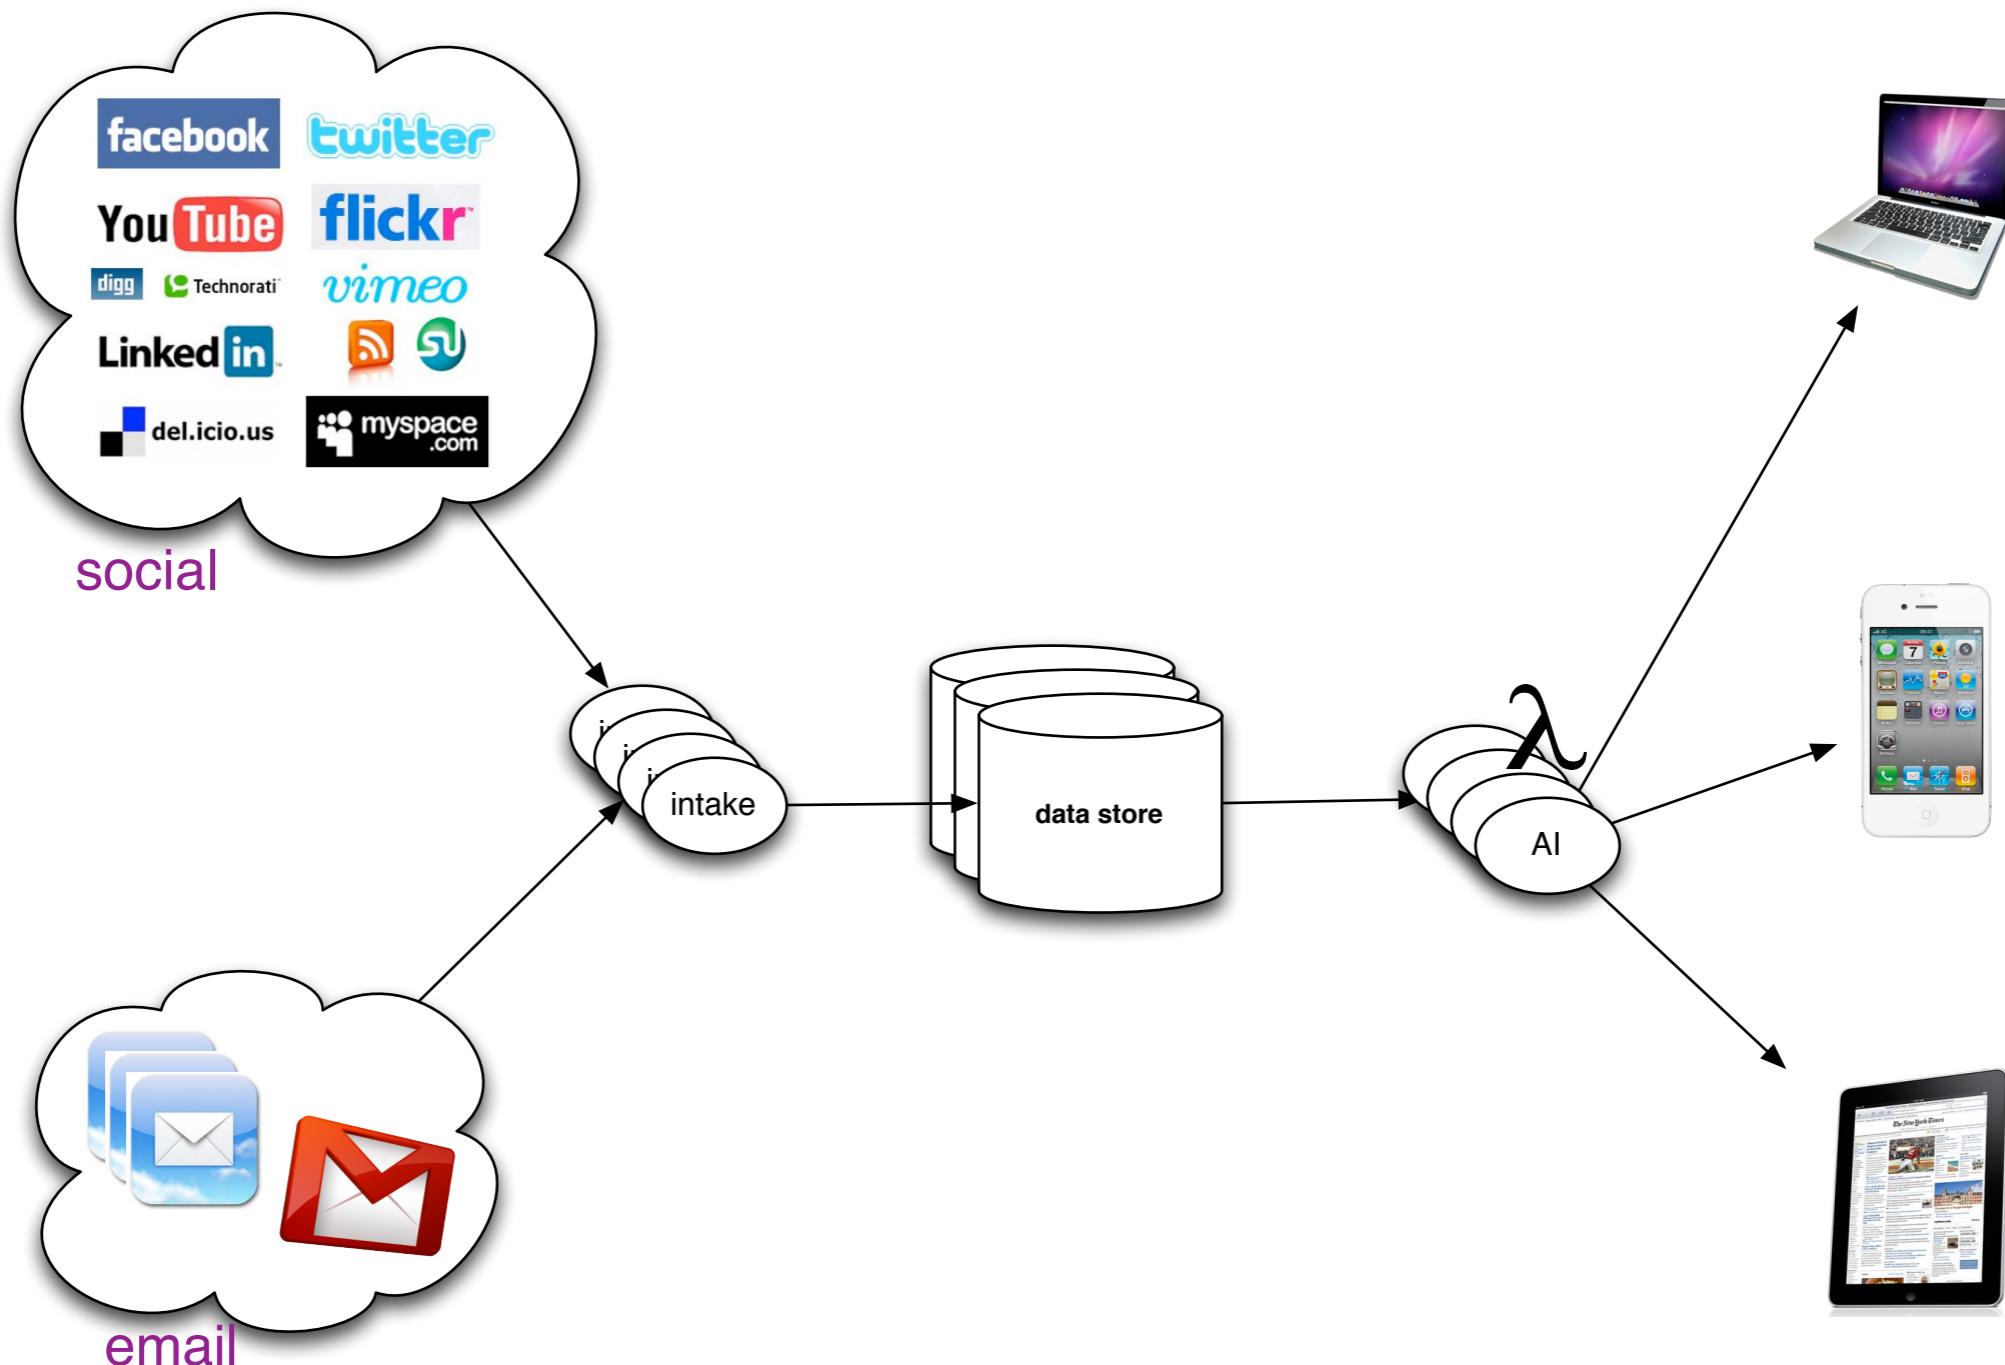
zolodeck - intelligent virtual assistant



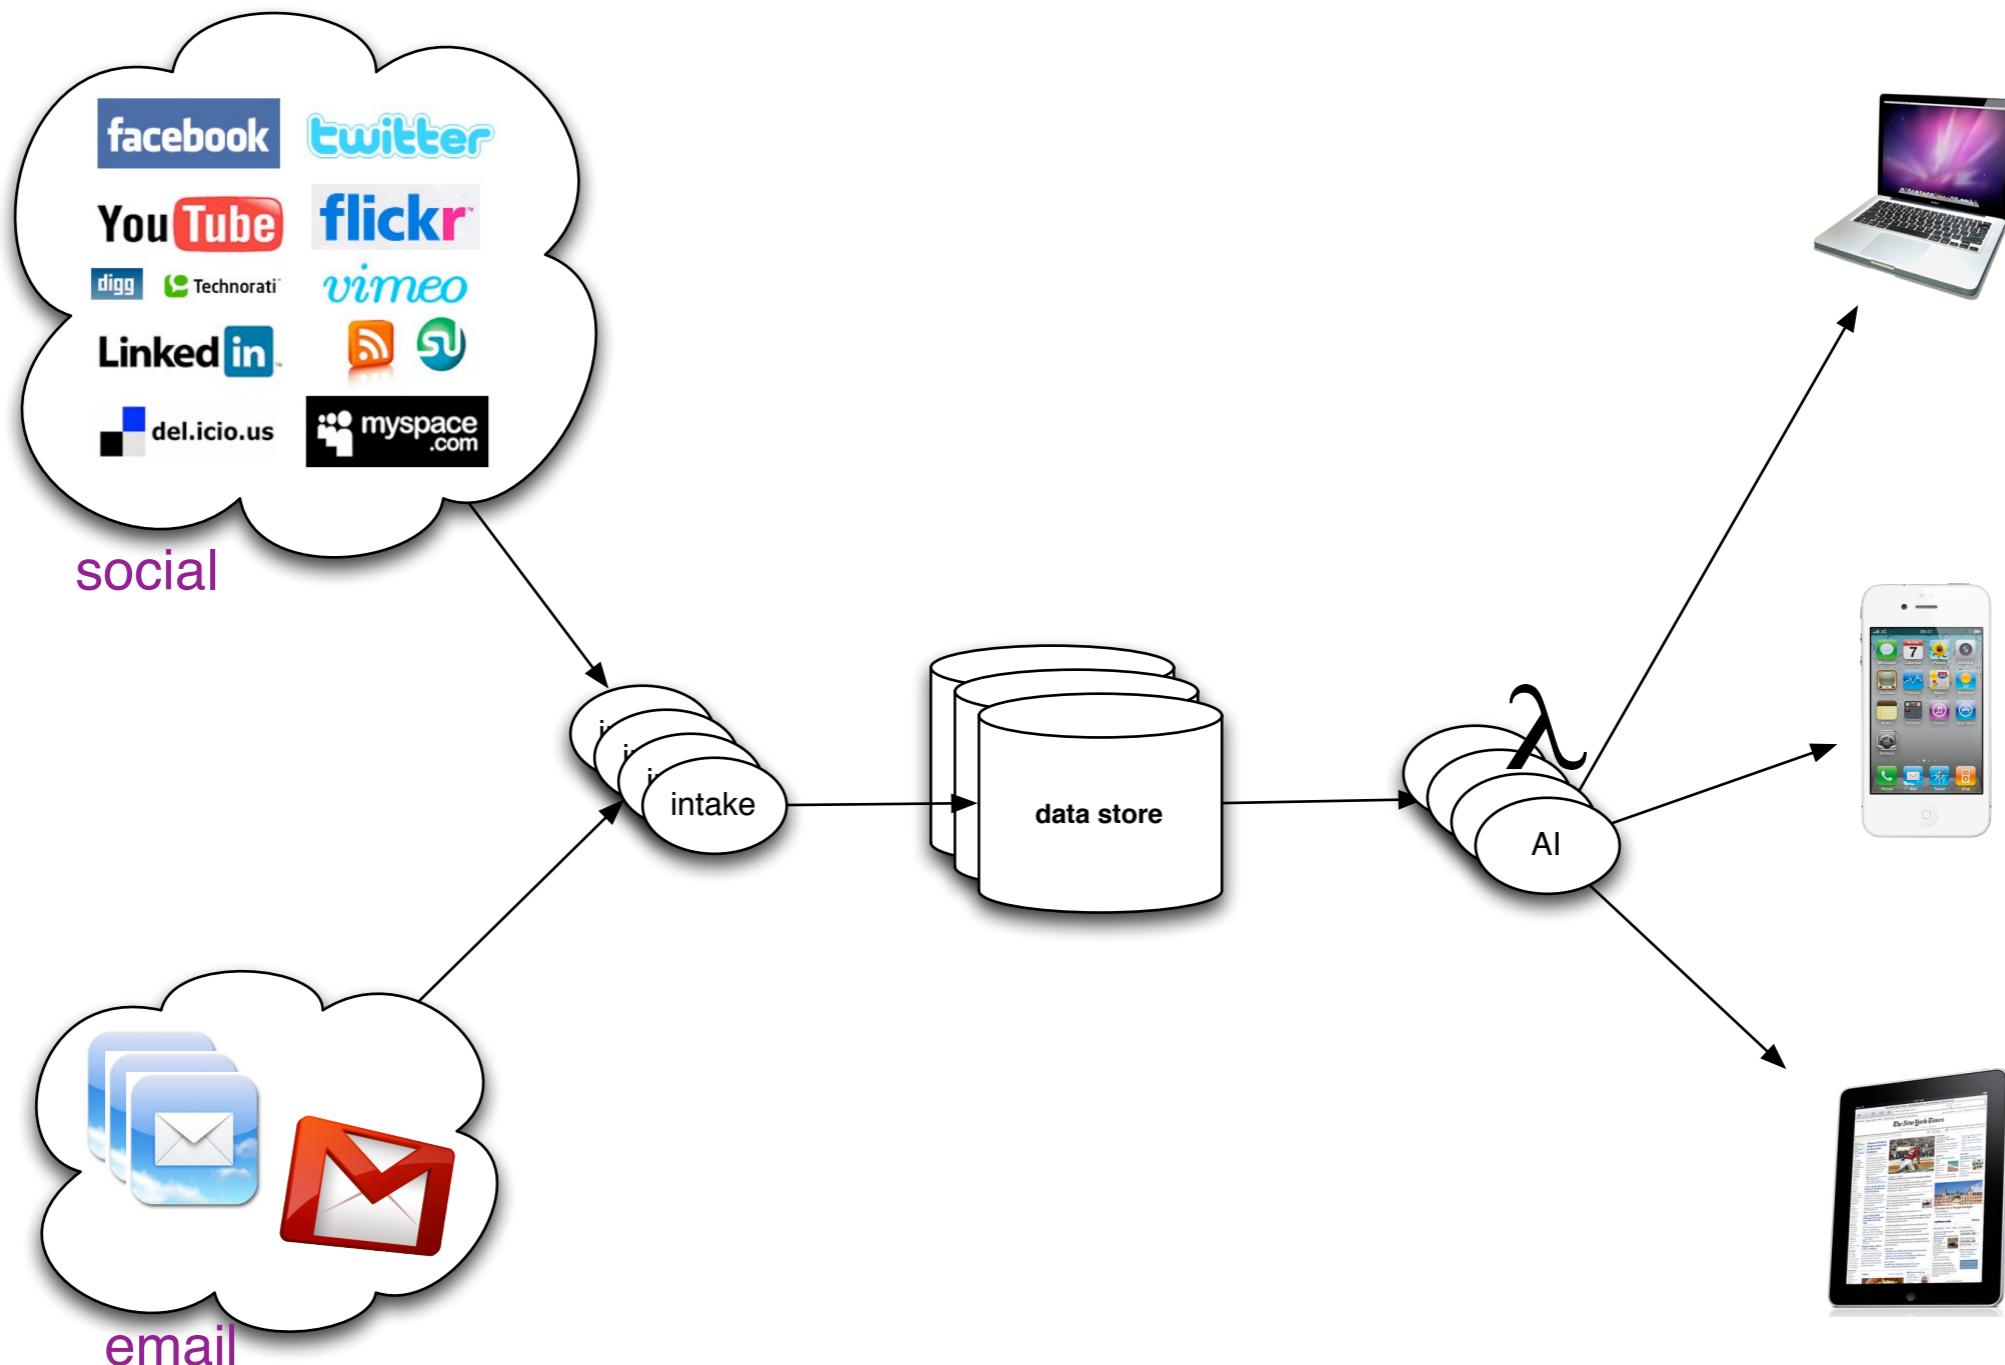
zolodeck

- monitor all social networks, email inboxes
- show where you are with your existing network (lots of metrics)
- watch out for you when a relationship is weakening
- remind you to reach out with the right person, at the right time, with the right content
- help you connect with people you should be connected to, and help you connect people who should be connected

zolodeck



zolodeck

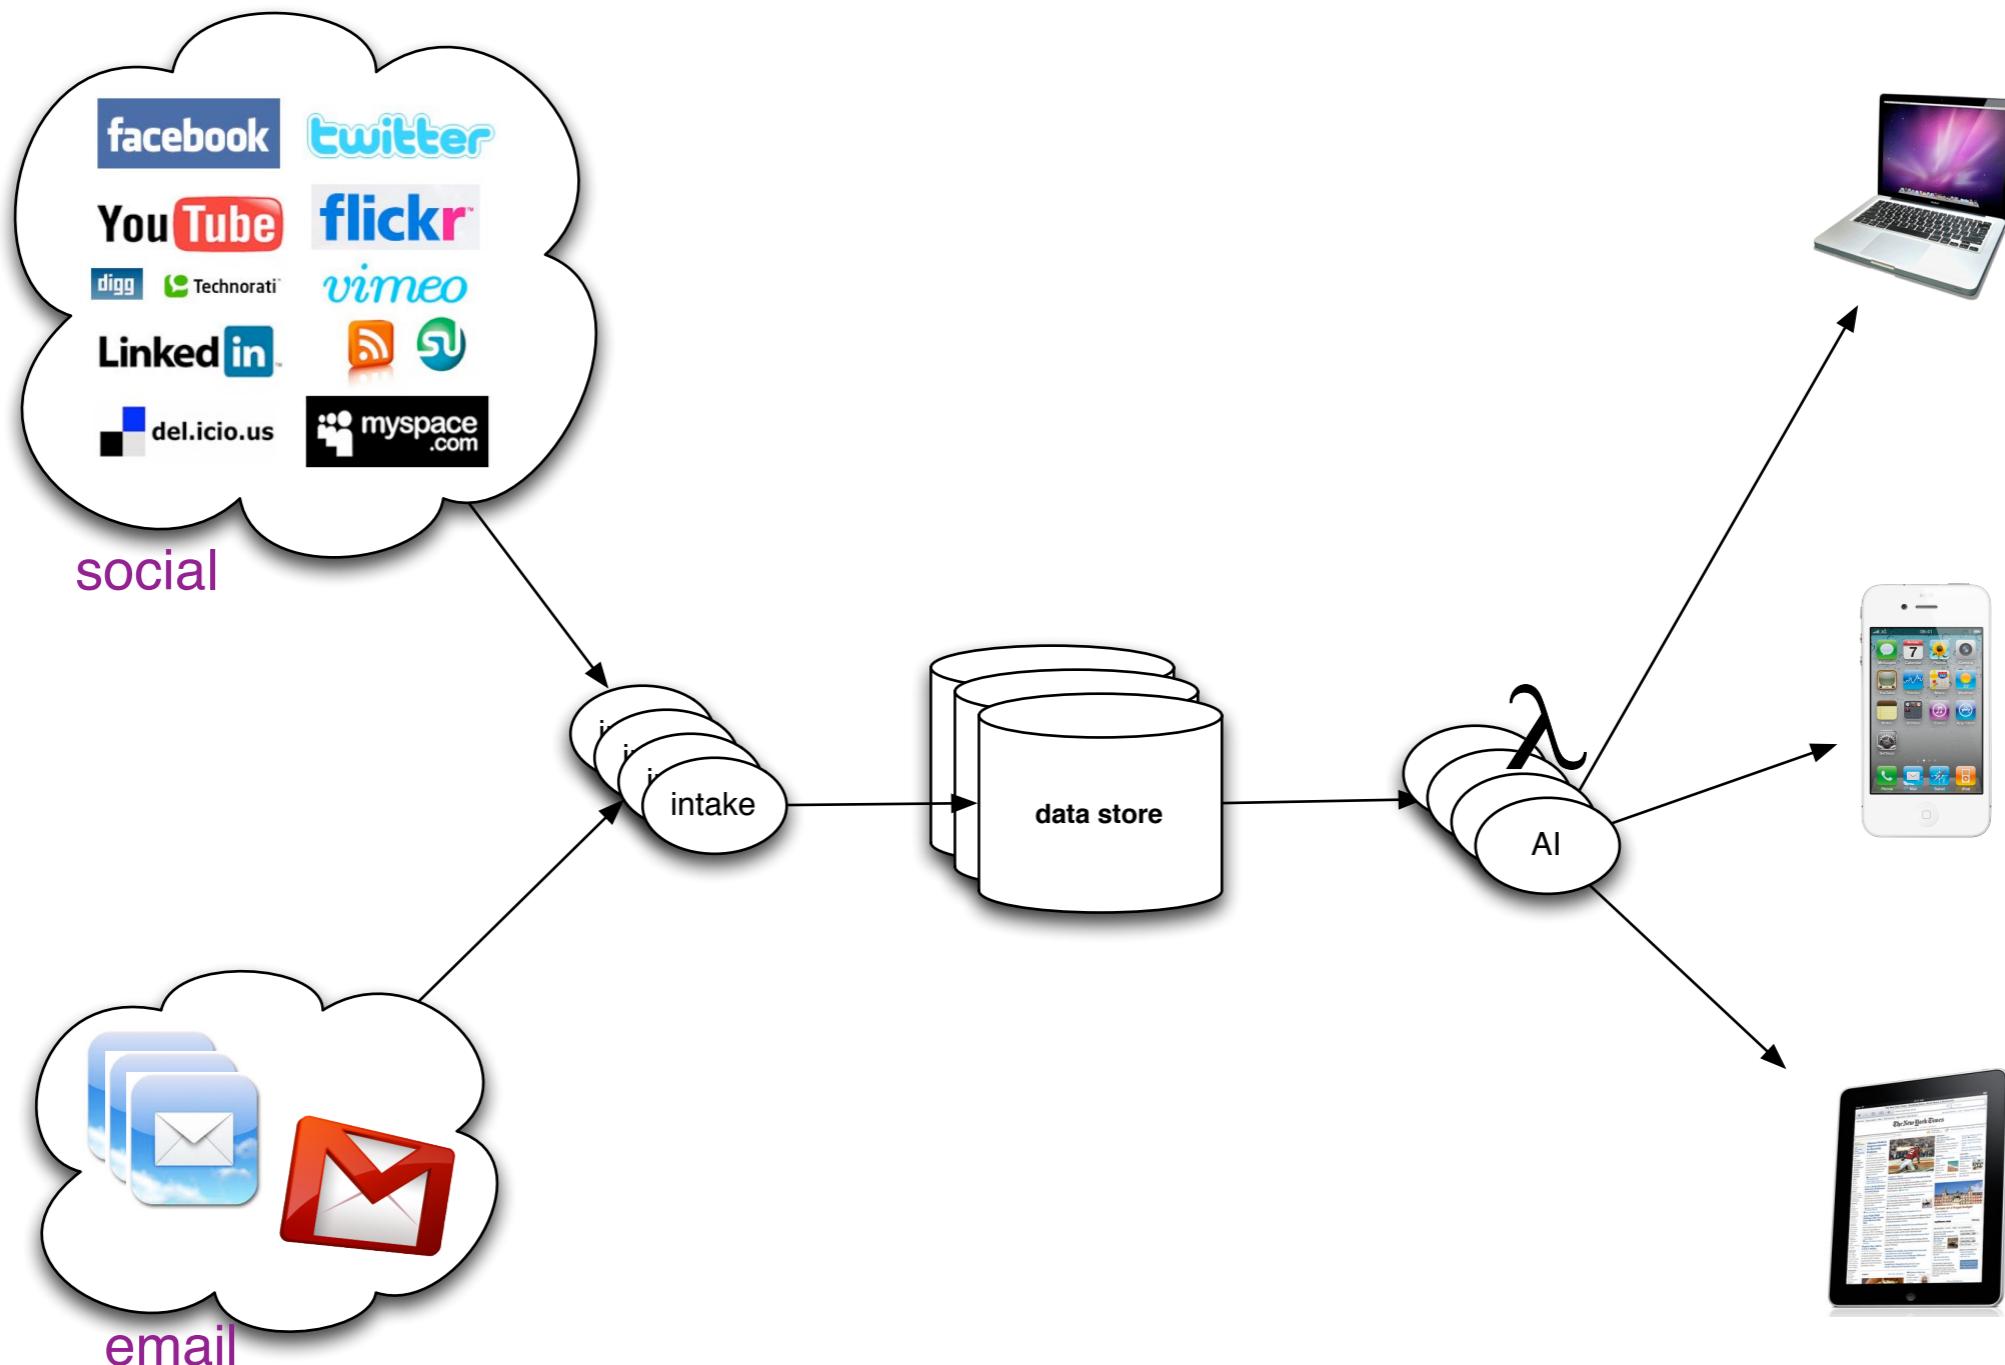


data



zolodeck

zolodeck

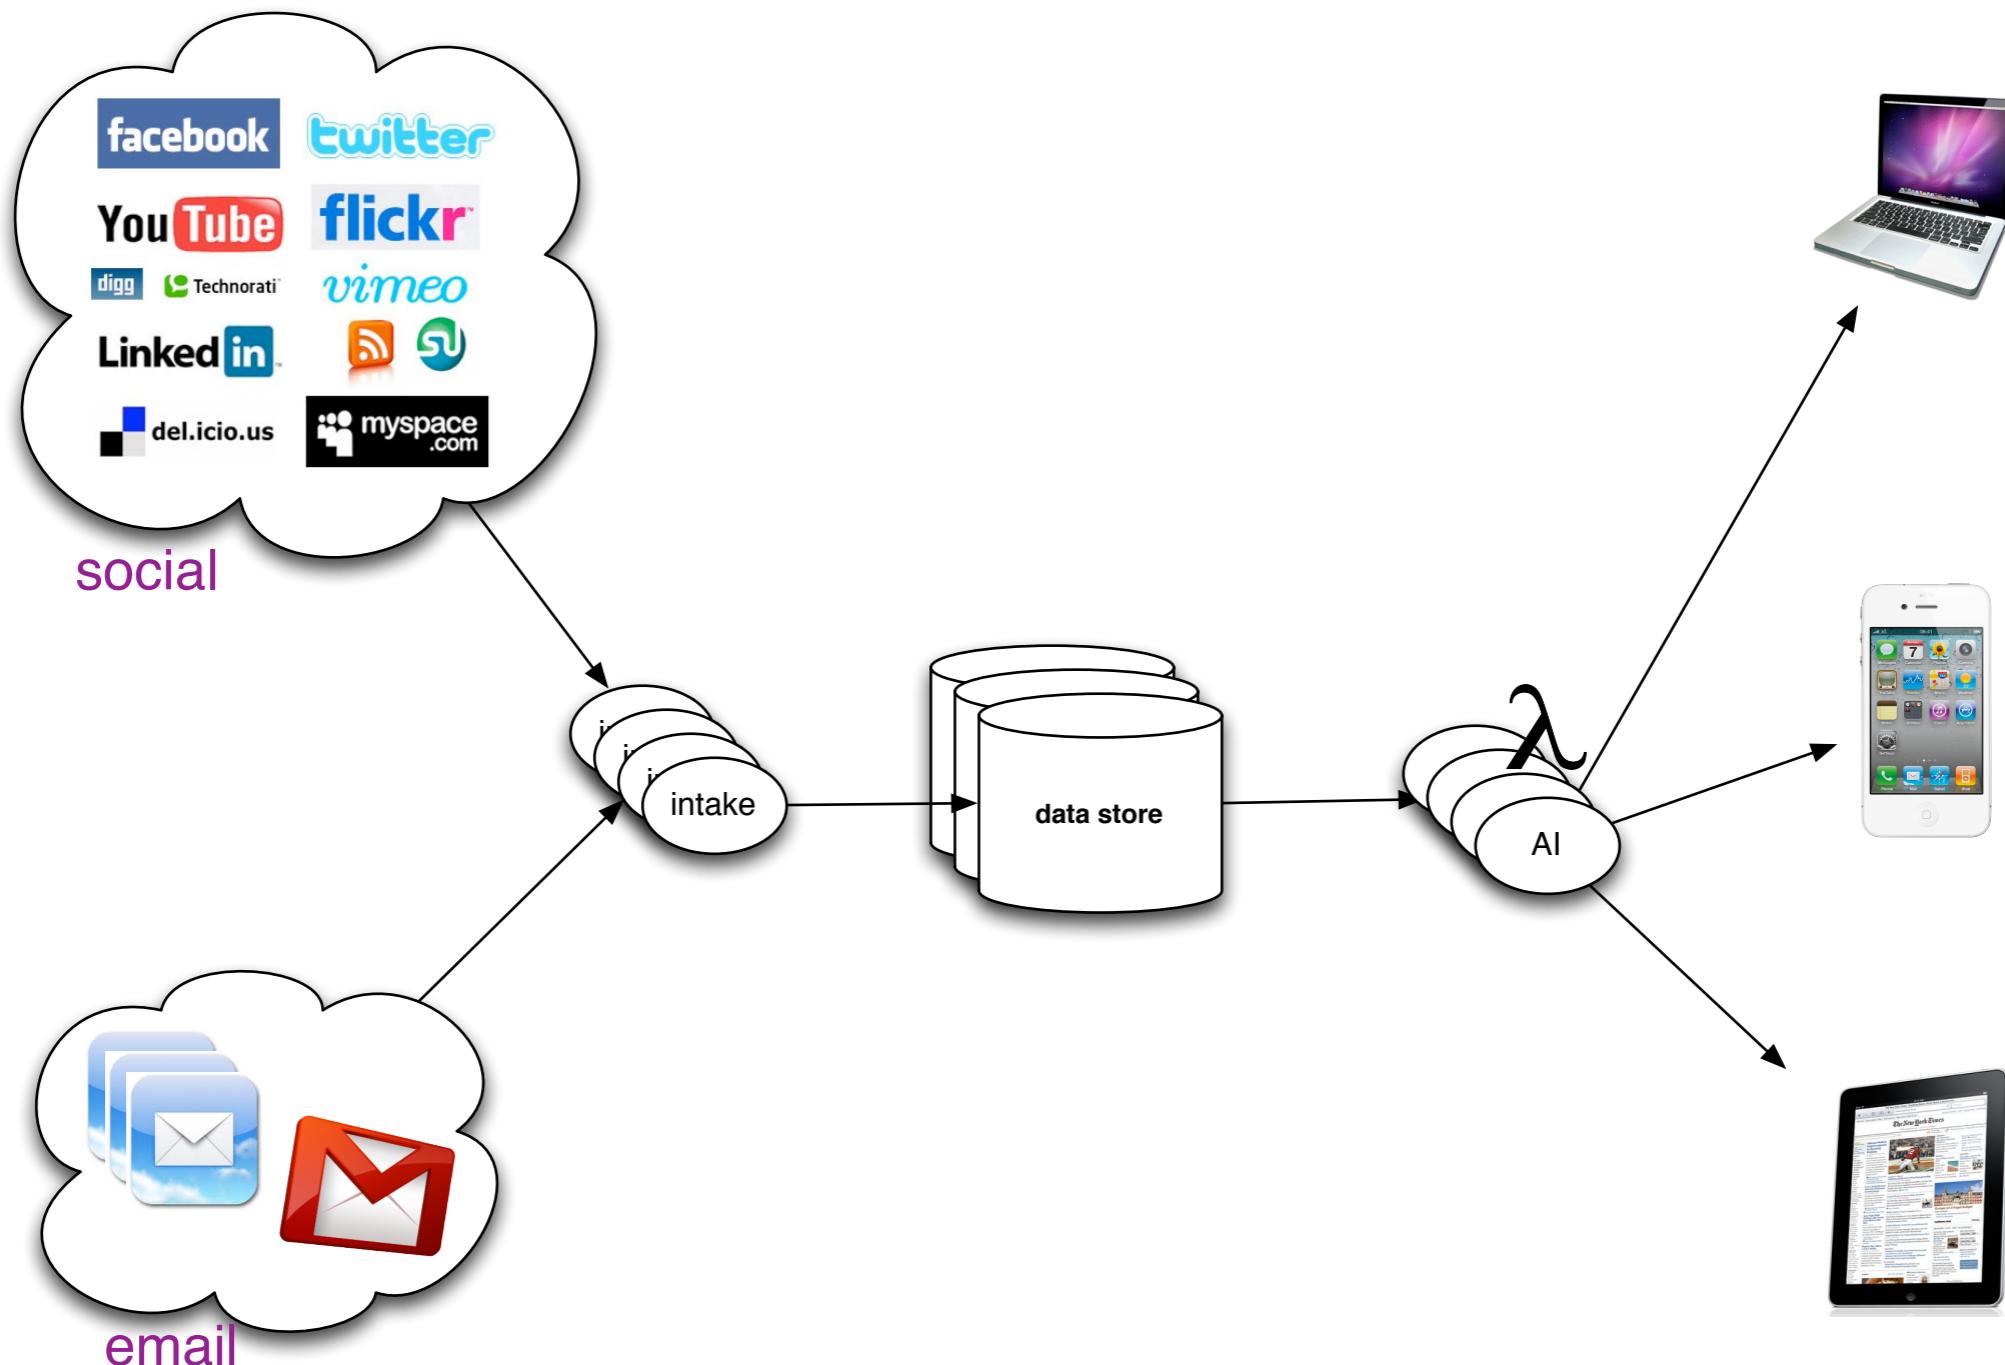


data volume



zolodeck

zolodeck

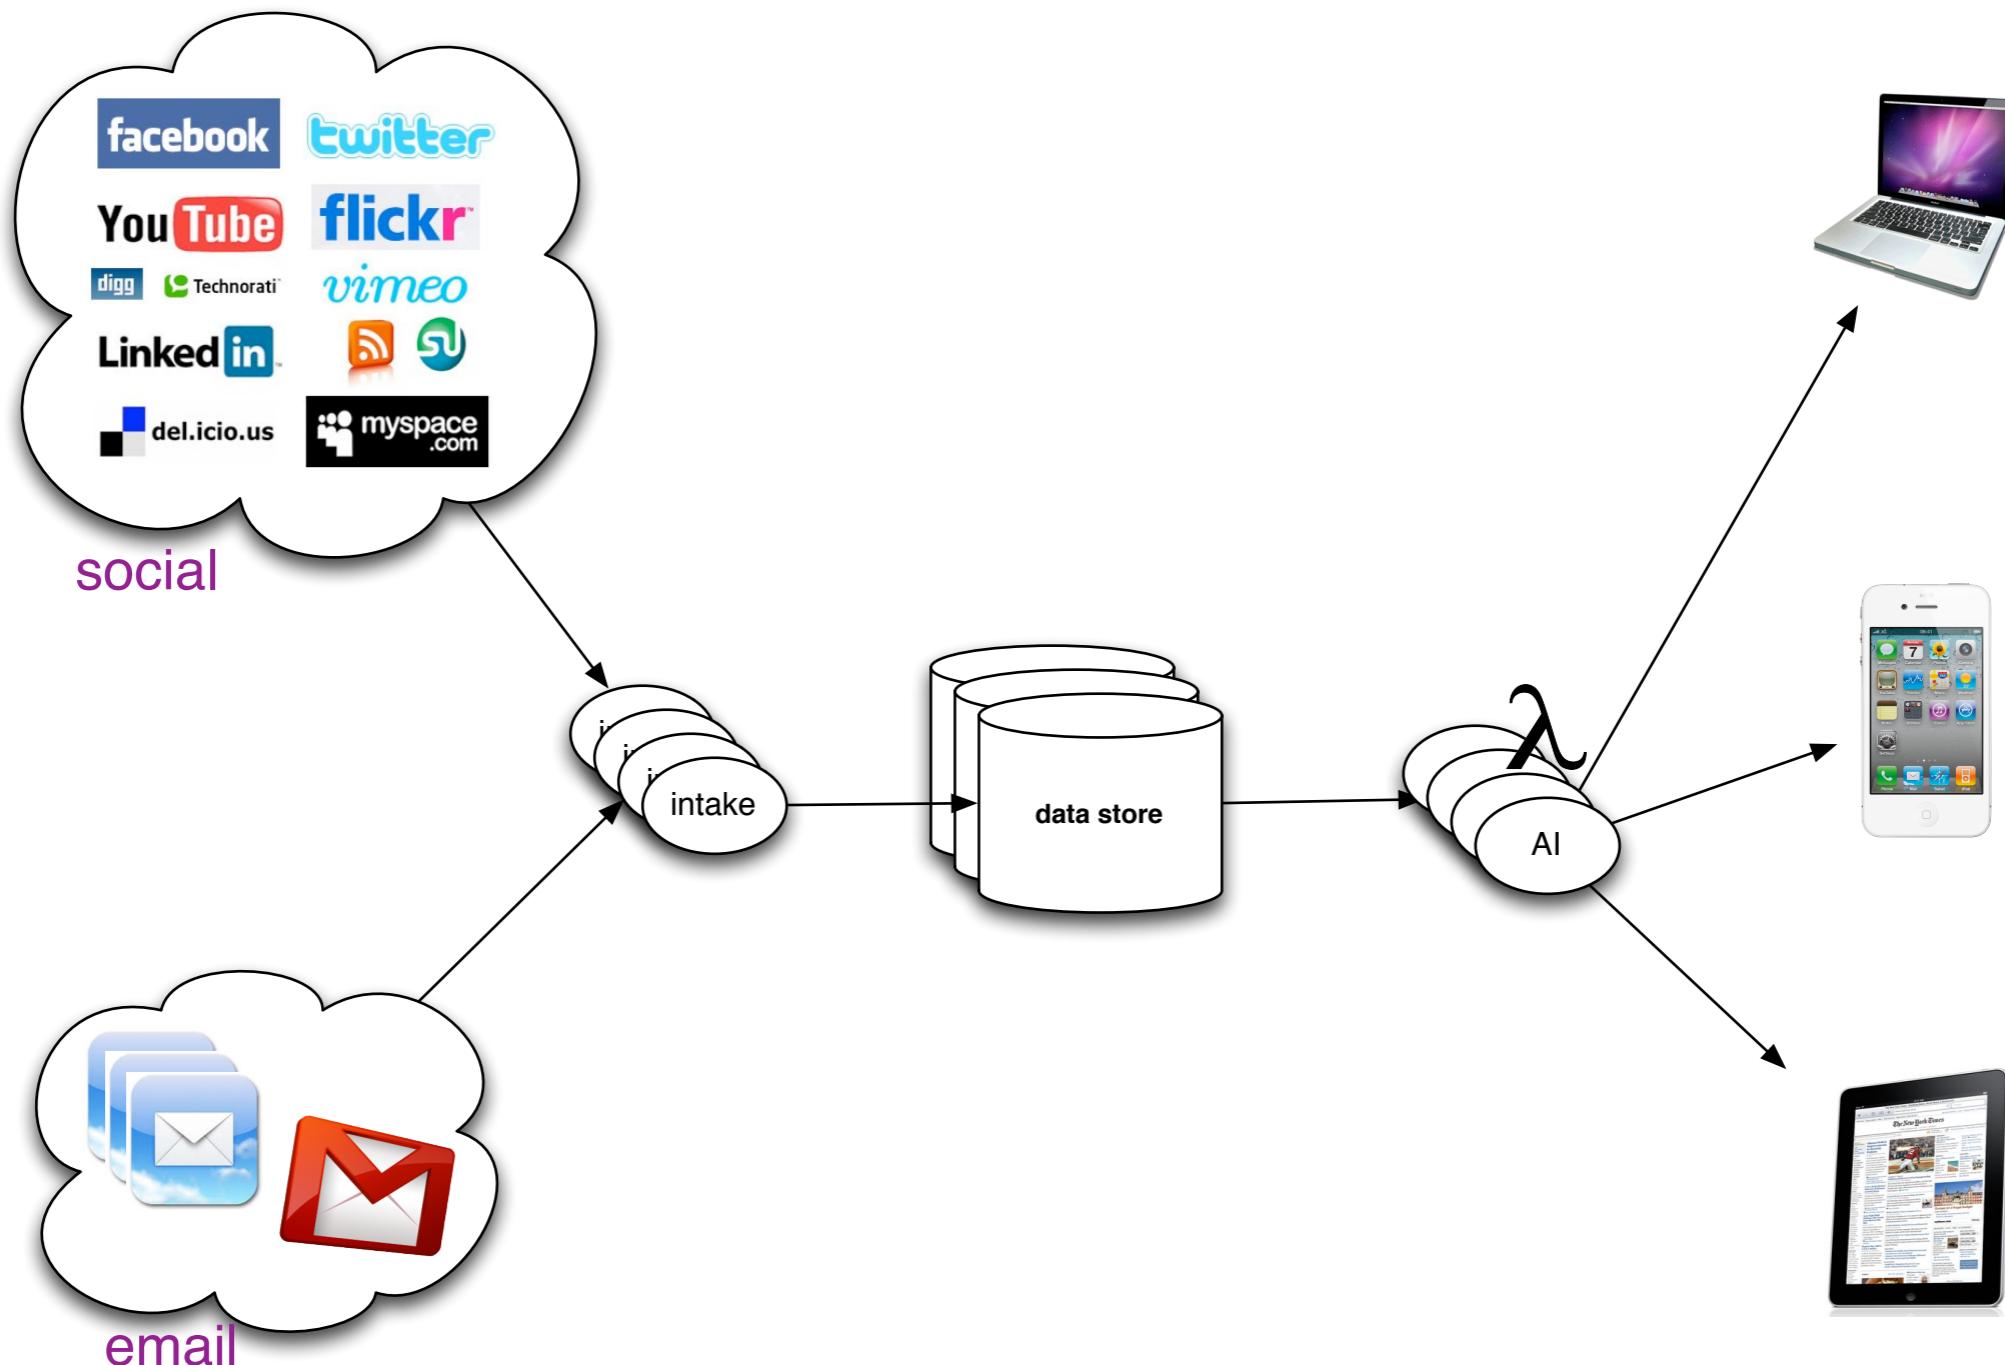


data volume

processing



zolodeck

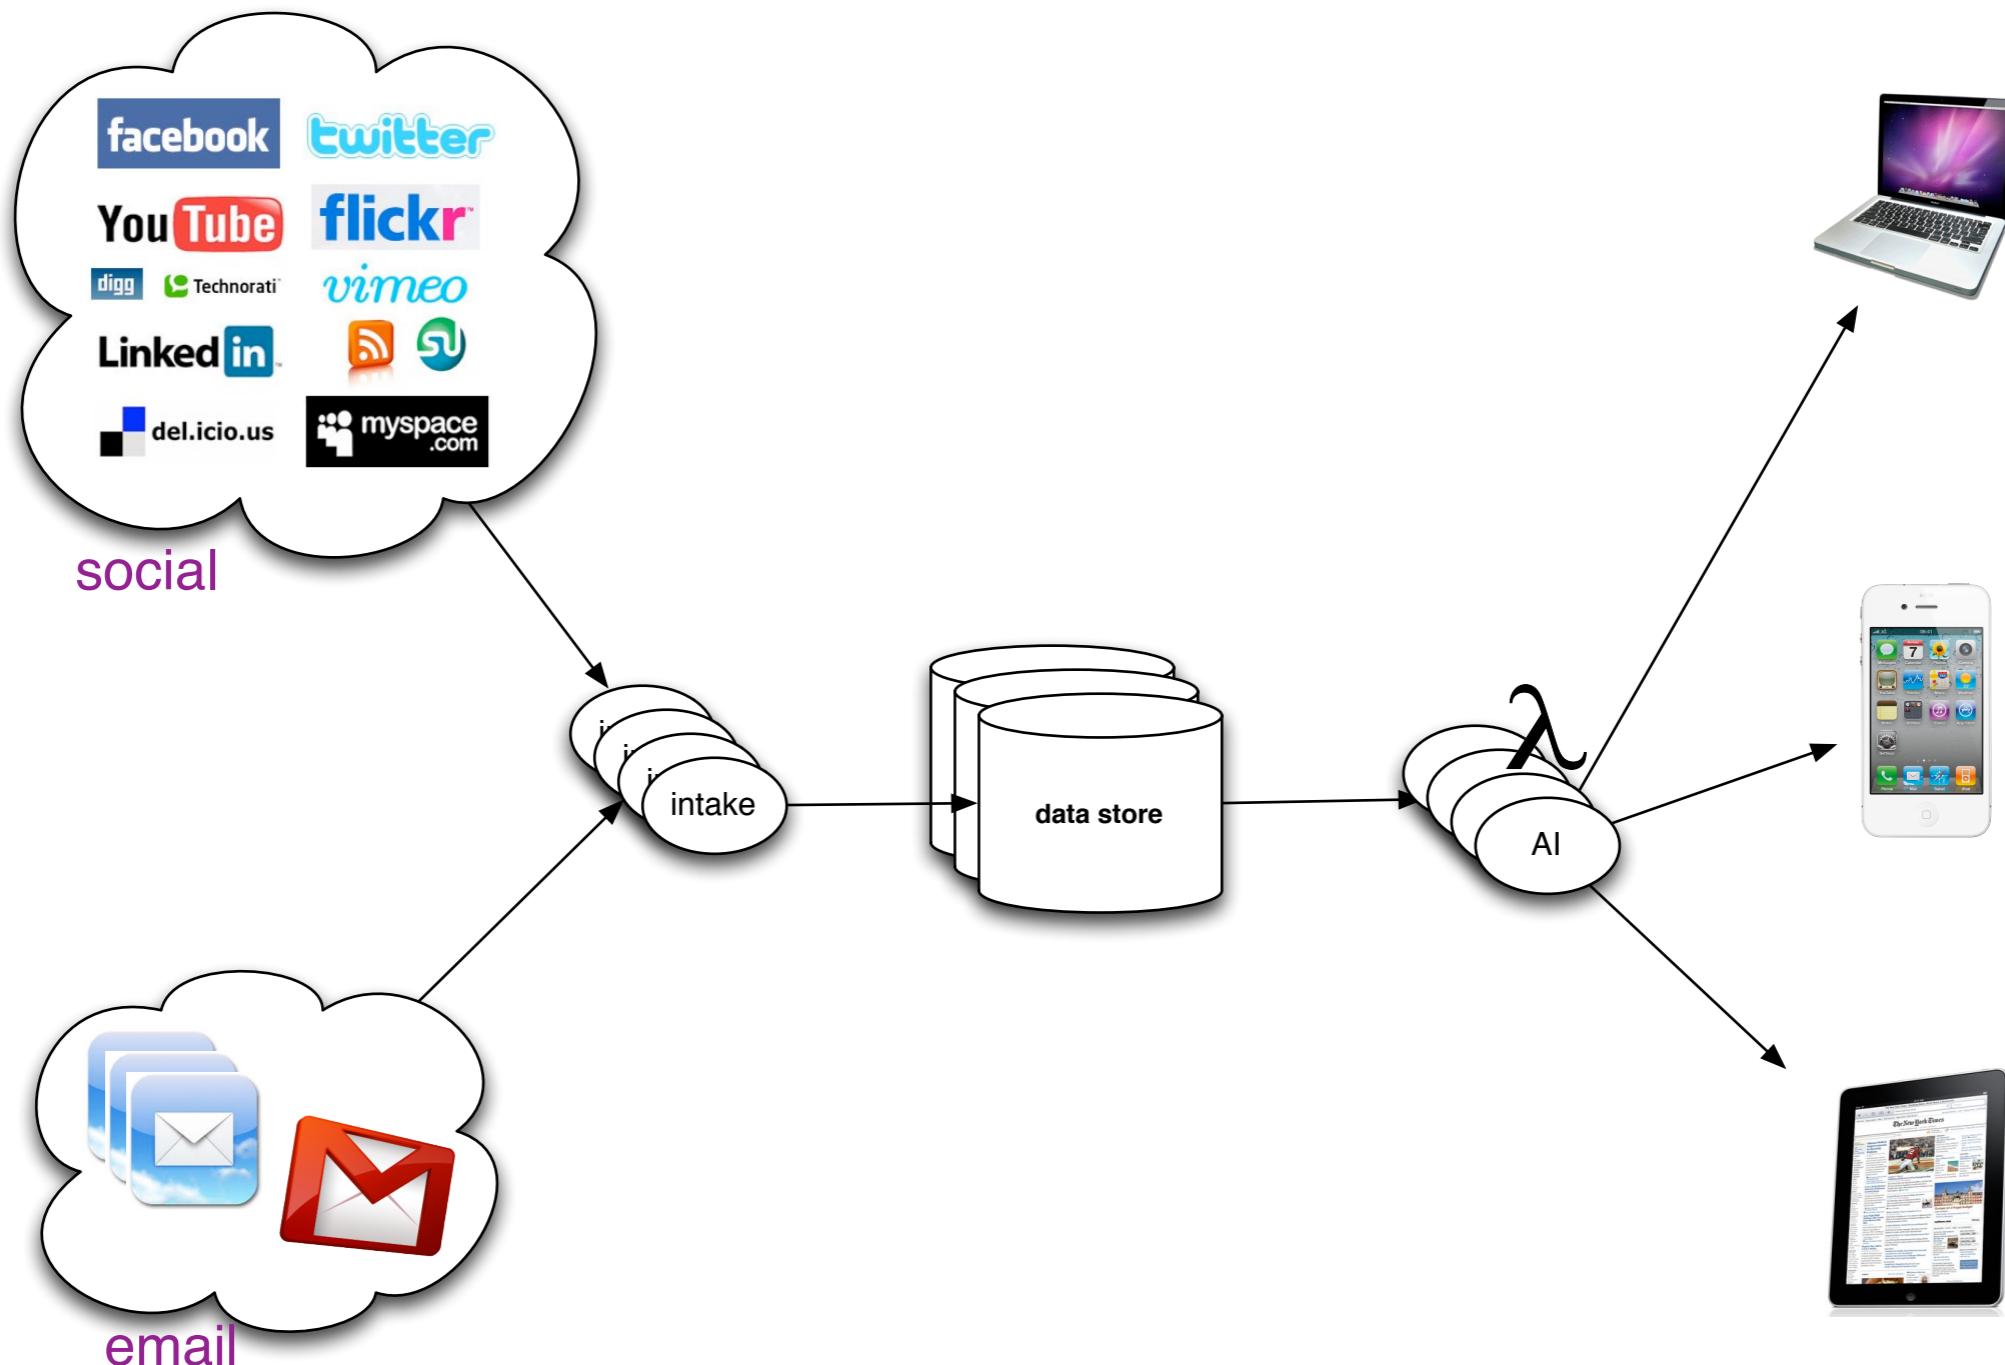


data volume

processing: scale



zolodeck



data volume

processing: scale, real-time



our choices



language: clojure



language: clojure
processing: storm



language: clojure
processing: storm
data: datomic



datomic



the inevitable path to sharding



the inevitable pain of sharding



lose the things a database does



lose the things a database does

consistency
transactions
queries



the new NoSQL stuff



gain vs loss



consistency
transactions
queries

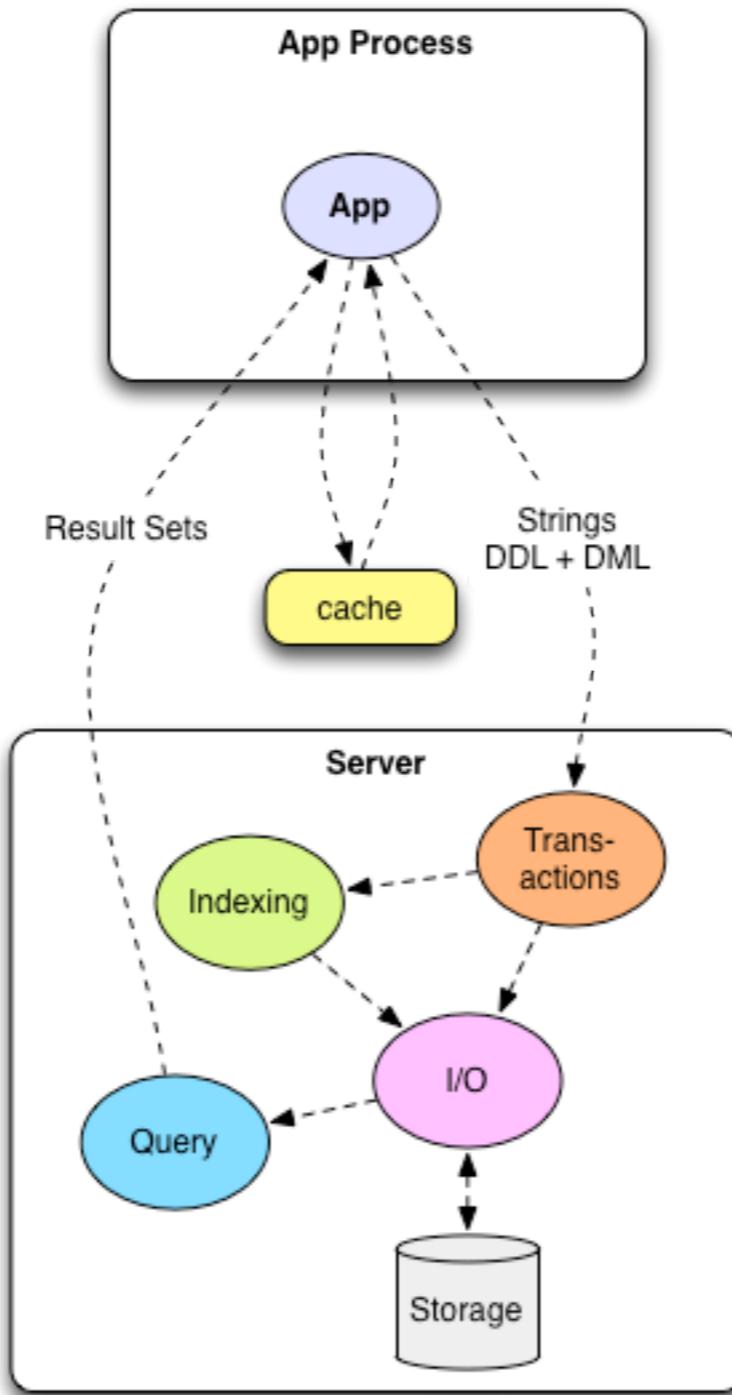


datomic gives you these back



datomic gives you these back
without manual sharding





the good old database



zolodeck

the database deconstructed



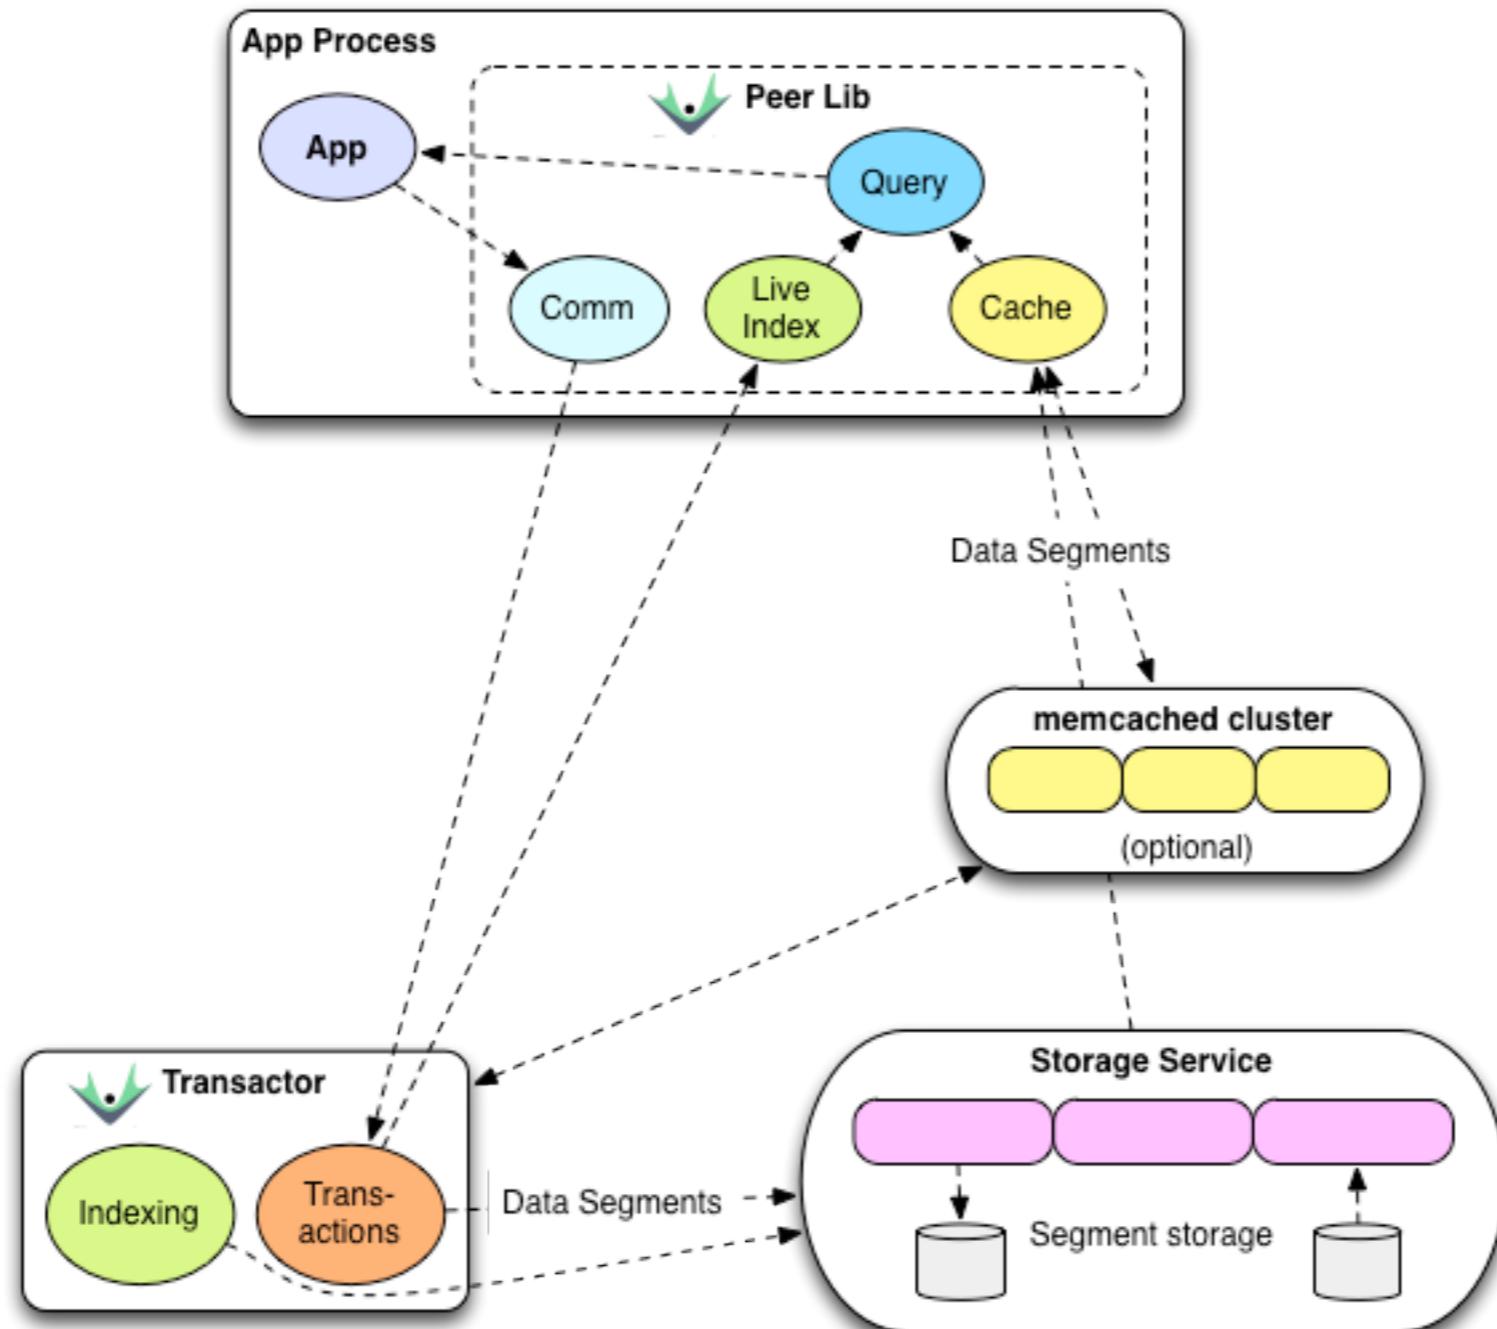
the database deconstructed

relocate subsystems
separate services each do one thing
simplification



peers
transactors
storage services





the datomic way



separating reads and writes



applications only read from storage
transactor – handles transactions, provides consistency, reflects changes to peers

integrated data distribution



integrated data distribution

built-in, in-memory caches



integrated data distribution

built-in, in-memory caches
self-tunes to working set



zolodeck

integrated data distribution

built-in, in-memory caches
self-tunes to working set
automatic



peers have the query engine
datalog



zolodeck

simple rules and data patterns
declarative: implicit joins, meaning is evident [:find ?e ?x :where [?e :age 42] [?e :likes ?x]]
locality: data + app

elastic



cloud-ready

elastic

add more peers



zolodeck

cloud-ready

elastic

add more peers
commodity hardware



cloud-ready

elastic

add more peers
commodity hardware
resilient to failures



cloud-ready

immutable



immutable

remember everything



immutable

remember everything
audit everything



immutable

remember everything
audit everything
automatic



Minimal schema



Minimal schema

fact oriented - entity, attribute, value, transaction



Minimal schema

fact oriented - entity, attribute, value, transaction
anya



Minimal schema

fact oriented - entity, attribute, value, transaction

anya age



Minimal schema

fact oriented - entity, attribute, value, transaction

anya age 3



Minimal schema

fact oriented - entity, attribute, value, transaction

anya	age	3	t100
------	-----	---	------



Minimal schema

fact oriented - entity, attribute, value, transaction

anya age 3 t|00

supports cardinality



Minimal schema

fact oriented - entity, attribute, value, transaction

anya age 3 tl

supports cardinality

structure is up to the app code



transactional



zolodeck

- consistent: without impeding other threads or peers
- consistent: reads/queries without transactions (thanks to immutability)
- peers get a queue of all transactions
- can use transactions in queries
- facilitates event-driven triggers without polling

transactional

ordered



zolodeck

- consistent: without impeding other threads or peers
- consistent: reads/queries without transactions (thanks to immutability)
- peers get a queue of all transactions
- can use transactions in queries
- facilitates event-driven triggers without polling

transactional ordered first-class transactions



zolodeck

- consistent: without impeding other threads or peers
- consistent: reads/queries without transactions (thanks to immutability)
- peers get a queue of all transactions
- can use transactions in queries
- facilitates event-driven triggers without polling

transactional
ordered
first-class transactions
annotated transactions



zolodeck

- consistent: without impeding other threads or peers
- consistent: reads/queries without transactions (thanks to immutability)
- peers get a queue of all transactions
- can use transactions in queries
- facilitates event-driven triggers without polling

transactional

ordered
first-class transactions
annotated transactions
as-of, since, or windowed queries



- consistent: without impeding other threads or peers
- consistent: reads/queries without transactions (thanks to immutability)
- peers get a queue of all transactions
- can use transactions in queries
- facilitates event-driven triggers without polling

full-text search



full-text search

built-in lucene



storage services



in-memory, in-process – great for testing

storage services

Amazon Dynamo
RDBMS
memory



in-memory, in-process – great for testing

storage services

Amazon Dynamo
RDBMS
memory

unit testing



in-memory, in-process – great for testing

so that's the data story



now, the processing story



how to process lots of data?



when machine is not enough?

start with a single box



typical progression:

add machines
assign work (mod hash?)
add more machines



next:

message queue
workers
dispatcher



managing work



managing work
assigning machines



managing work

assigning machines
monitoring workers
restart workers



managing work

assigning machines
monitoring workers
restart workers
re-assign work



managing work

assigning machines
monitoring workers
restart workers
re-assign work
idempotent?



pipelined work



- multiple steps, workers send messages, multiple queues...
- for scale, or failover, or contracting elastic cluster

pipelined work

Aaaaaa!!



- multiple steps, workers send messages, multiple queues...
- for scale, or failover, or contracting elastic cluster

painful



painful
not fault-tolerant



painful
not fault-tolerant
tedious



abstract away:



higher-level abstraction than message passing

abstract away:

guaranteed processing
horizontal scalability
fault-tolerance



higher-level abstraction than message passing

abstract away:

guaranteed processing
horizontal scalability
fault-tolerance

focus on business logic



higher-level abstraction than message passing

abstract away:

guaranteed processing
horizontal scalability
fault-tolerance

focus on business logic
not like Hadoop



higher-level abstraction than message passing

abstract away:

guaranteed processing
horizontal scalability
fault-tolerance

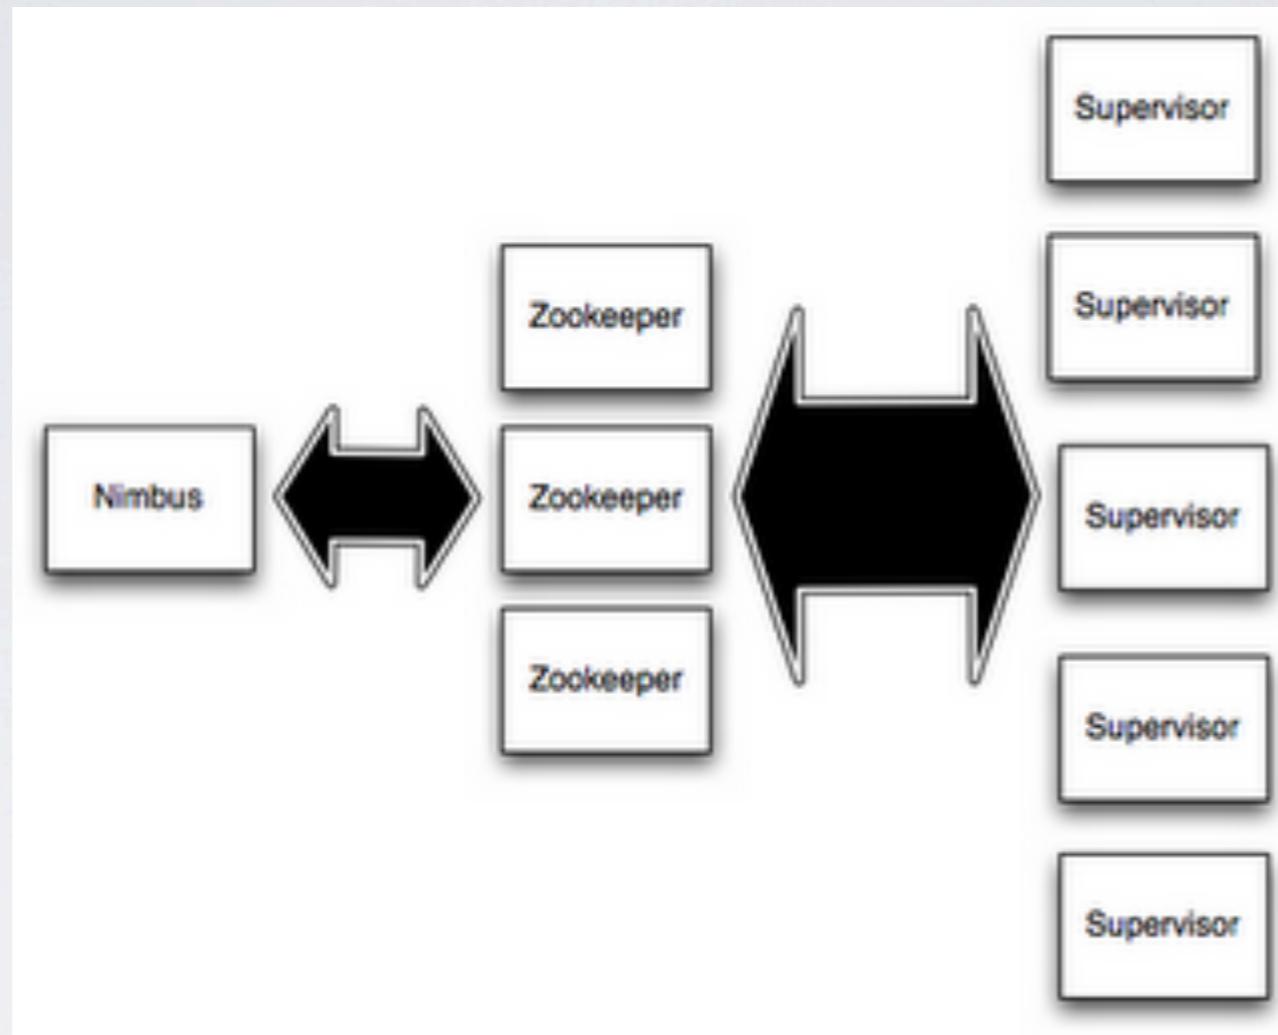
focus on business logic
not like Hadoop

storm



higher-level abstraction than message passing

components:



nimbus
zookeeper
supervisor



topology

```
storm jar something.jar classname arg1 arg2
```



zolodeck

- stream: unbounded sequence of tuples (named list of values)
- tuple: any type of value, with custom serializer
- spout: sources of streams (eg. reads from twitter API)
- bolt: process input streams, and produce output streams
- bolt: does all the work
- topology definitions are thrift structs (so any programming language can be used)

topology

```
storm jar something.jar classname arg1 arg2
```

stream



zolodeck

- stream: unbounded sequence of tuples (named list of values)
- tuple: any type of value, with custom serializer
- spout: sources of streams (eg. reads from twitter API)
- bolt: process input streams, and produce output streams
- bolt: does all the work
- topology definitions are thrift structs (so any programming language can be used)

topology

storm jar something.jar classname arg1 arg2

stream
spout



zolodeck

- stream: unbounded sequence of tuples (named list of values)
- tuple: any type of value, with custom serializer
- spout: sources of streams (eg. reads from twitter API)
- bolt: process input streams, and produce output streams
- bolt: does all the work
- topology definitions are thrift structs (so any programming language can be used)

topology

storm jar something.jar classname arg1 arg2

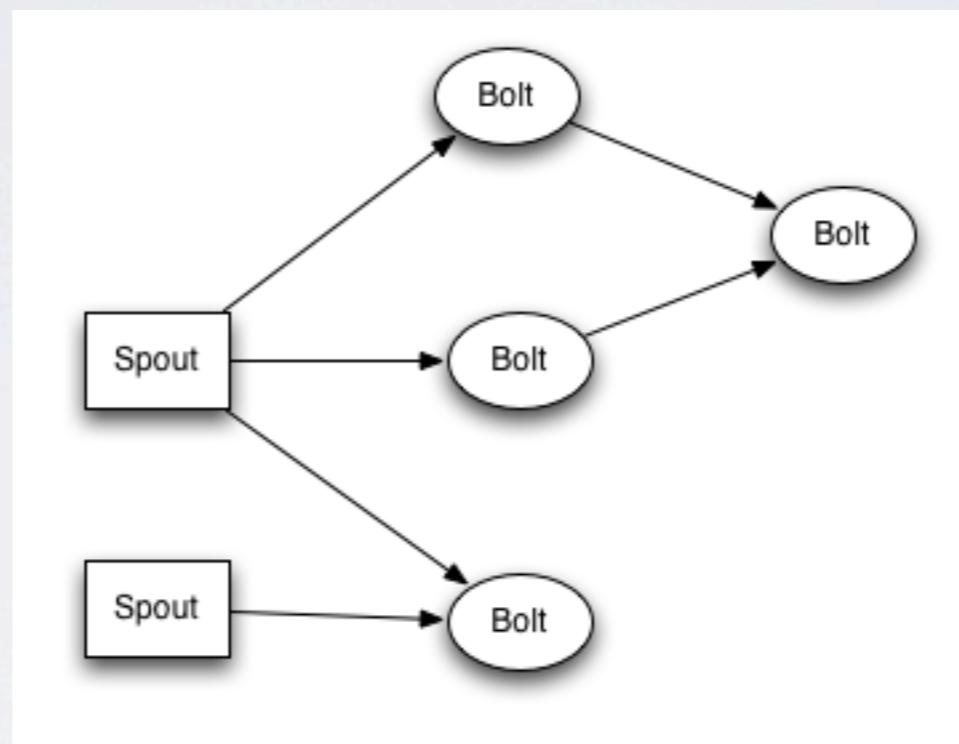
stream
spout
bolt



zolodeck

- stream: unbounded sequence of tuples (named list of values)
- tuple: any type of value, with custom serializer
- spout: sources of streams (eg. reads from twitter API)
- bolt: process input streams, and produce output streams
- bolt: does all the work
- topology definitions are thrift structs (so any programming language can be used)

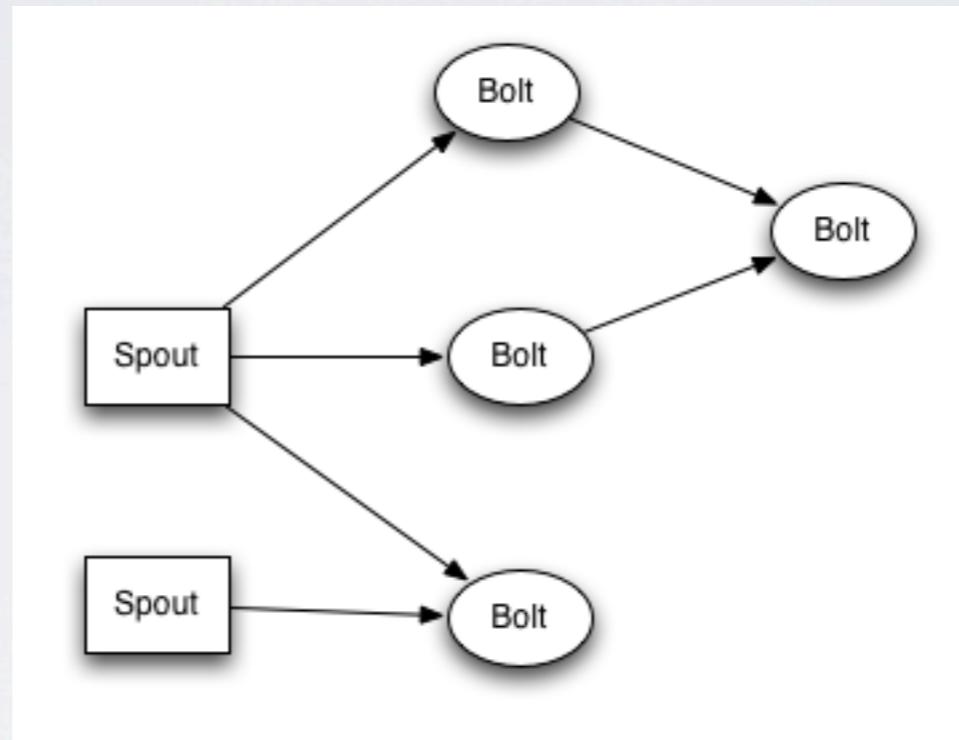
a topology



zolodeck

- multiple nodes
- multiple workers on each node
- multiple tasks will run on each worker, in parallel
- exchanging data as needed
- stream grouping: how to partition stream of messages (e.g. shuffle, fields (mod hash), all, global)

a topology



grouping



zolodeck

- multiple nodes
- multiple workers on each node
- multiple tasks will run on each worker, in parallel
- exchanging data as needed
- stream grouping: how to partition stream of messages (e.g. shuffle, fields (mod hash), all, global)

topology builder



zolodeck

topology builder

```
TopologyBuilder builder = new TopologyBuilder();
builder.setSpout("words", new TestWordSpout(), 10);
builder.setBolt("exclaim1", new ExclamationBolt(), 3)
    .shuffleGrouping("words");
builder.setBolt("exclaim2", new ExclamationBolt(), 2)
    .shuffleGrouping("exclaim1");
```



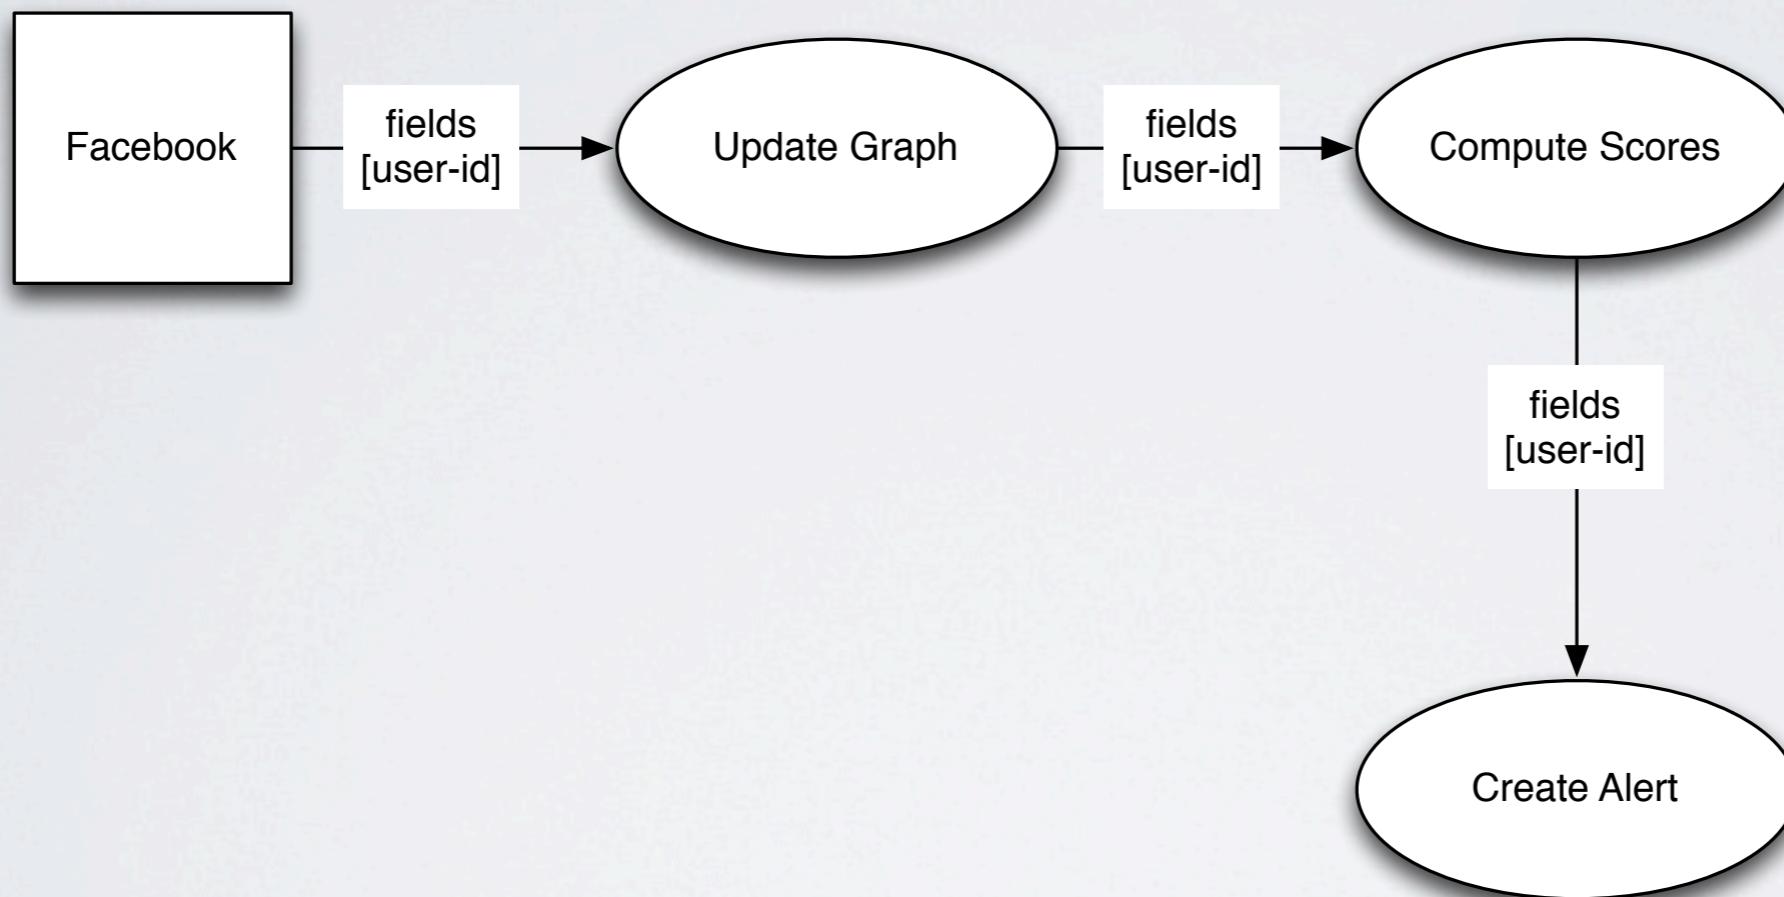
zolodeck

- “words” is the name, 10 is parallelism, # of tasks across the cluster
- “exclaim1” is name of bolt, 3 is parallelism, # of tasks, reads from “words” via shuffle grouping
- consumer decides what to accept

zolodeck snippet



zolodeck snippet



zolodeck

submit the topology



config: specify the number of JVM process on the cluster

submit the topology

nimbus manages the cluster and the job



zolodeck

config: specify the number of JVM process on the cluster

local mode

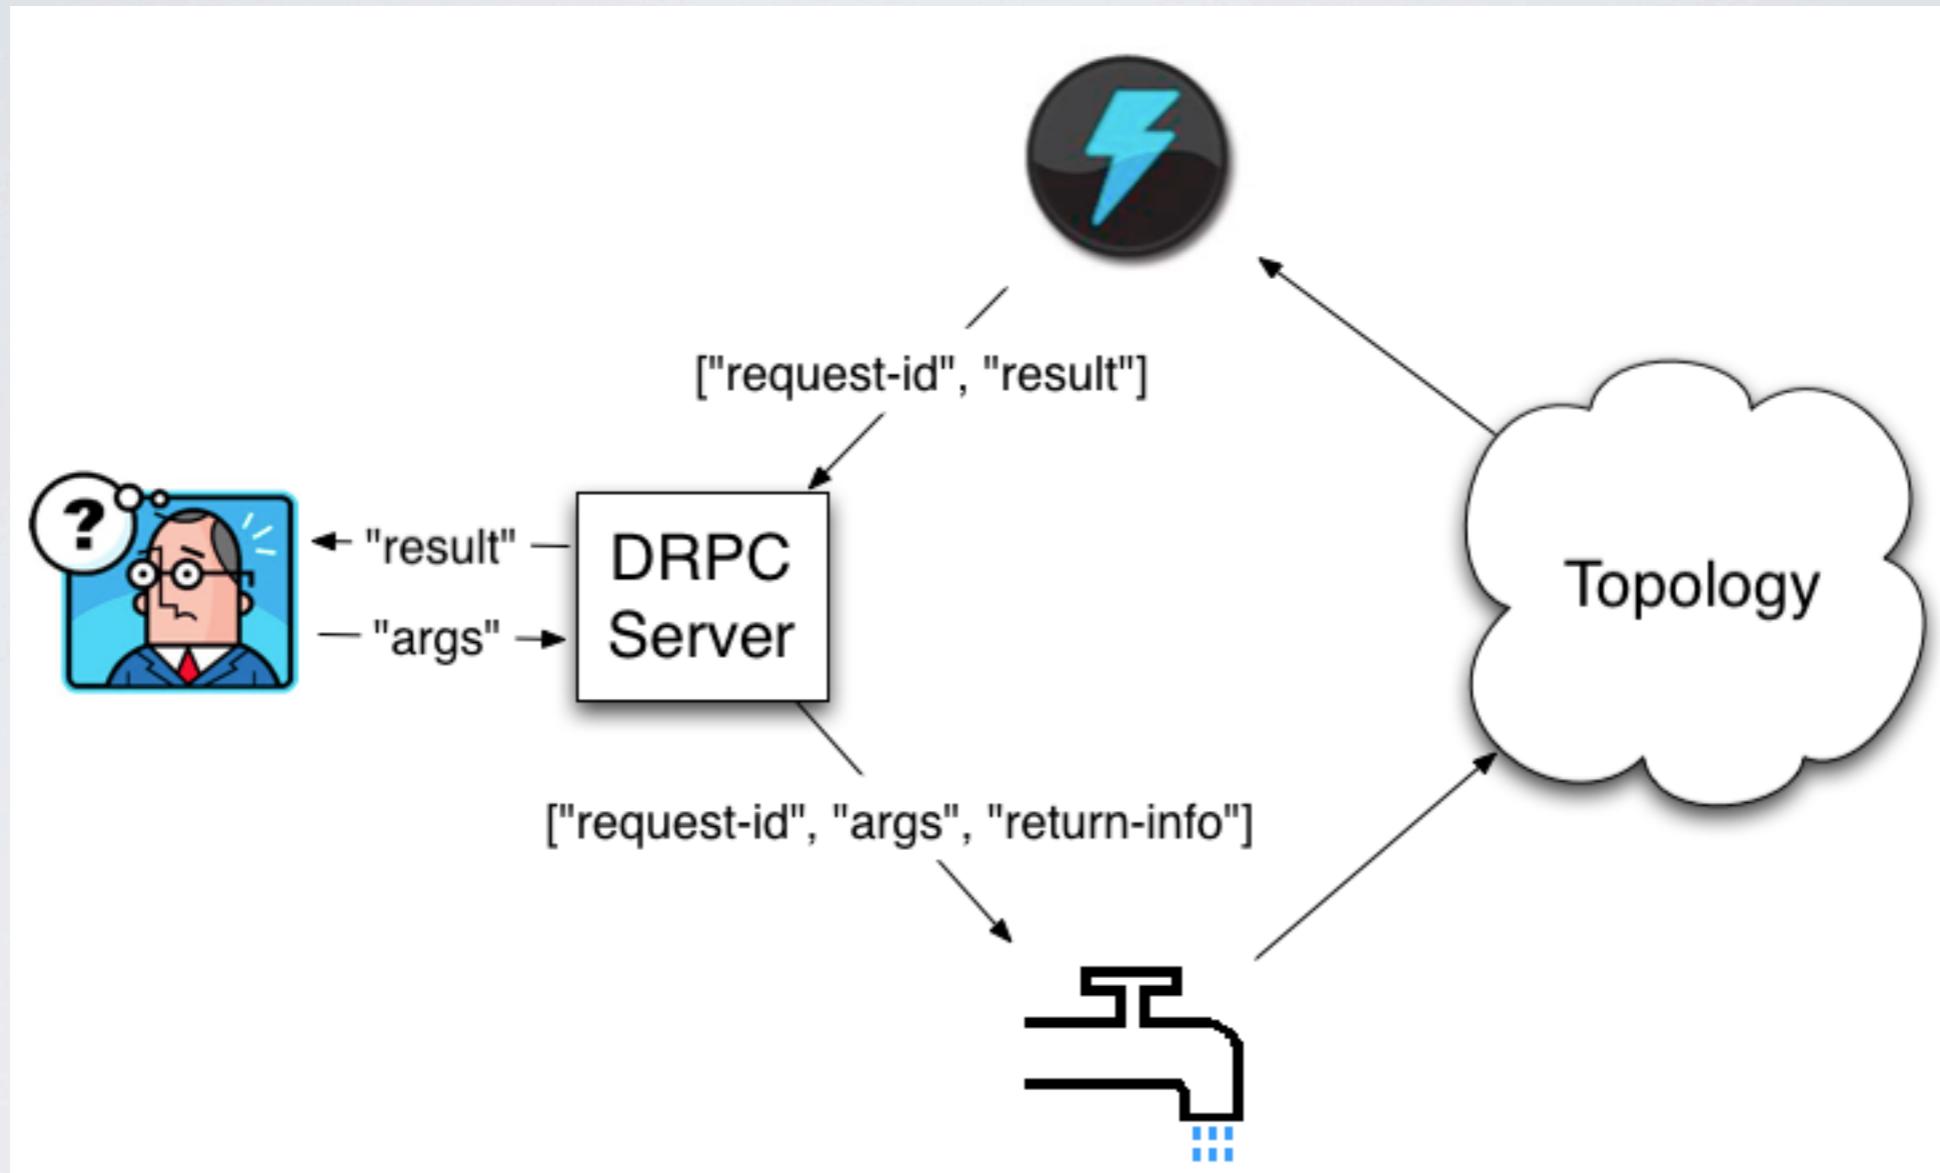


local mode

unit testing



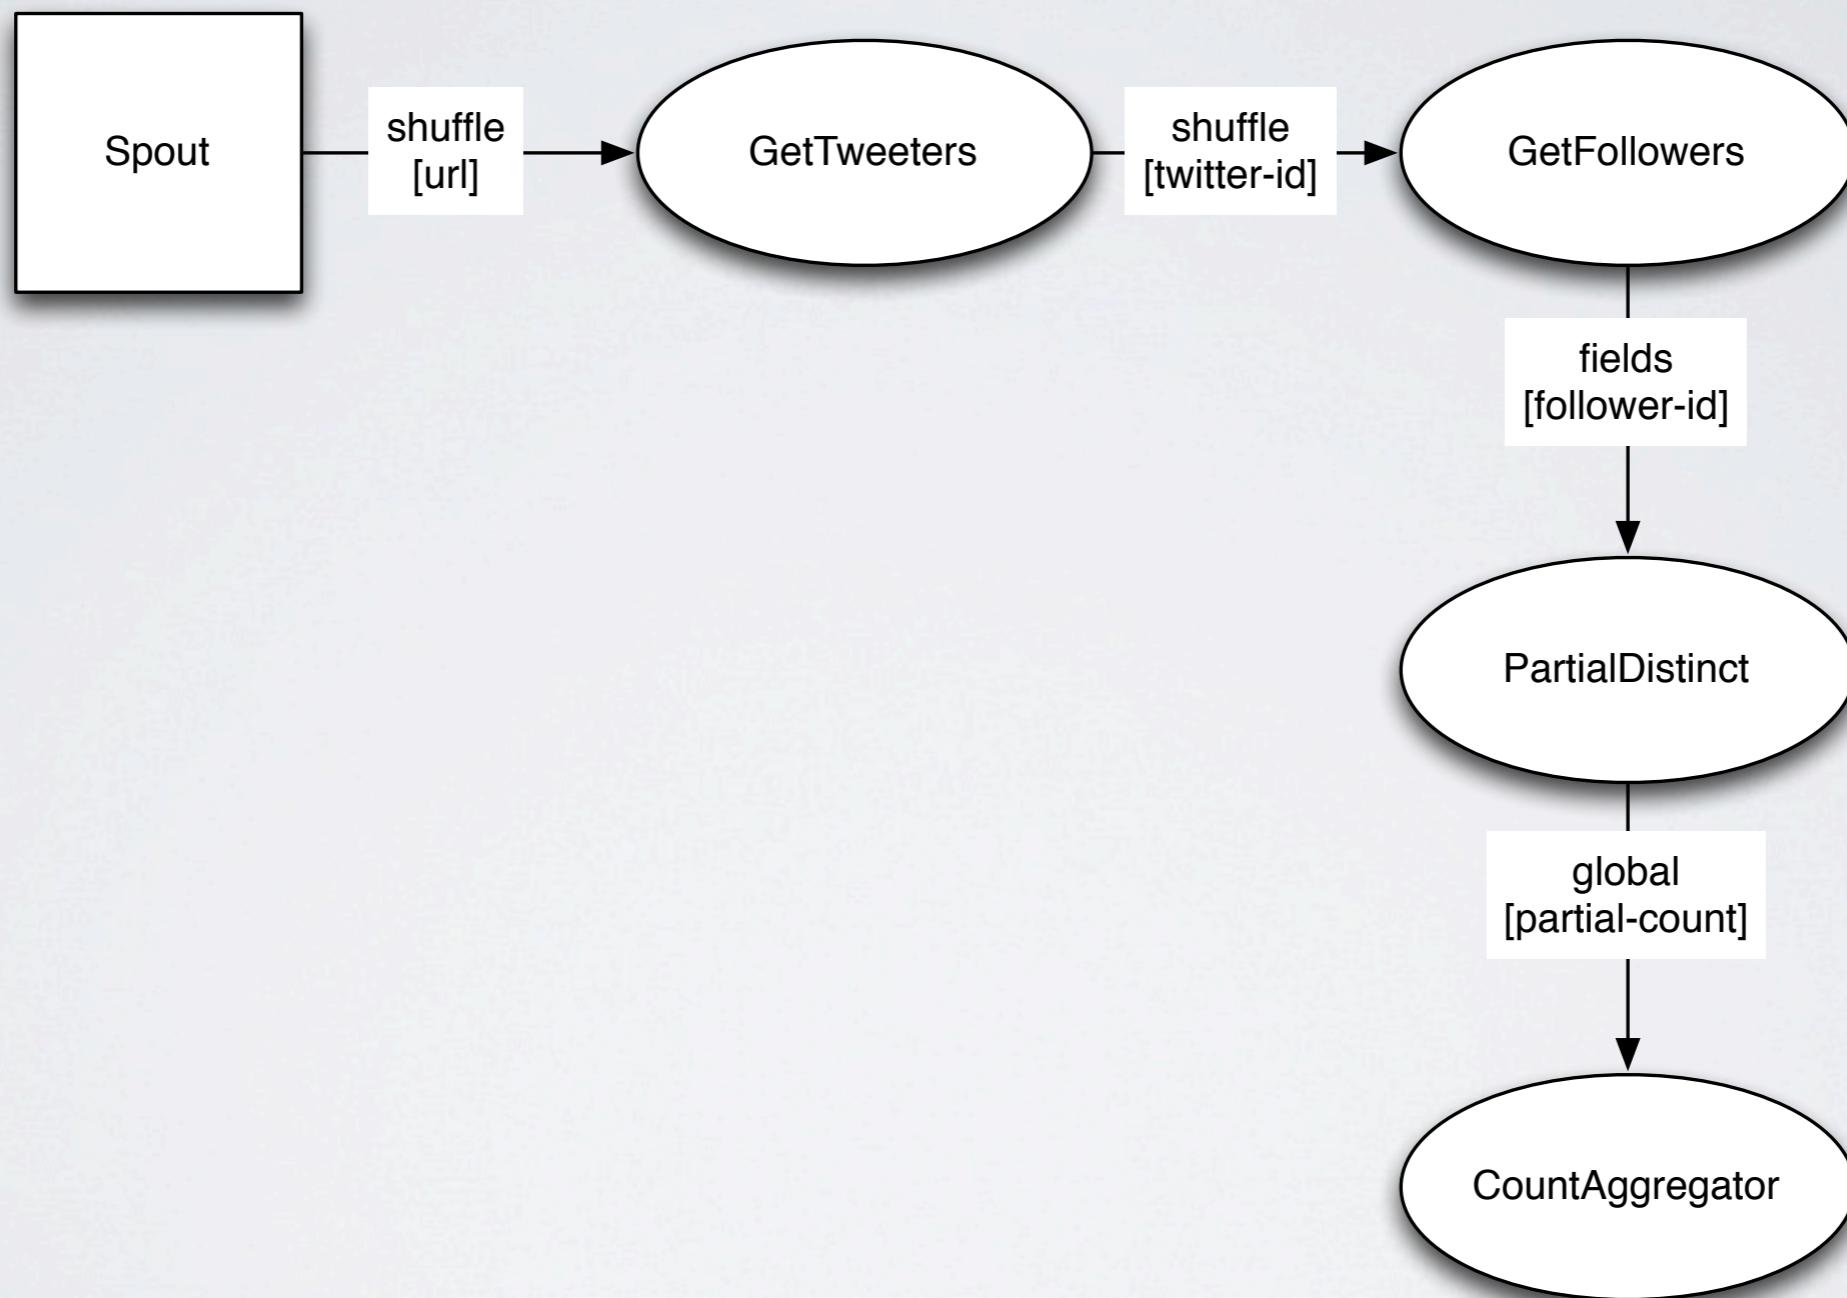
distributed RPC



zolodeck

- server acts as spout
- creates a stream of function invocations (tuple - id of req, args, return info (host, port))
-

get-reach(url)



get-reach(url)

```
LinearDRPCTopologyBuilder builder = new  
LinearDRPCTopologyBuilder("reach");  
builder.addBolt(new GetTweeters(), 3);  
builder.addBolt(new GetFollowers(), 12)  
    .shuffleGrouping();  
builder.addBolt(new PartialUniquer(), 6)  
    .fieldsGrouping(new Fields("id", "follower"));  
builder.addBolt(new CountAggregator(), 2)  
    .fieldsGrouping(new Fields("id"));
```



zolodeck

- request ID is the parameter based on which an initial request is broken up, and sub tasks are run
- linearDRPCTopologyBuilder knows when to call “finish” on each batch

storm-deploy



single script deployment of storm cluster on AWS

storm



zolodeck

at-least once guarantee
tuple tree
tell storm – adding a new edge to tuple tree
tell storm – acknowledging a tuple is processed
tuple tree – 20 bytes for each, regardless of size of tuple tree

storm

guaranteed processing
easy horizontal scalability
fault-tolerance



zolodeck

at-least once guarantee
tuple tree
tell storm – adding a new edge to tuple tree
tell storm – acknowledging a tuple is processed
tuple tree – 20 bytes for each, regardless of size of tuple tree

storm

guaranteed processing
easy horizontal scalability
fault-tolerance

focus on business logic



zolodeck

at-least once guarantee
tuple tree
tell storm – adding a new edge to tuple tree
tell storm – acknowledging a tuple is processed
tuple tree – 20 bytes for each, regardless of size of tuple tree

storm

guaranteed processing
easy horizontal scalability
fault-tolerance

focus on business logic

stream processing
distributed RPC
continuous computation



at-least once guarantee
tuple tree
tell storm – adding a new edge to tuple tree
tell storm – acknowledging a tuple is processed
tuple tree – 20 bytes for each, regardless of size of tuple tree

so that's the processing story



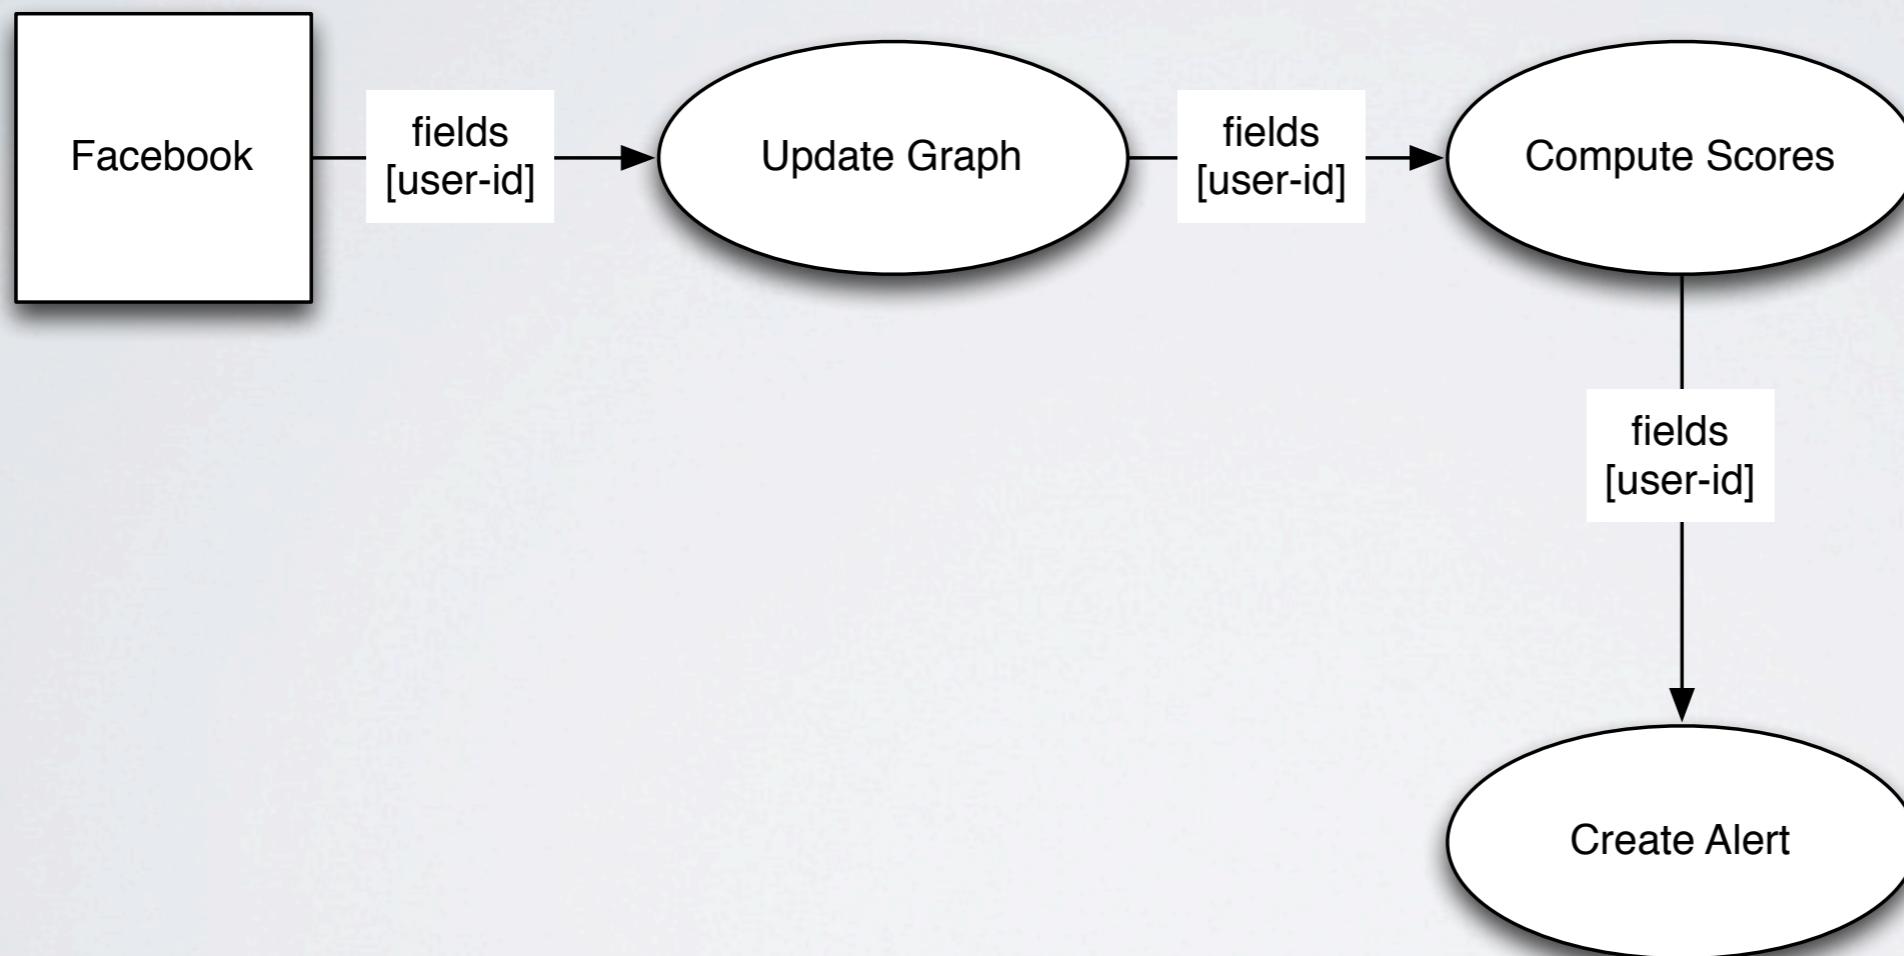
used by Twitter, and a whole bunch of other companies
1M messages/sec

zolodeck example

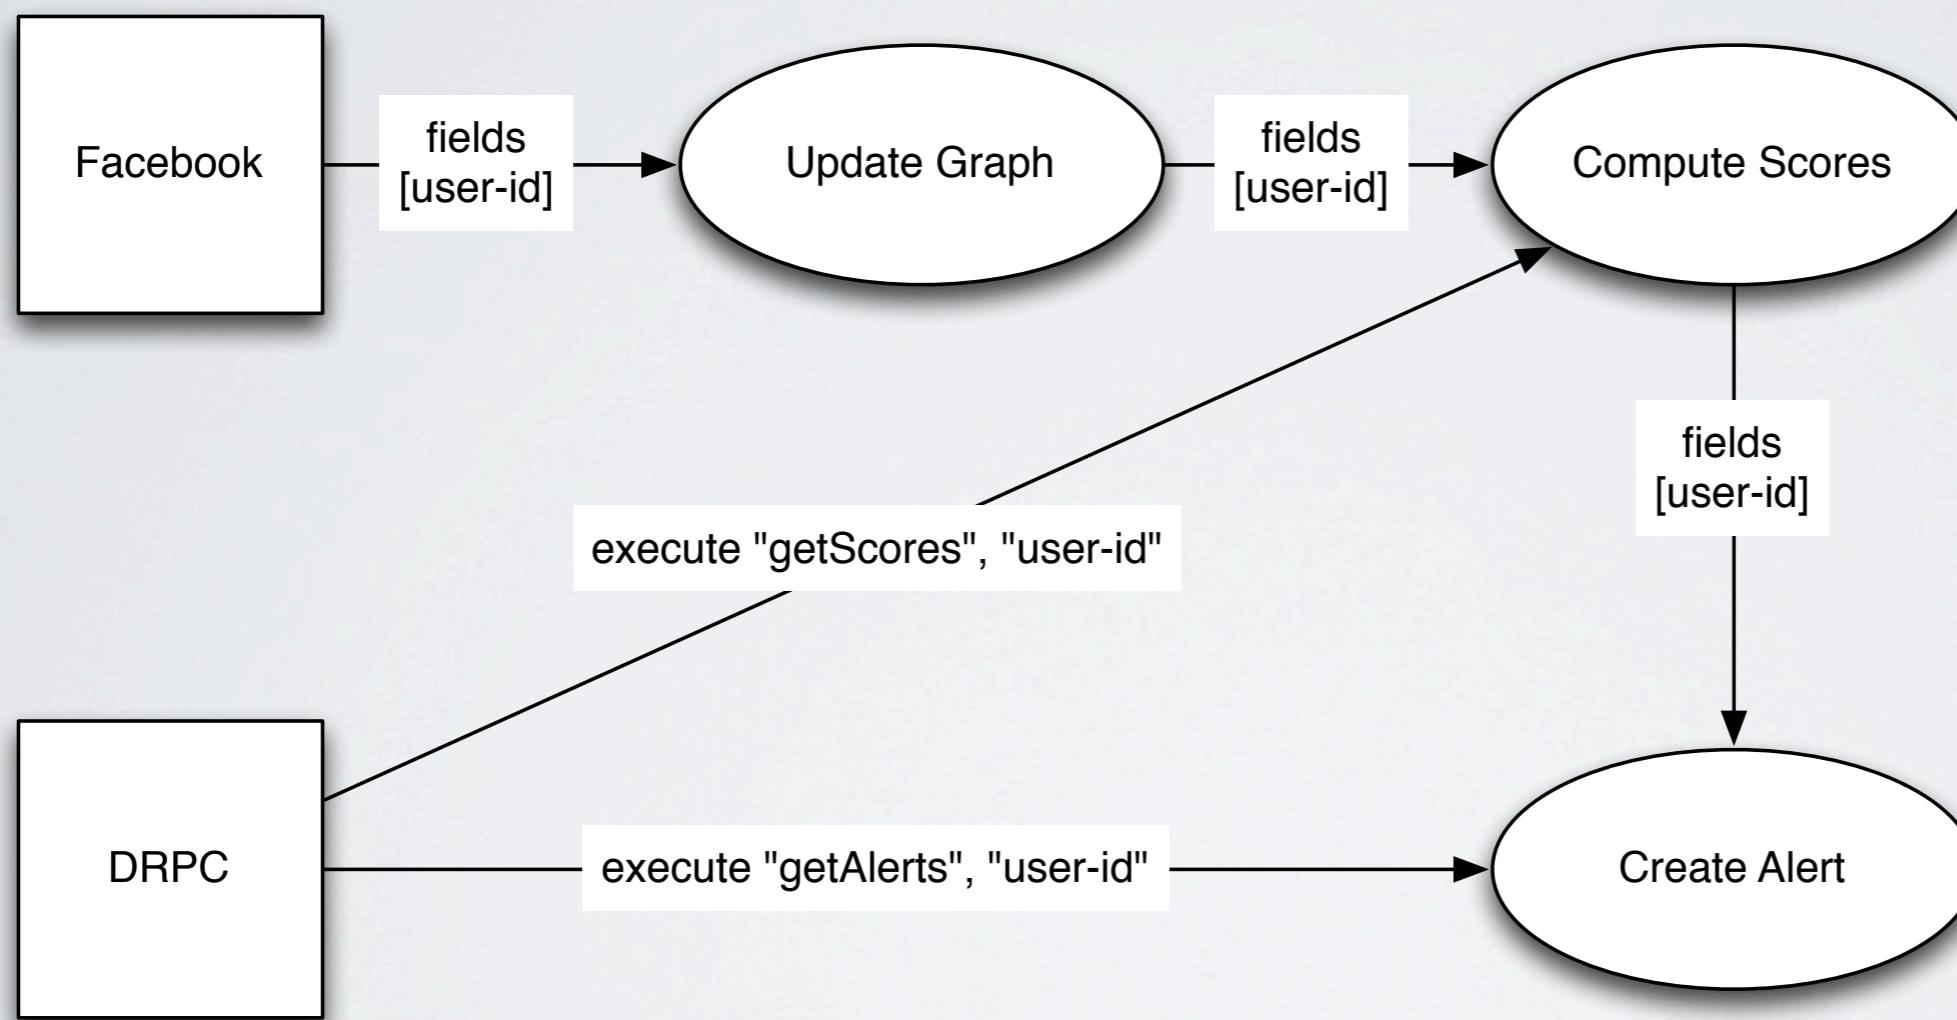


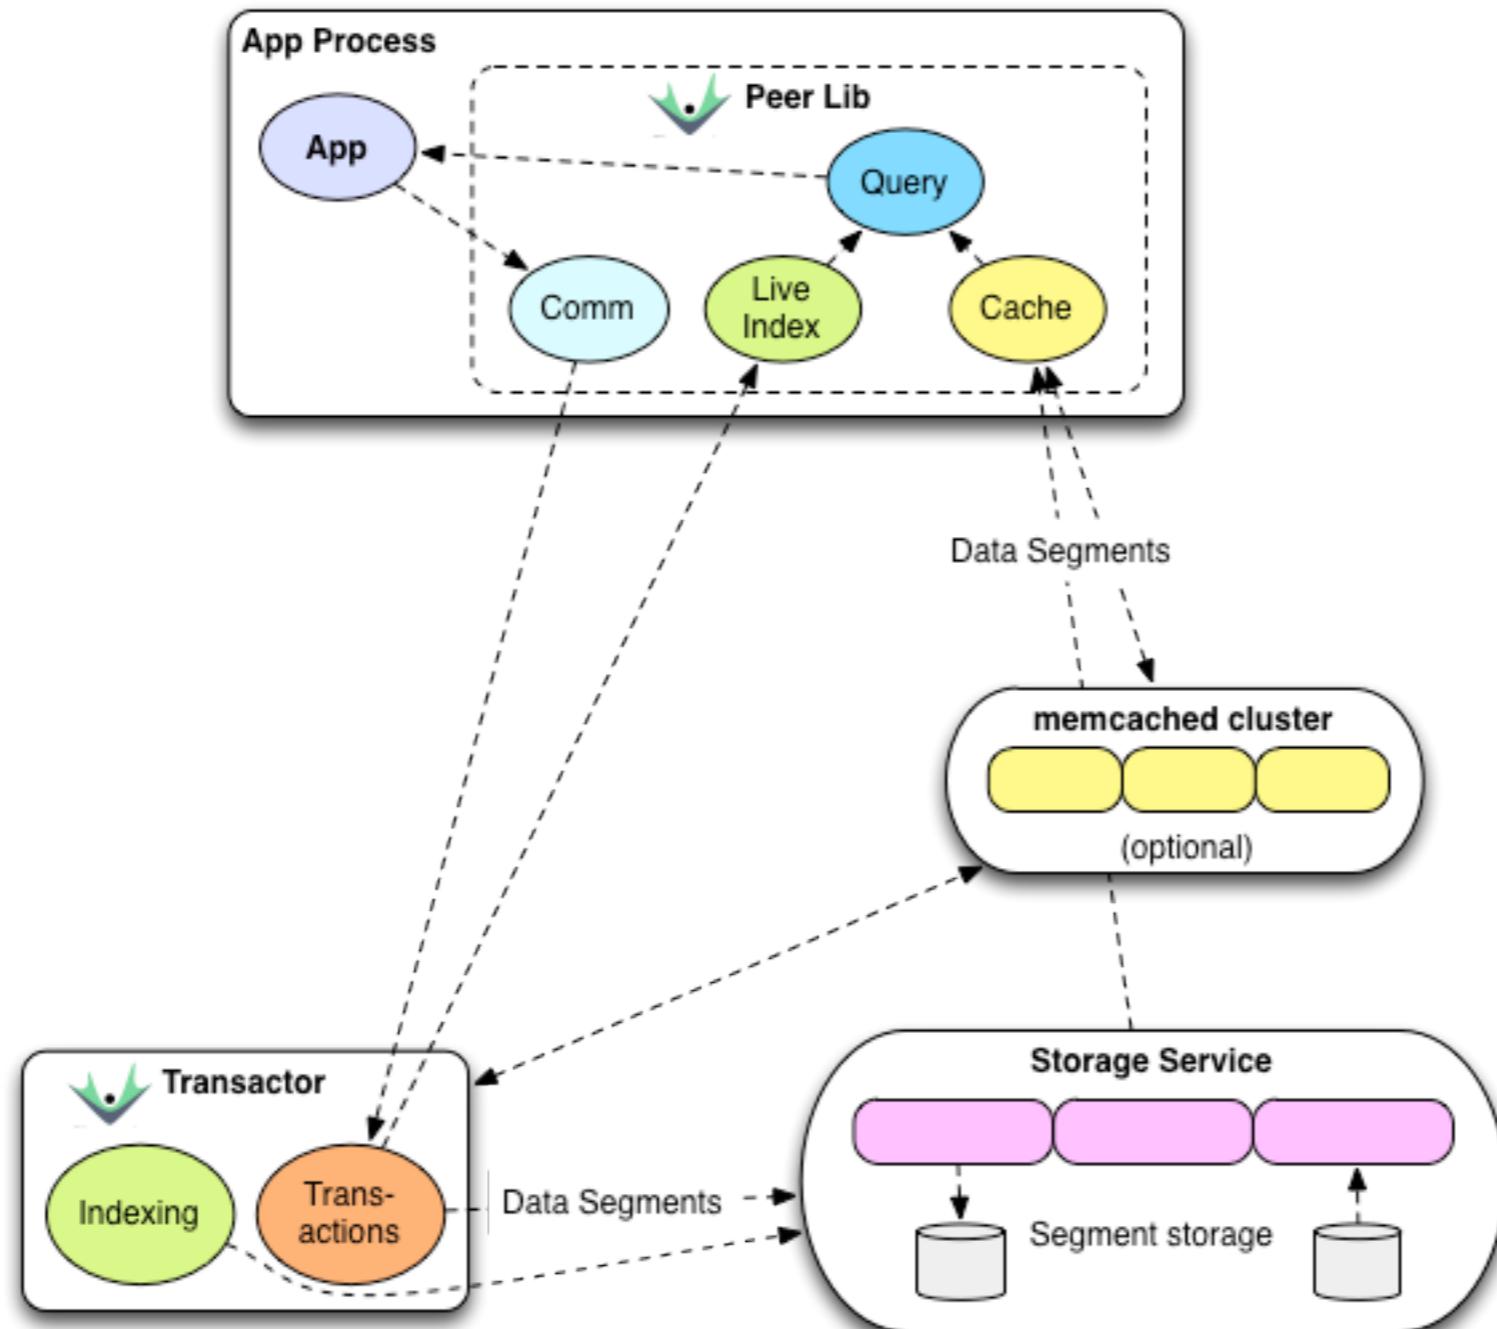
zolodeck

zolodeck snippet



zolodeck snippet

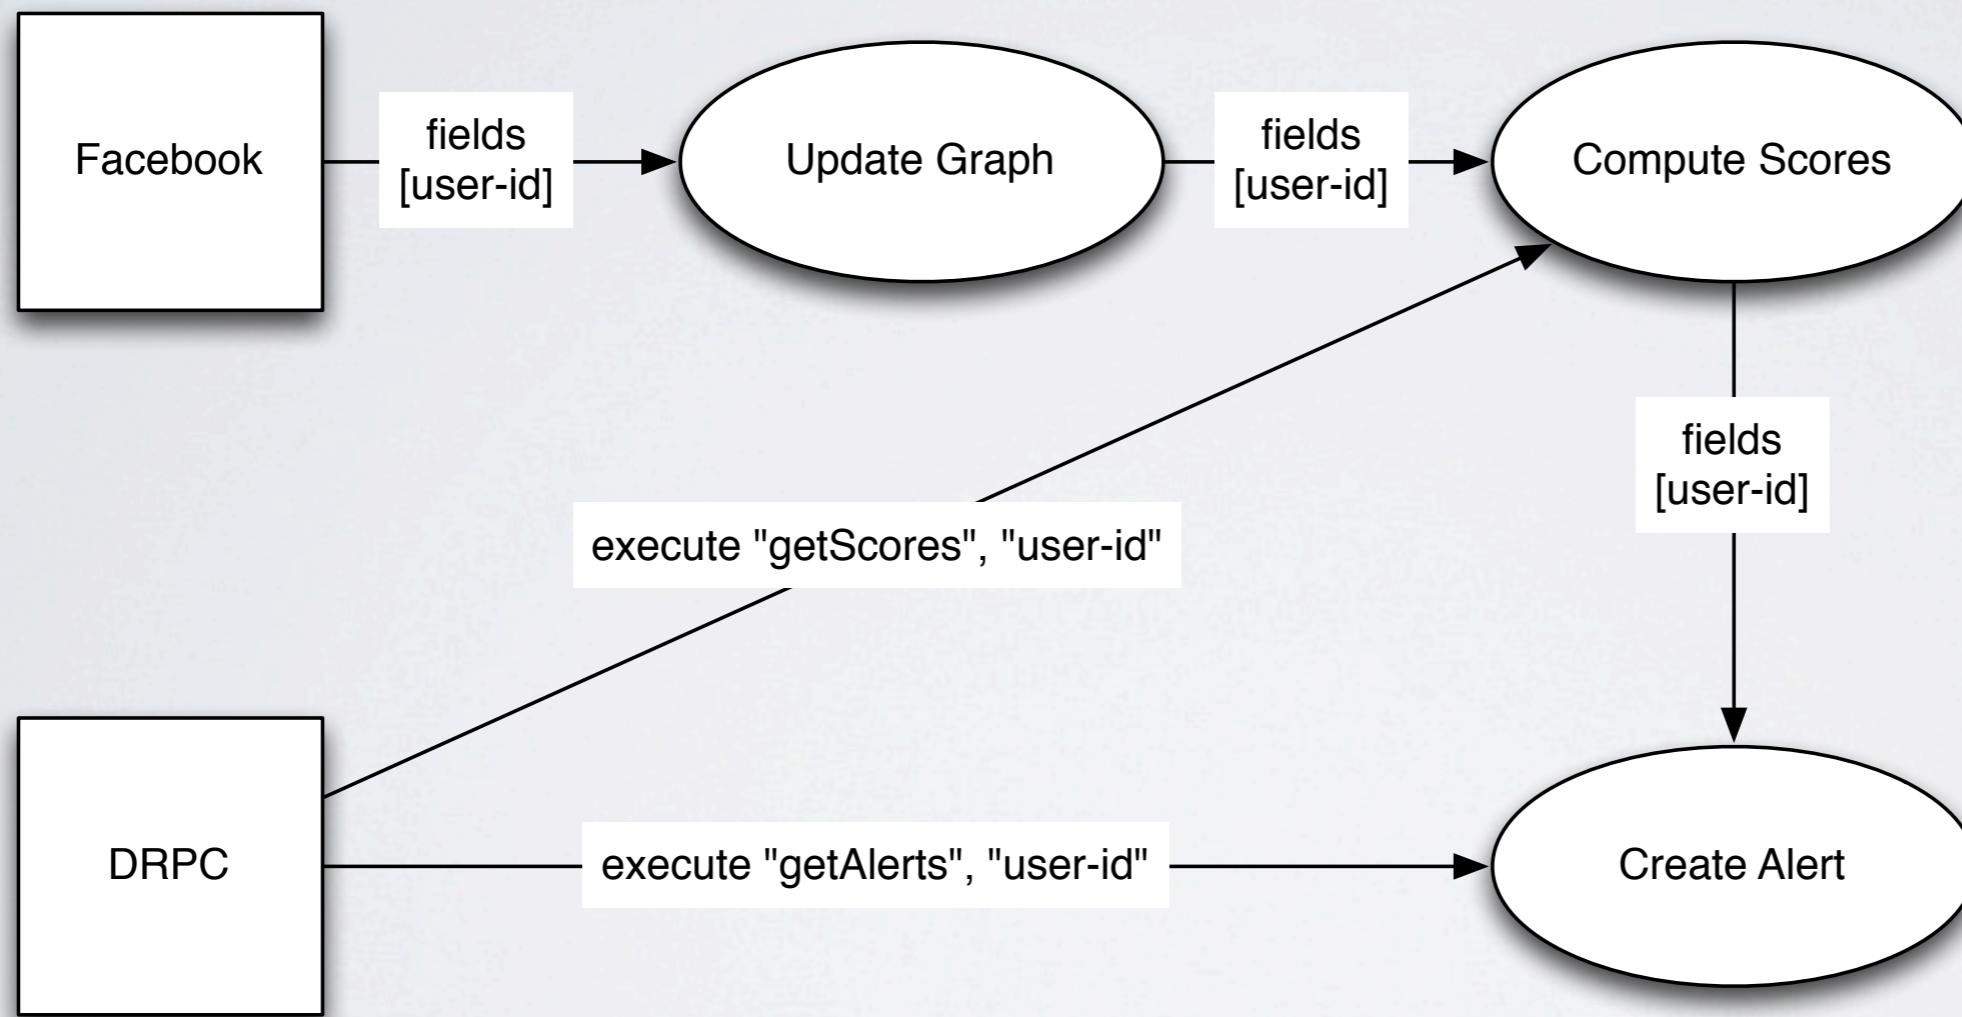




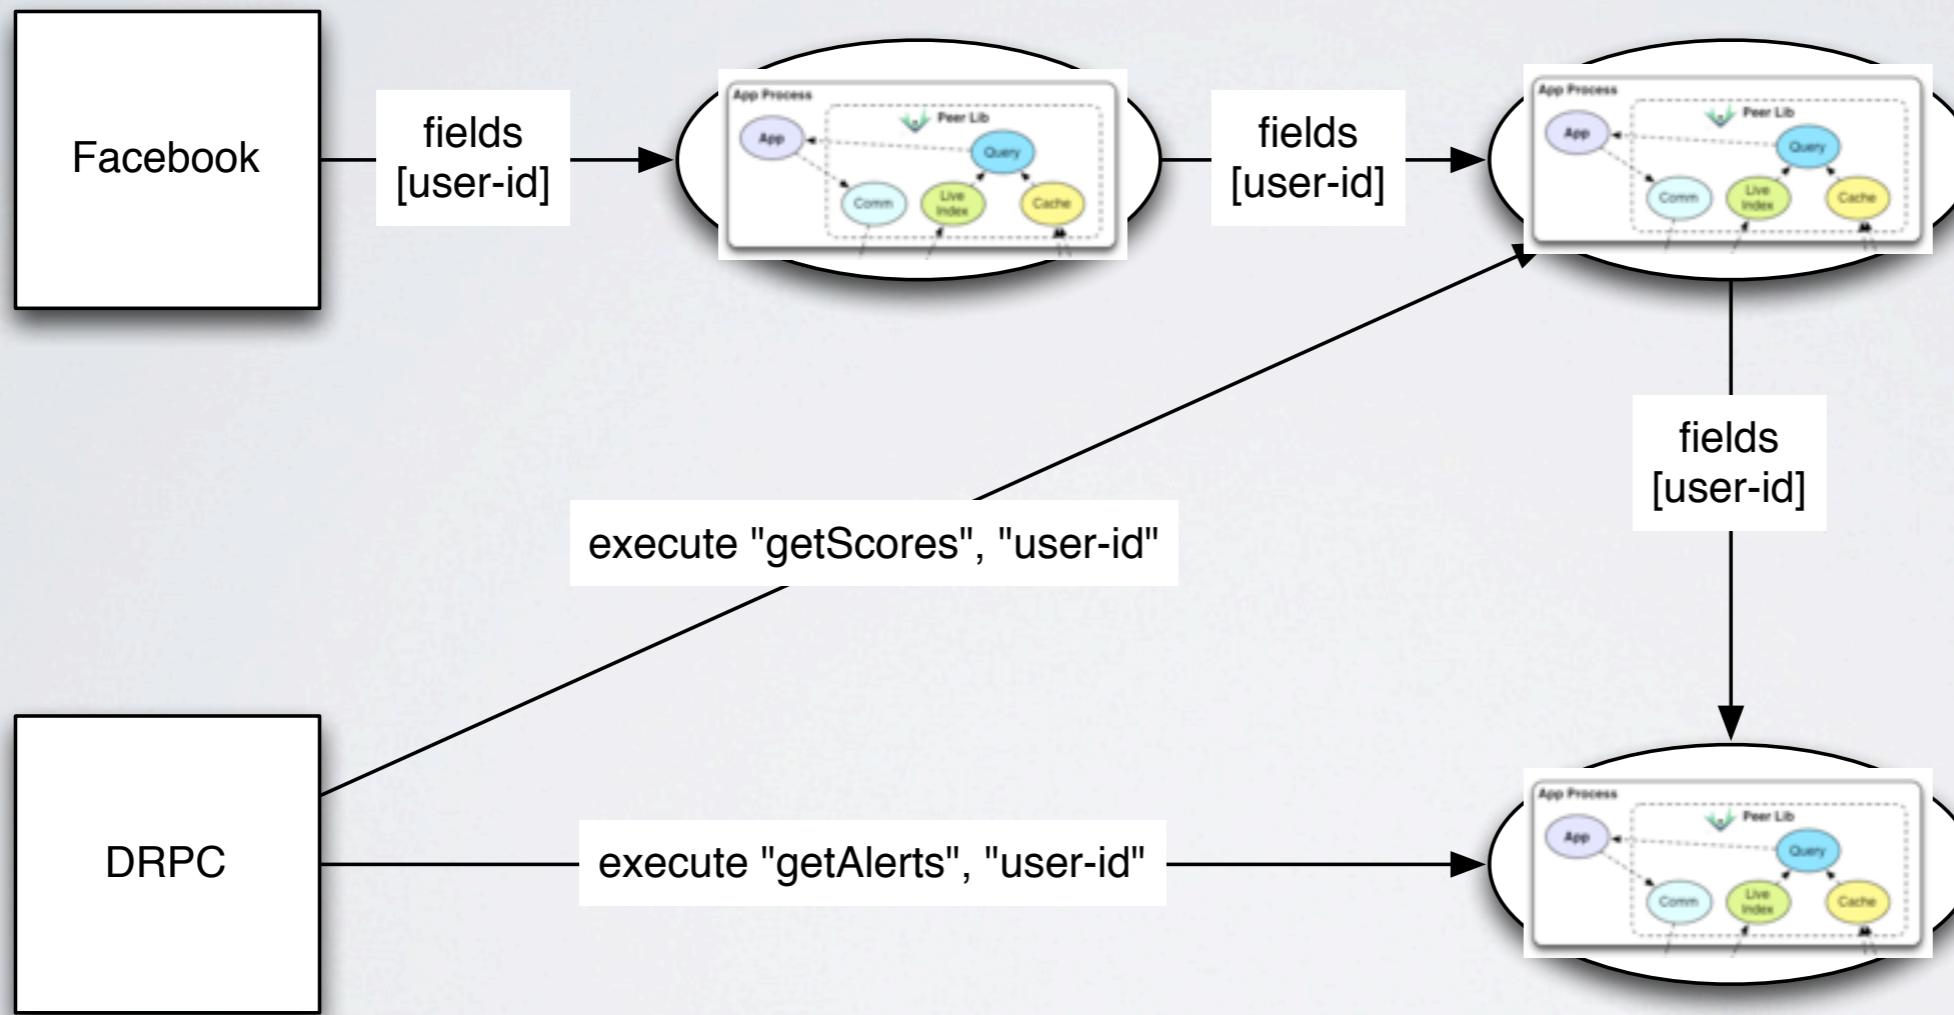
recap: the datomic way



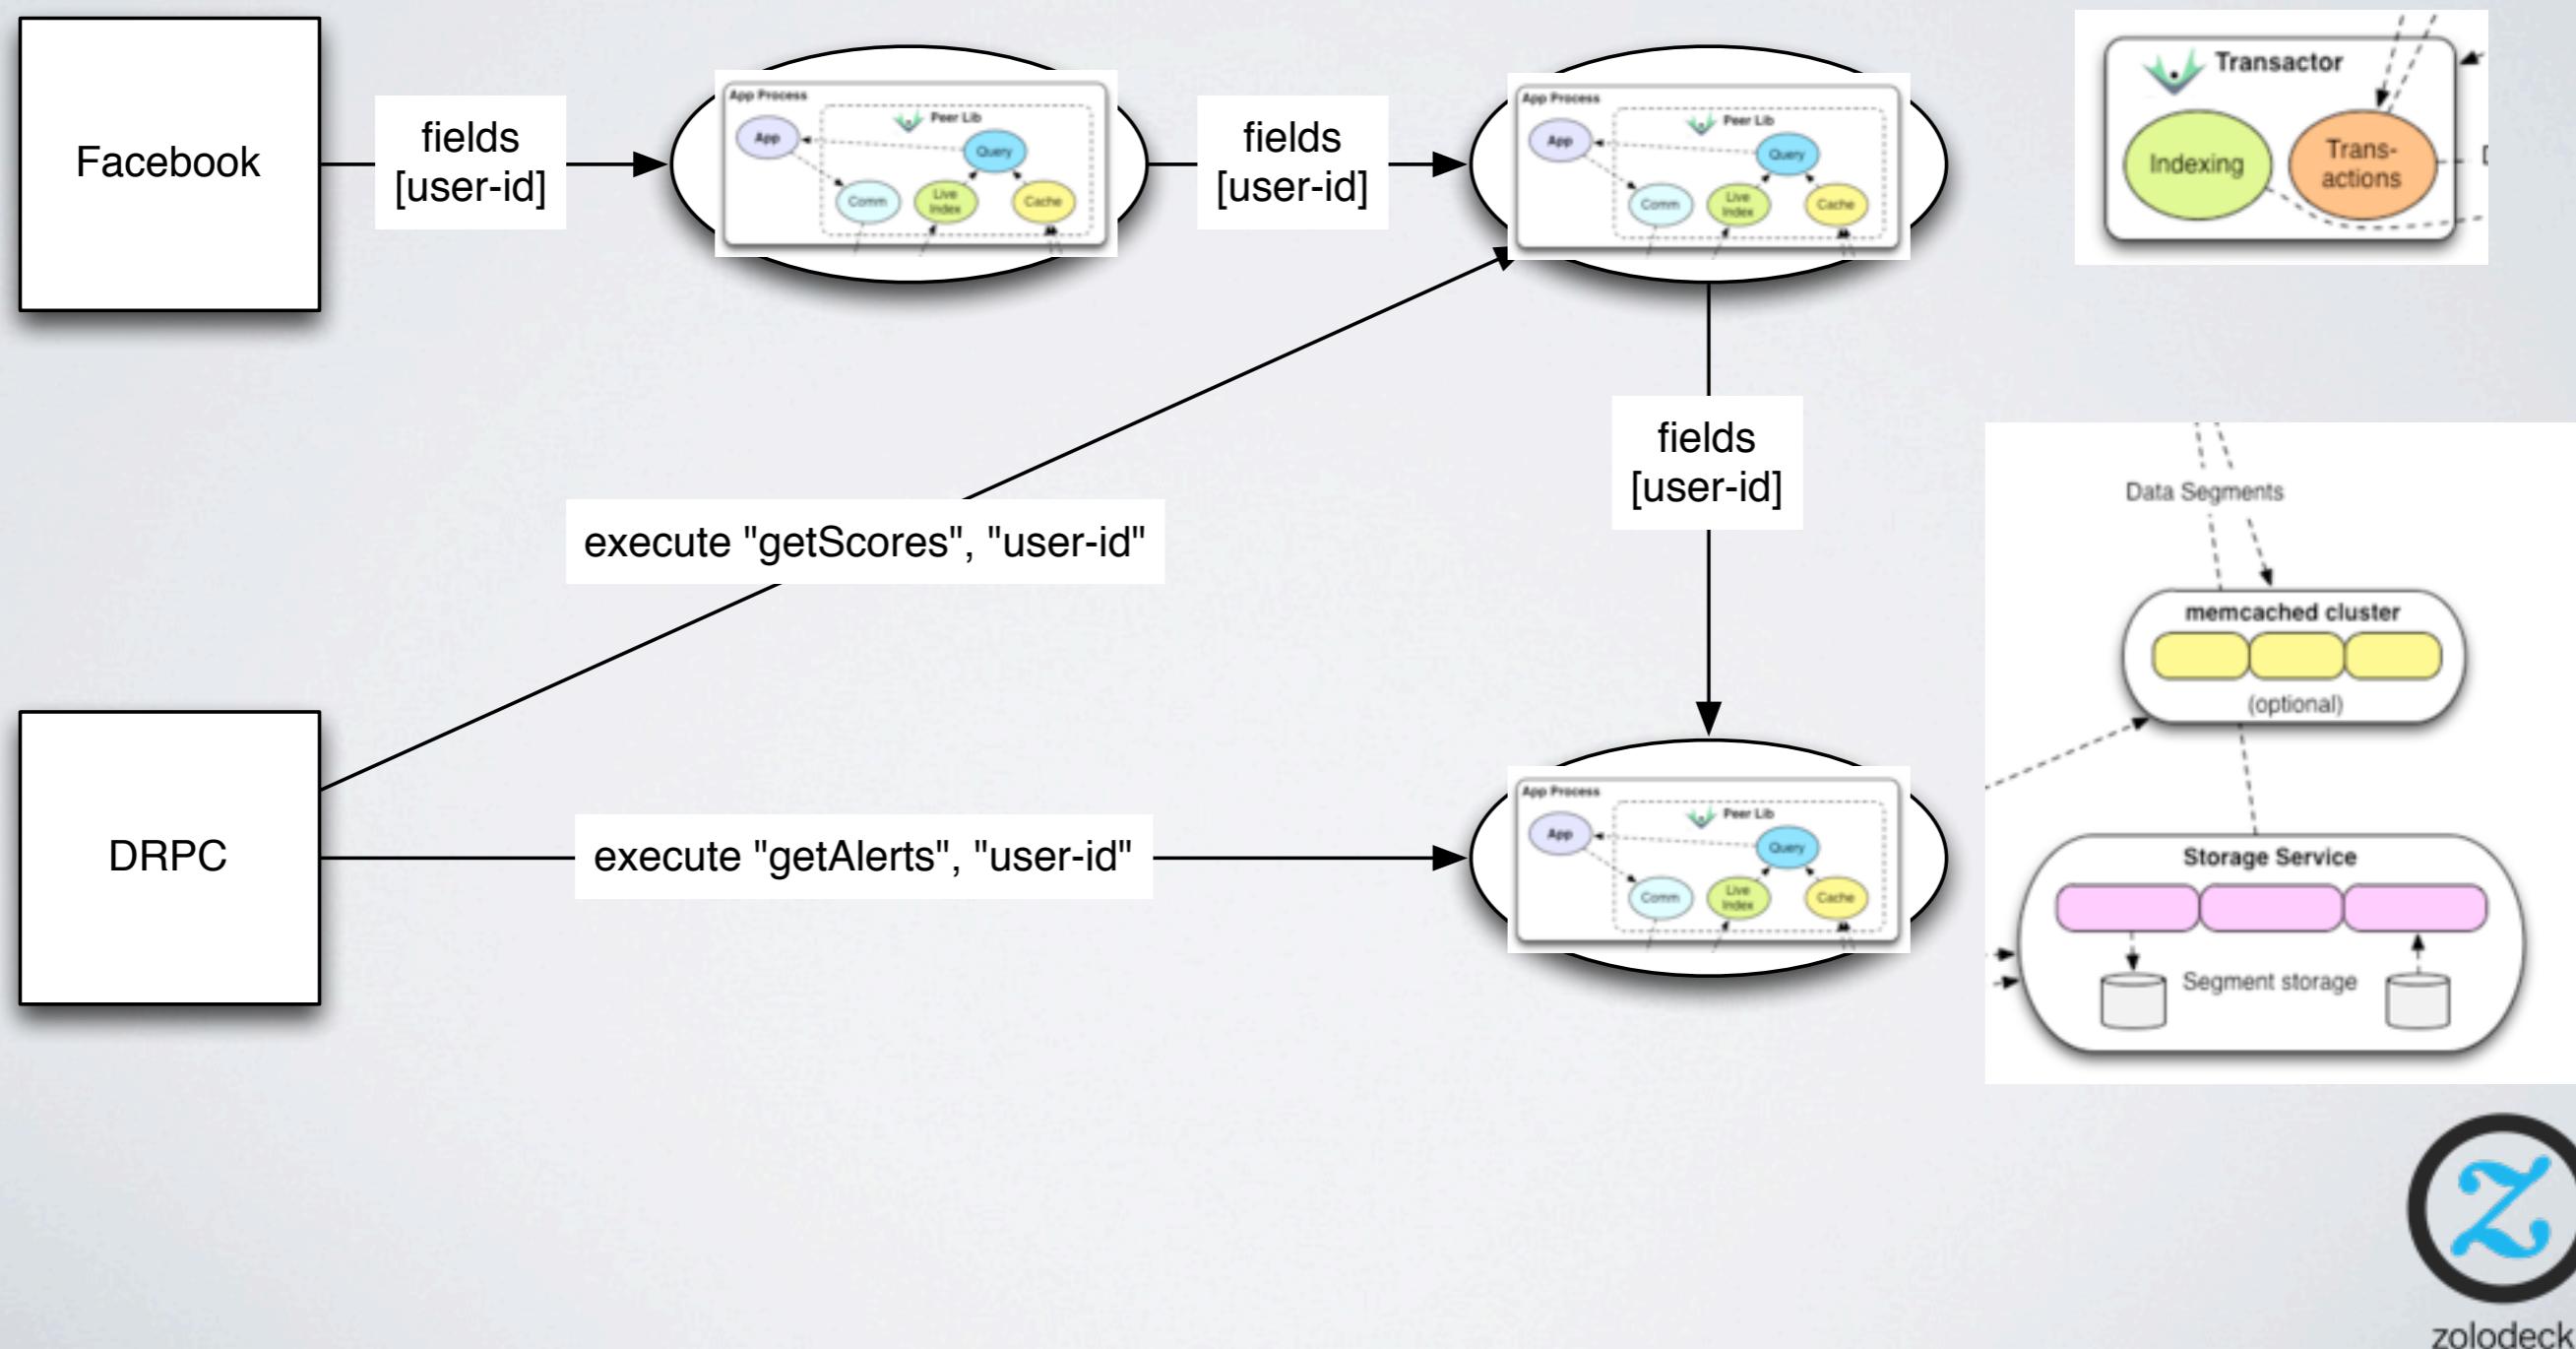
zolodeck snippet



zolodeck snippet



zolodeck snippet



clojure



functional language
modern Lisp
concurrent programming
JVM



simplicity



simplicity

functional abstractions



zolodeck

simplicity

functional abstractions
easy to test



simplicity

functional abstractions
easy to test
immutability



simplicity

functional abstractions
easy to test
immutability
macro system



datomic speaks clojure

written in Clojure
schema, transactions, datalog are data-structures



storm speaks Clojure



storm speaks Clojure

50% written in Clojure
beautiful DSL



datomic

consistent, scalable
no sharding
built-in caching

datalog power
db is local
peer isolation
schema flexibility
time travel

configuration-free
cloud-ready

storm

guaranteed processing
easy horizontal scalability
fault-tolerance

focus on business logic

stream processing
distributed RPC
continuous computation



- analytics peers won't bog down transaction processing
- datalog is truly powerful, whole another topic

zolodeck = datomic + storm + clojure



zolodeck

zolodeck = datomic + storm + clojure

work locally



zolodeck

zolodeck = datomic + storm + clojure

work locally
start small



zolodeck = datomic + storm + clojure

work locally

start small

scale effortlessly



zolodeck

zolodeck = datomic + storm + clojure

work locally

start small

scale effortlessly

~ no ops



zolodeck

zolodeck = datomic + storm + clojure

work locally

start small

scale effortlessly

~ no ops

focus on business logic



zolodeck

zolodeck = datomic + storm + clojure

work locally

start small

scale effortlessly

~ no ops

focus on business logic

become an networking extraordinaire



thanks

zolodeck.com

Your personal virtual assistant

