


# Graph: composable production systems in Clojure



Jason Wolfe (@w01fe)  
Strange Loop '12

# Motivation


- Interesting software has:
  - many components
  - **complex** web of **dependencies**
- Developers want:
  - simple, factored code
  - easy testability
  - tools for monitoring and debugging



# Graph

- **Graph** is a **simple, declarative** way to express system **composition**
- A Graph is just a map of functions that can depend on previous outputs
- Graphs are easy to create, reason about, test, and build upon

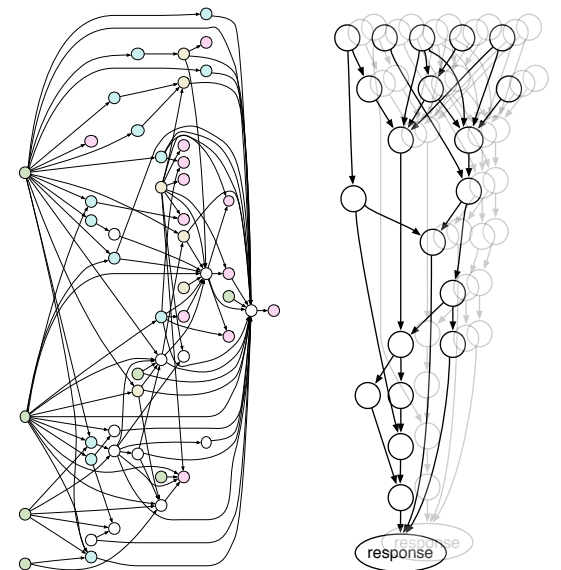
```
{:x (fnk [i] ...)  
 :y (fnk [j x] ...)  
 :z (fnk [x y] ...)}
```



input	output
{:i 1 :j 2}	{:x 2 :y 5 :z 12}

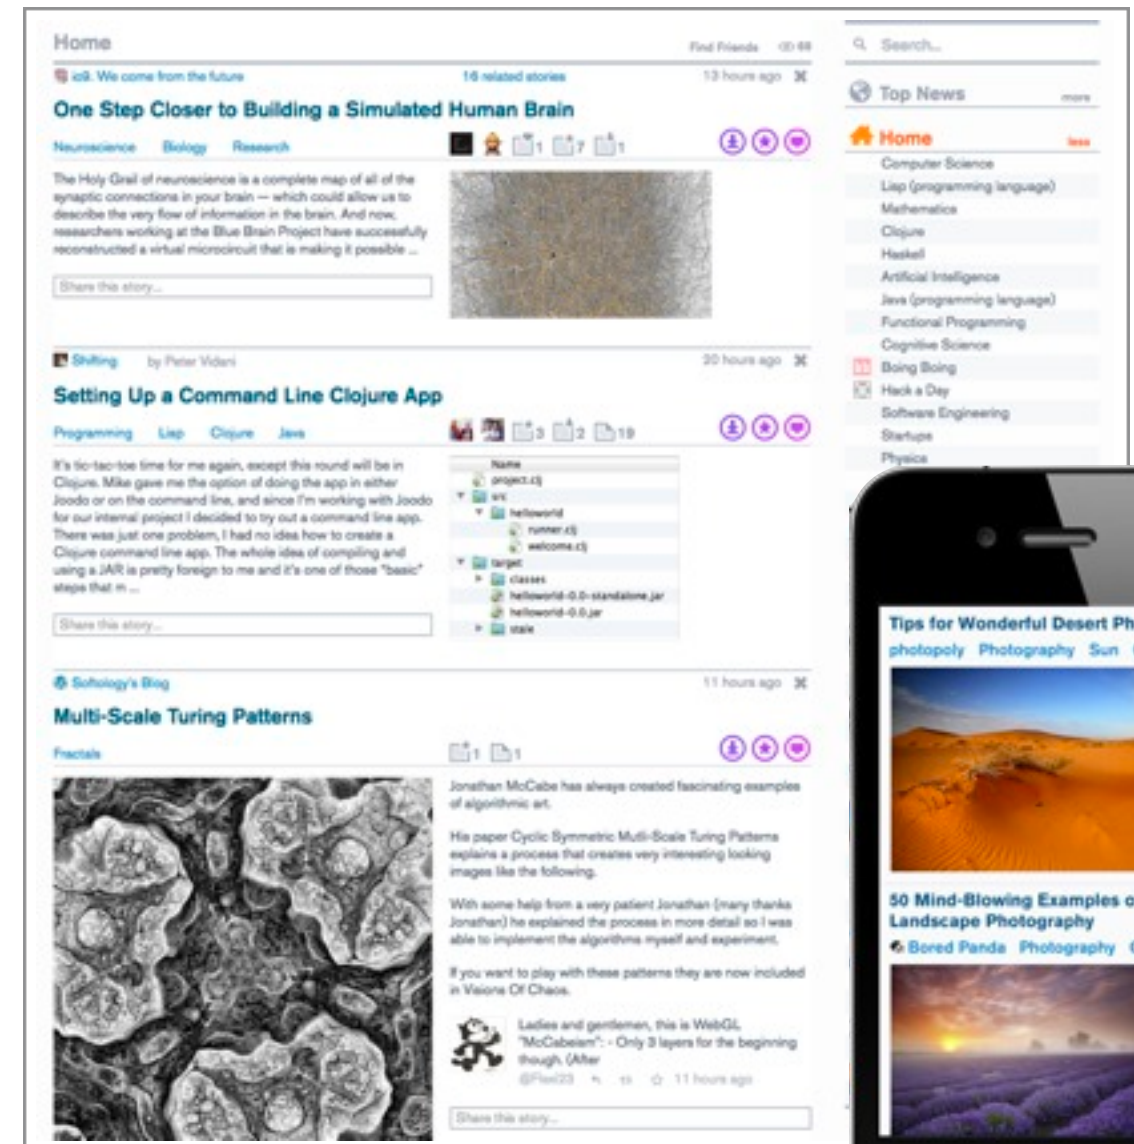
# Outline

- Prismatic
- Design Goals
- Graph: specs and compilation
- Applications
  - newsfeed generation
  - production services


$$\begin{aligned} &{:x} \text{ (fnk [i] \dots)} \\ &{:y} \text{ (fnk [j x] \dots)} \\ &{:z} \text{ (fnk [x y] \dots)} \end{aligned}$$


# Prismatic

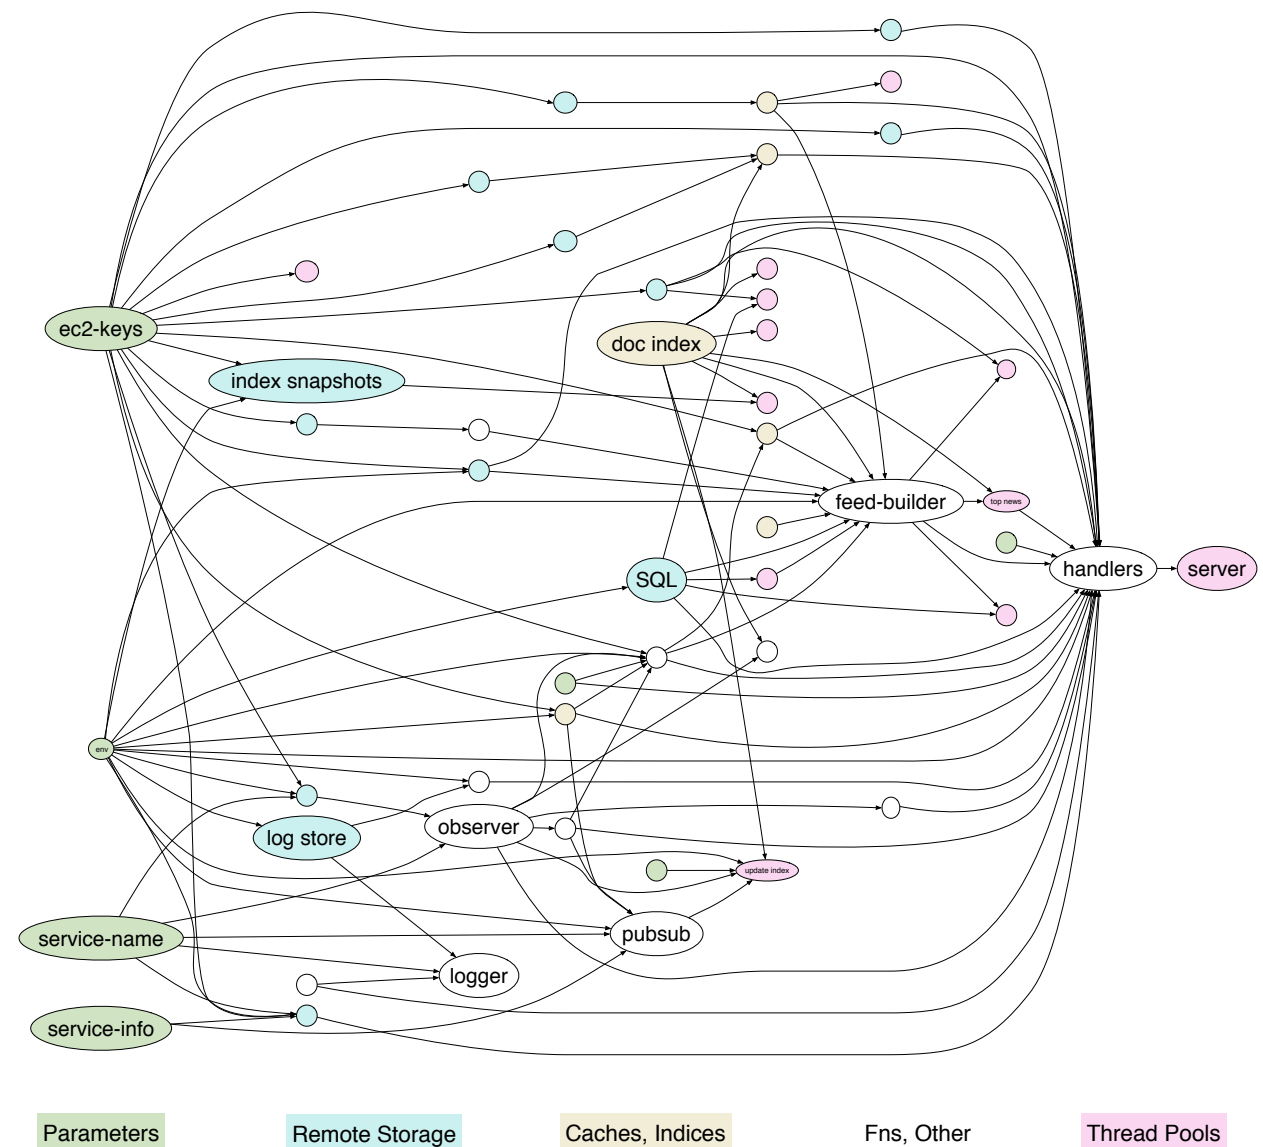
- Personalized, interest-based **newsfeeds**
- Build crawlers, topic models, graph analysis, story clustering, ...
- Backend 99.9% **Clojure**
- Personalized ranked feeds in real-time (~200ms)



getprismatic.com


# Prismatic's production API service

- >100 components
  - storage systems
  - caches & indices
  - ranking algorithms
- Coordinate in intricate dance to serve feeds *fast*
- Relentlessly refactored
- Still dozens of top-level components in **complex dependency network**




# The feed builder

- 20+ steps from query to personalized ranking, 20+ parameters
- Not a simple pipeline



# The feed builder

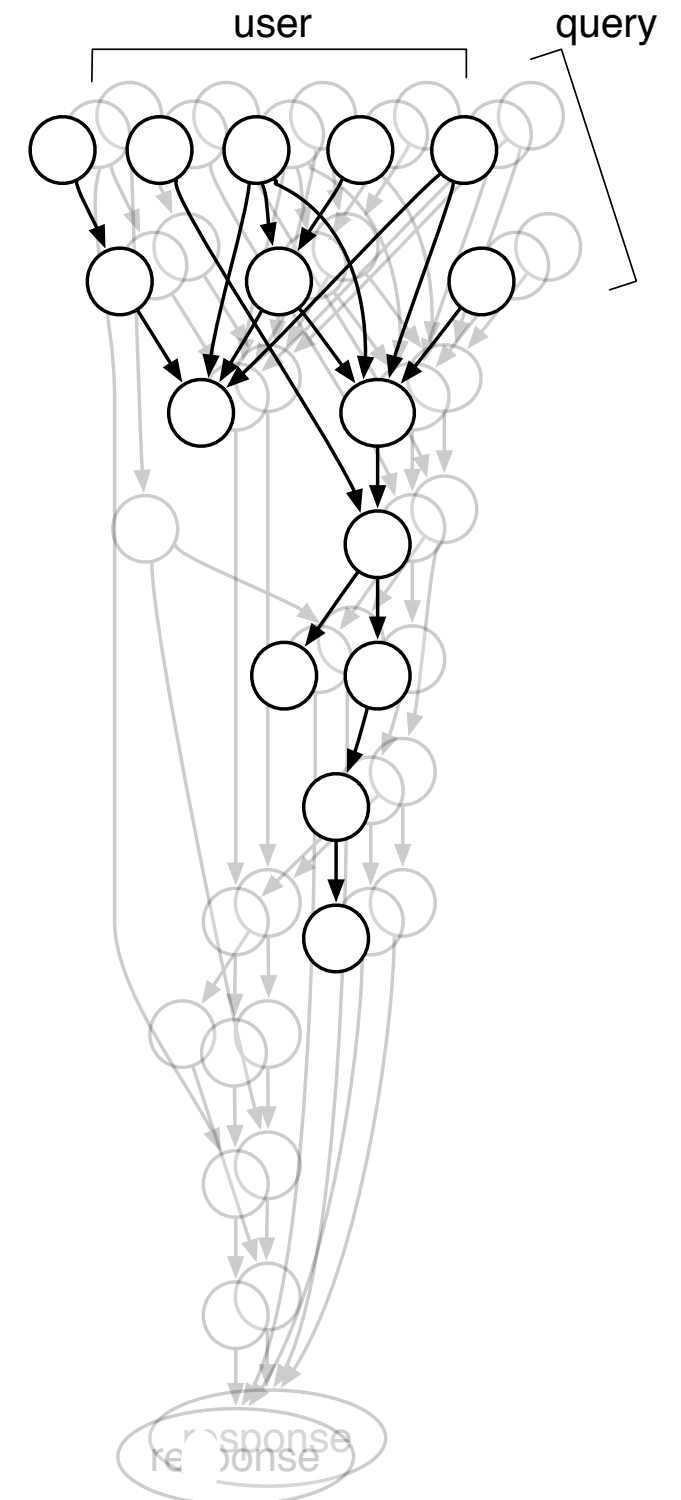
- 20+ steps from query to personalized ranking, 20+ parameters
- Not a simple pipeline
- > 10 feed types w/ slightly different steps, configurations







# The feed builder

- 20+ steps from query to personalized ranking, 20+ parameters
- Not a simple pipeline
- > 10 feed types w/ slightly different steps, configurations
- Support for early stopping




# Theme: complexity of composition

- Previous implementations: defns with huge  $\lambda$ ets
- **Unwieldy** for large systems with complex or polymorphic dependencies
- Hard to **test, debug, and monitor**



# The 'monster let'

- Tens of parameters, not compositional
- Mocks/polymorphic flow difficult
- Ad hoc monitoring & shutdown logic per item
- Core issue: structure of (de)composition is locked up in an **opaque** function



```
(defn start [{:keys [a z]}]  
  (let [s1 (store a ...)  
        s2 (store b ...)  
        db (sql-db c)  
        t2 (cron s2 db...)  
        ...  
        srv (server ...)]  
    (fn shutdown []  
      (.stop srv)  
      ...  
      (.flush s1))))
```

# Prismatic software engineering philosophy

- Fine-grained, composable abstractions (FCA)



Libraries >> Frameworks

- Strive for simplicity, work with the language
- Graph is a FCA for composition

# Goal: declarative


- **Declarative** specifications fix ‘monster Let’
  - Explicitly list **components, dependencies**
  - Enable abstractions over components, reasoning about composition
- Not new: Pregel, Dryad, Storm, ...

# Goal: simple

- Distill this idea to its **simplest**, most **idiomatic** expression
  - a Graph spec is just a (Clojure) map
  - no XML files or interface hell
- Graphs are ordinary data
  - manipulate them ‘for free’
  - --> unexpected applications

It is better to have 100 functions operate on one data structure than 10 functions on 10 data structures. - Alan Perlis

# From 'Let' to Graph



```
(defn stats [{:keys [xs]}]
  (let [n (count xs)
        m (/ (sum xs) n)
        m2 (/ (sum sq xs) n)
        v (- m2 (* m m))]
    {:n n :m m :m2 m2 :v v}))
```

```
{:n (fnk [xs] (count xs))
 :m (fnk [xs n] (/ (sum xs) n))
 :m2 (fnk [xs n] (/ (sum sq xs) n))
 :v (fnk [m m2] (- m2 (* m m)))}
```

# Bring on the fnk

- **fnk** = keyword function
- Similar to `{:keys []}` destructuring
  - nicer opt. arg. support
  - asserts that keys exist
  - **metadata** about args
- Quite useful in itself
- Only macros in Graph

```
(defnk foo [x y [s 1]]  
  (+ x (* y s)))
```

```
(= 8 (foo {:x 2 :y 3 :s 2}))
```

```
(= 5 (foo {:x 2 :y 3}))
```


```
(thrown? Ex. (foo {:x 2}))
```

```
(= (meta foo)  
  {:req-ks #{:x :y}}  
  :opt-ks #{:s}))
```



# A Graph Specification


- A Graph is just a **map from keywords to fnks**
- Required keys of each fnk specify graph relationships
- Entire **graph specifies a fnk** to map of results



```
{:n (fnk [xs]
      (count xs))
:m  (fnk [xs n]
      (/ (sum xs) n))
:m2 (fnk [xs n]
      (/ (sum sq xs) n))
:v  (fnk [m m2]
      (- m2 (* m m)))}
```

# A Graph Specification

- A Graph is just a **map from keywords to fnks**
- Required keys of each fnk specify graph relationships
- Entire **graph specifies a fnk** to map of results



```
{:xs [1 2 3 6]}
```

```
{:n (fnk [xs]  
      (count xs))
```


```
:m (fnk [xs n]  
      (/ (sum xs) n))
```

```
:m2 (fnk [xs n]  
      (/ (sum sq xs) n))
```

```
:v (fnk [m m2]  
      (- m2 (* m m))))}
```

# A Graph Specification

- A Graph is just a **map from keywords to fnks**
- Required keys of each fnk specify graph relationships
- Entire **graph specifies a fnk** to map of results



```
{:xs [1 2 3 6]}
```

```
{:n 4}
```


```
:m (fnk [xs n]  
      (/ (sum xs) n))
```

```
:m2 (fnk [xs n]  
       (/ (sum sq xs) n))
```

```
:v (fnk [m m2]  
      (- m2 (* m m)))}
```

# A Graph Specification

- A Graph is just a **map from keywords to fnks**
- Required keys of each fnk specify graph relationships
- Entire **graph specifies a fnk** to map of results



```
{:xs [1 2 3 6]}
```

```
{:n 4
```


```
:m 3
```

```
:m2 (fnk [xs n]  
      (/ (sum sq xs) n))
```

```
:v (fnk [m m2]  
      (- m2 (* m m))))}
```

# A Graph Specification

- A Graph is just a **map from keywords to fnks**
- Required keys of each fnk specify graph relationships
- Entire **graph specifies a fnk** to map of results



`{:xs [1 2 3 6]}`

`{:n 4}`


`:m 3`

`:m2 12.5`

`:v (fnk [m m2]  
(- m2 (* m m)))}`

# A Graph Specification

- A Graph is just a **map from keywords to fnks**
- Required keys of each fnk specify graph relationships
- Entire **graph specifies a fnk** to map of results



`{:xs [1 2 3 6]}`

`{:n 4`

`:m 3`

`:m2 12.5`

`:v 3.5`

`}`

# Compiling Graphs

- Compile graph to fnk that returns map of outputs

```
(def g  
  {:n (fnk [xs] ...)  
   :m (fnk [xs n] ...)  
   :m2 (fnk [xs n] ...)  
   :v (fnk [m m2] ...)})
```

```
(def stats  
  (compile g))
```

```
(= (stats {:xs [1 2 3 6]})  
   {:n 4      :m 3  
    :m2 12.5  :v 3.5})
```

# Compiling Graphs

- Compile graph to fnk that returns map of outputs
- error checked

```
(def g  
  {:n (fnk [xs] ...)  
   :m (fnk [xs n] ...)  
   :m2 (fnk [xs n] ...)  
   :v (fnk [m m2] ...)})
```

```
(def stats  
  (compile g))
```

```
(thrown?  
  (Ex. "missing :xs")  
  (stats {:x 1}))
```



# Compiling Graphs

- Compile graph to fnk that returns map of outputs
- error checked
- can return lazy map

```
(def g  
  {:n (fnk [xs] ...)  
   :m (fnk [xs n] ...)  
  :m2 (fnk [xs n] ...)  
  :v (fnk [m m2] ...)})
```

```
(def stats  
  (lazy-compile g))
```

```
(= (:m (stats {:xs [1 5]}))  
   3)
```

# Compiling Graphs

- Compile graph to fnk that returns map of outputs
- error checked
- can return lazy map
- can auto-parallelize

```
(def g
  {:n (fnk [xs] ...)}
  :m (fnk [xs n] ...)}
  :m2 (fnk [xs n] ...)}
  :v (fnk [m m2] ...))
```

```
(def stats
  (par-compile g))
```

```
(= (:v (stats {:xs [1 5]}))
   3.5)
```

# Compiling Graphs

- Compile graph to fnk that returns map of outputs
  - error checked
  - can return lazy map
  - can auto-parallelize

```
(def g
  {:n 2
   :m (fnk [xs n] ...)
   :m2 (fnk [xs n] ...)
   :v (fnk [m m2] ...)})
```

```
(def stats
  (par-compile g))
```

```
(= (:v (stats {:xs [1 5]}))
   3.5)
```

# Compiling Graphs

- Compile graph to fnk that returns map of outputs
  - error checked
  - can return lazy map
  - can auto-parallelize

```
(def g
  {:n 2
   :m 3
   :m2 13
   :v (fnk [m m2] ...)})
```

```
(def stats
  (par-compile g))
```

```
(= (:v (stats {:xs [1 5]}))
   3.5)
```

# Compiling Graphs

- Compile graph to fnk that returns map of outputs
  - error checked
  - can return lazy map
  - can auto-parallelize

```
(def g  
  {:n 2  
   :m 3  
   :m2 13  
   :v 4  
  })
```

```
(def stats  
  (par-compile g))
```

```
(= (:v (stats {:xs [1 5]}))  
   3.5)
```

# Compiling Graphs

- Compile graph to fnk that returns map of outputs
  - error checked
  - can return lazy map
  - can auto-parallelize
- With more tooling, also compile graphs to production services
- Could compile to cross-machine topologies, ...


```
(def g  
  {:n 2  
   :m 3  
   :m2 13  
   :v 4  
  })
```

```
(def stats  
  (par-compile g))
```

```
(= (:v (stats {:xs [1 5]}))  
   3.5)
```


# Before: feed builder

- Real-time personally ranked feeds
- 100-line fn expressed core composition logic, ~20 params
  - several nested lets, escape hatches
- Component polymorphism (10 flavors of feeds)
  - kludge of cases
  - ball of multimethods
  - protocols + hacks



# Feed builder in Graph

- Default parameters
- Graph with 'holes' captures shared logic



```
(def default-params  
  {:alpha 0.7  
   ...  
   :phasers :stun})
```

```
(def partial-graph  
  {:query (fnk ...)  
   ...  
   :y (fnk [a x] ..)  
   ...  
   :resp (fnk ...)})
```



# Feed builder in Graph

- Each feed type specifies
  - updated parameters
  - missing/new graph nodes
- To make feed fn, just
  - merge in updates
  - compile resulting graph

```
(def default-params ..)
(def partial-graph ..)
```

```
(def topic-feed
  (compile-feed-fn
    {:alpha 0.2}
    {:x (fnk ...)}
    :q (fnk ...))))
```

```
(defn compile-feed-fn [params nodes]
  (let [p (merge default-params params)
        g (compile (merge partial-graph nodes))]
    (fn feed [req] (g (merge p req)))))
```

# After: feed builder

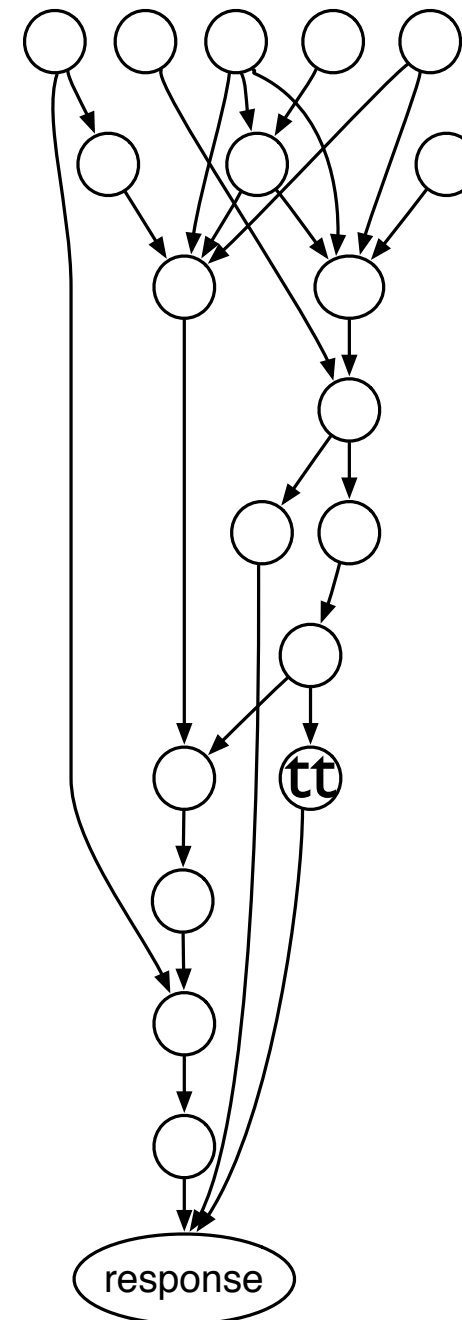
- Simpler, cleaner code
- Polymorphism is trivial

```
(def topic-feed  
  (compile-feed-fn  
    {:alpha 0.2}  
    {:x (fnk ...)  
     :q (fnk ...)})))
```

```
(def home-feed  
  (compile-feed-fn  
    {:alpha 0.4}  
    {:x (fnk ...)  
     :r (fnk ...)  
     :s (fnk ...)})))
```

# After: feed builder


- Simpler, cleaner code
- Polymorphism is trivial
- Early stopping for free via lazy compilation



# After: feed builder

- Simpler, cleaner code
- Polymorphism is trivial
- Early stopping for free via lazy compilation


```
(let [h (home-feed req)]  
  (:tt h))
```



# After: feed builder

- Simpler, cleaner code
- Polymorphism is trivial
- Early stopping for free via lazy compilation


```
(let [h (home-feed req)]  
  [(:tt h)  
   (:v h)])
```



# Also: easy to analyze

- Detect mis-wirings at graph compile time
  - positional constructor
- Avoid wrong # of args errors, arg ordering bugs
- Visualize graphs in 5 loc

```
(defn edges [graph]
  (for [[k f] graph
        :let [{:keys [req-ks opt-ks]} (meta f)]
        parent (concat req-ks opt-ks)]
    [parent k]))
```




# Also: easy to monitor

- Add **monitoring** and error reporting by mapping over fnks
- Since a Graph is a Map, can just use map-vals

node	n	avg ms	errors
<b>:fetch</b>	2500	1.5	0
<b>:rank</b>	1001	150.0	1
<b>:client</b>	1000	70.0	0

```
(defn observe-graph [g]
  (into {}
    (for [[k f] g]
      [k
        (with-meta
          (fn [m]
            (let [v (f m)]
              (print k m v)
              v))
          (meta f))]))))
```

# Example 2: production API service



```
(def api-service
  (service
    {:service-name "api"
     :backend-port 42424
     :server-threads 100}
    {:store1 (instance store
                      {:type :s3 ...})
     :memo (fnk [store1]
               {:resource ...})
     ...
     :api-server (...)}))
```



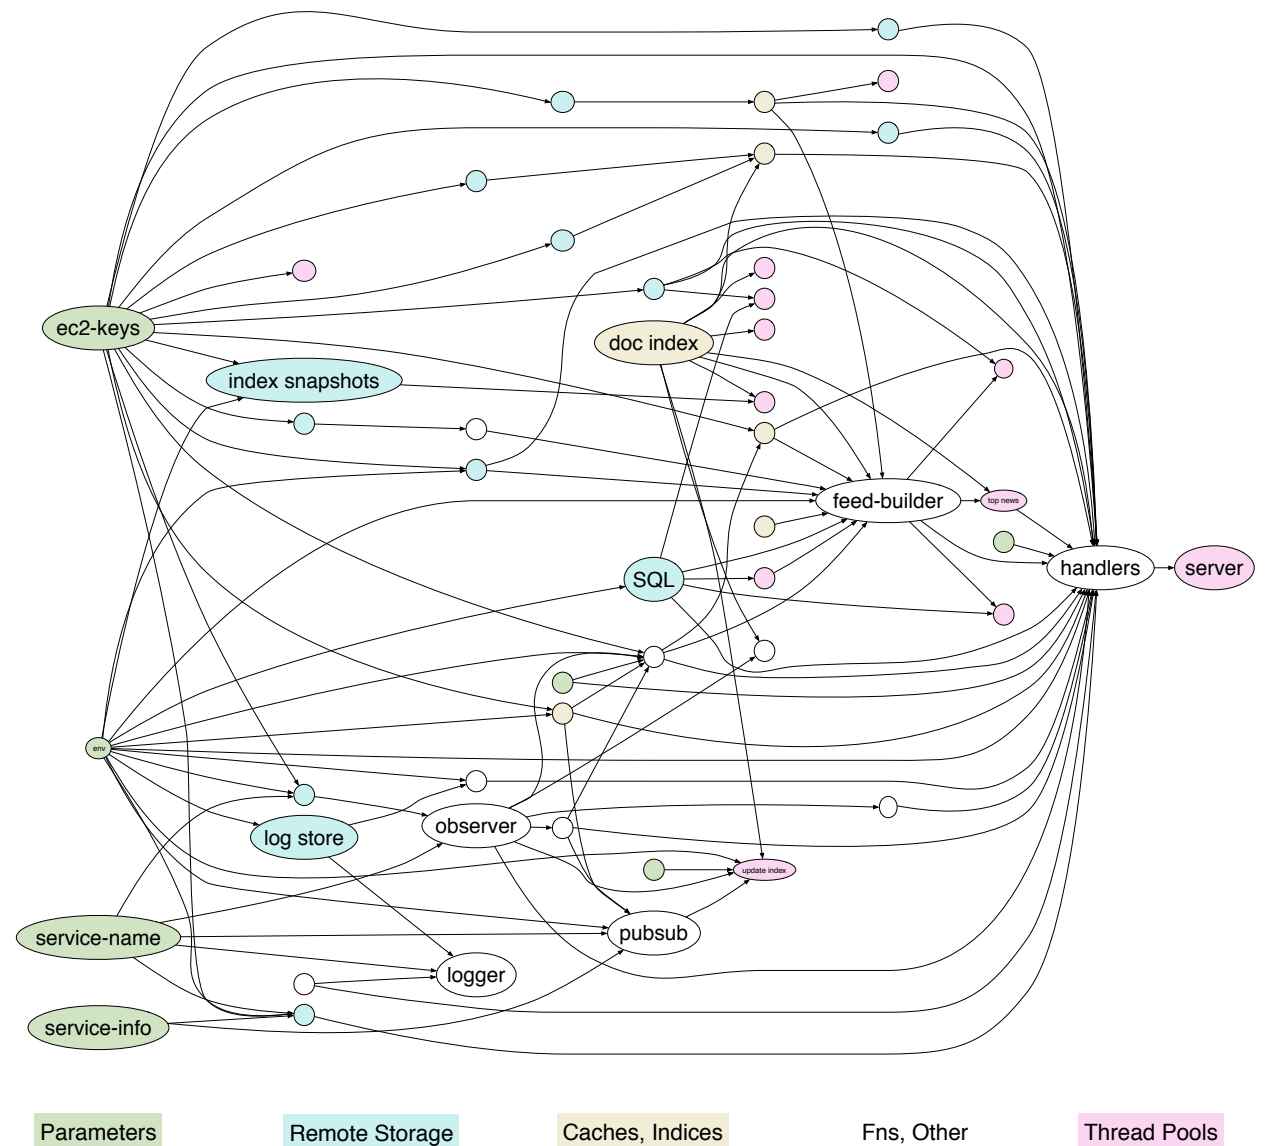
# Service definitions

- Service definition =
  - parameter map +
  - resource graph
- Crane reads params for provisioning, deployment
- Graph = service code
  - parameters are args
  - cron jobs, handlers at leaves

```
(def api-service
  (service
    {:service-name "api"
     :backend-port 42424
     :server-threads 100}
    {:store1 (instance store
                      {:type :s3 ...})
     :memo   (fnk [store1]
                {:resource ...})
     ...
    :api-server (...))}))
```

# Service definitions

- Service definition =
  - parameter map +
  - resource graph
- Crane reads params for provisioning, deployment
- Graph = service code
  - parameters are args
  - cron jobs, handlers at leaves



# Service built-ins

- Parameters and graph nodes available by convention
- Interface with deployment, other services, dashboard
- Smartly reconfigure with `env -- test/staging/prod`

## parameters

```
{:env           :prod  
 :instance-id  "i-123abc"  
 :ec2-keys     ...      }
```

## resources

```
{:nameserver   ...  
 :observer     ...  
 :pubsub       ...      }
```

# Nodes build Resources

- Resource = component
  - e.g., database, cache, fn
  - Plus metadata for shutdown, handlers, ...
  - Represent as a map
- Library of resources that work with builtins
  - data stores
  - processing queues
  - recurring tasks
  - ...

```
(defnk refreshing-atom
  [f period]
  (let [a (atom (f))
        e (Exec/newExec)]
    (.schedAtFixedRate e
      #(reset! a (f))
      period)
    {:res a
     :shutdown #(.sd e)})))
```

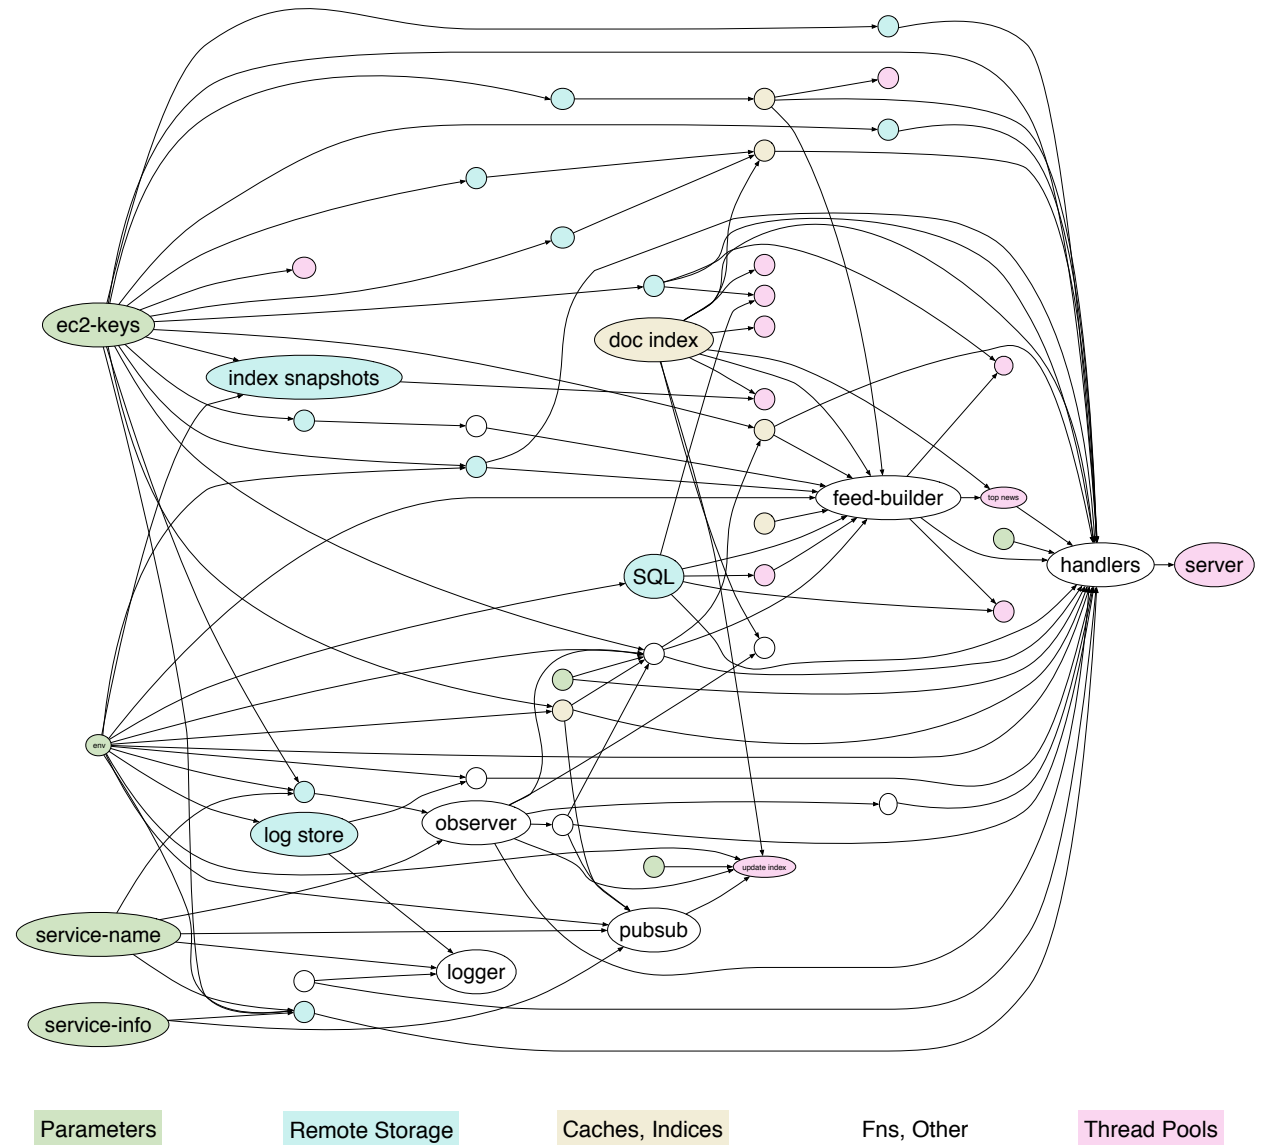
# Starting and Stopping

- Transform resource graph to ordinary graph
  - map over leaves, pull out `:resource`
  - assoc new `:shutdown` key
- Run graph to start service, get clean shutdown hook

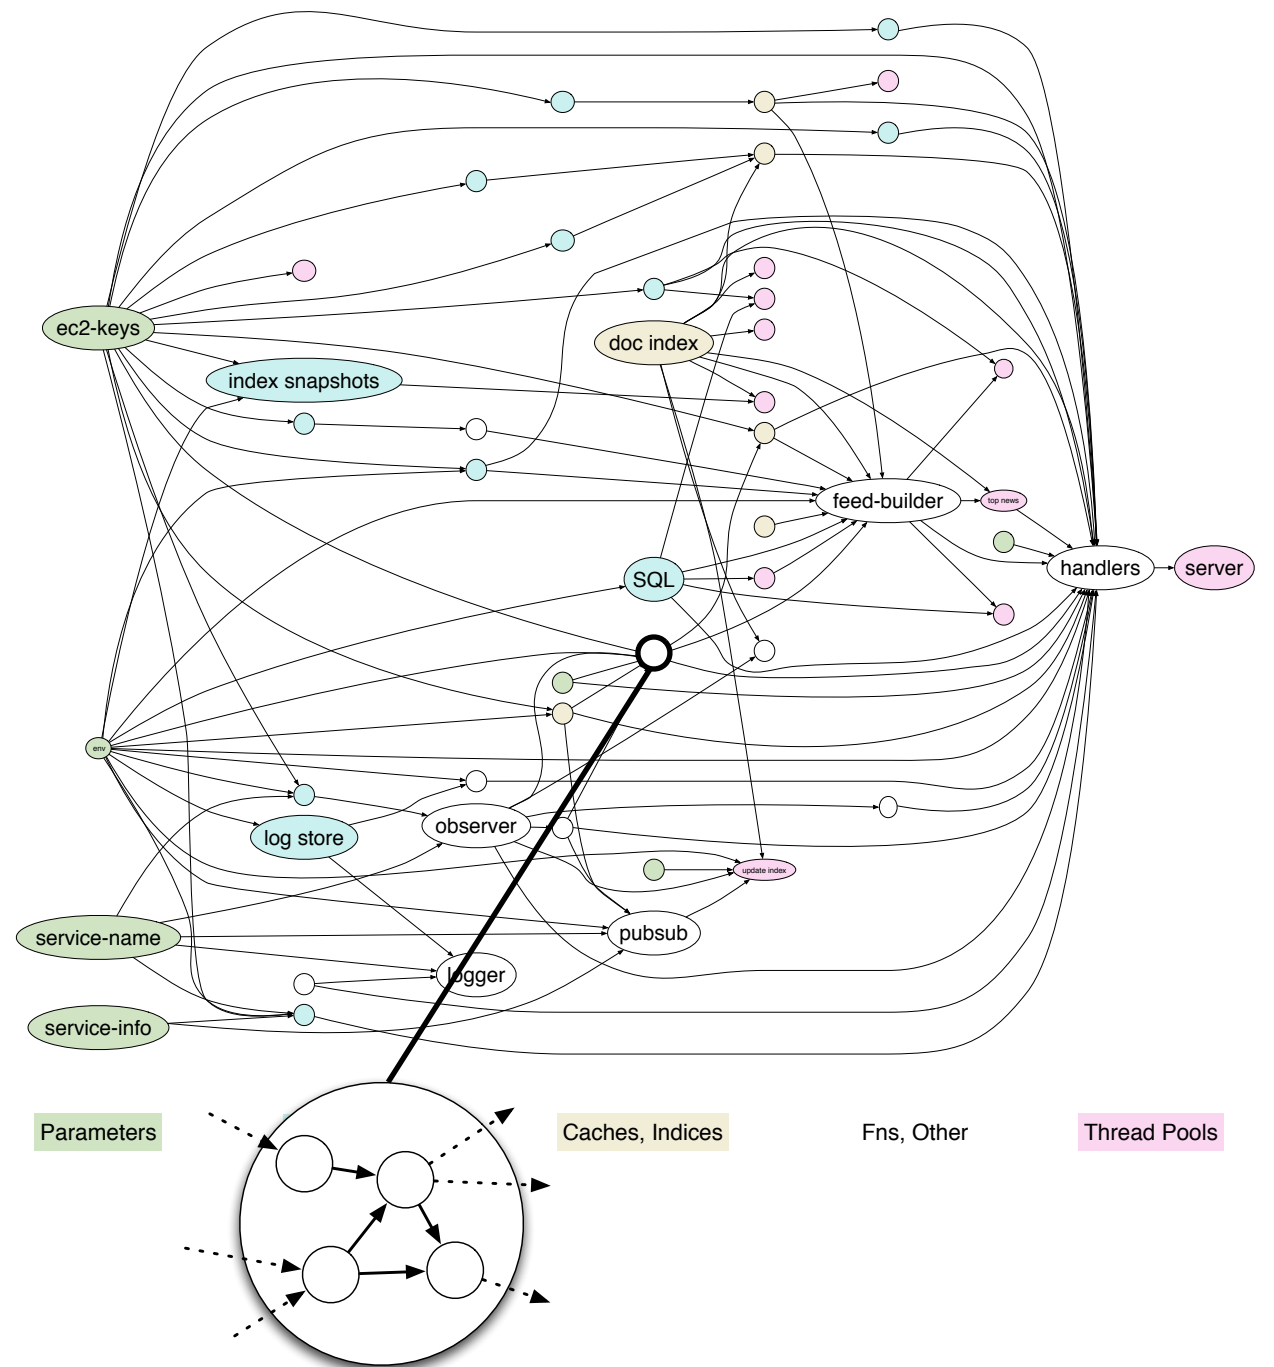
```
(defn start-service [spec]  
  ((->> (:graph spec)  
         resource-transform  
         compile)  
   (:parameters spec)))
```

```
(def api  
  (start-service  
   api-service))  
  
((:shutdown api))
```

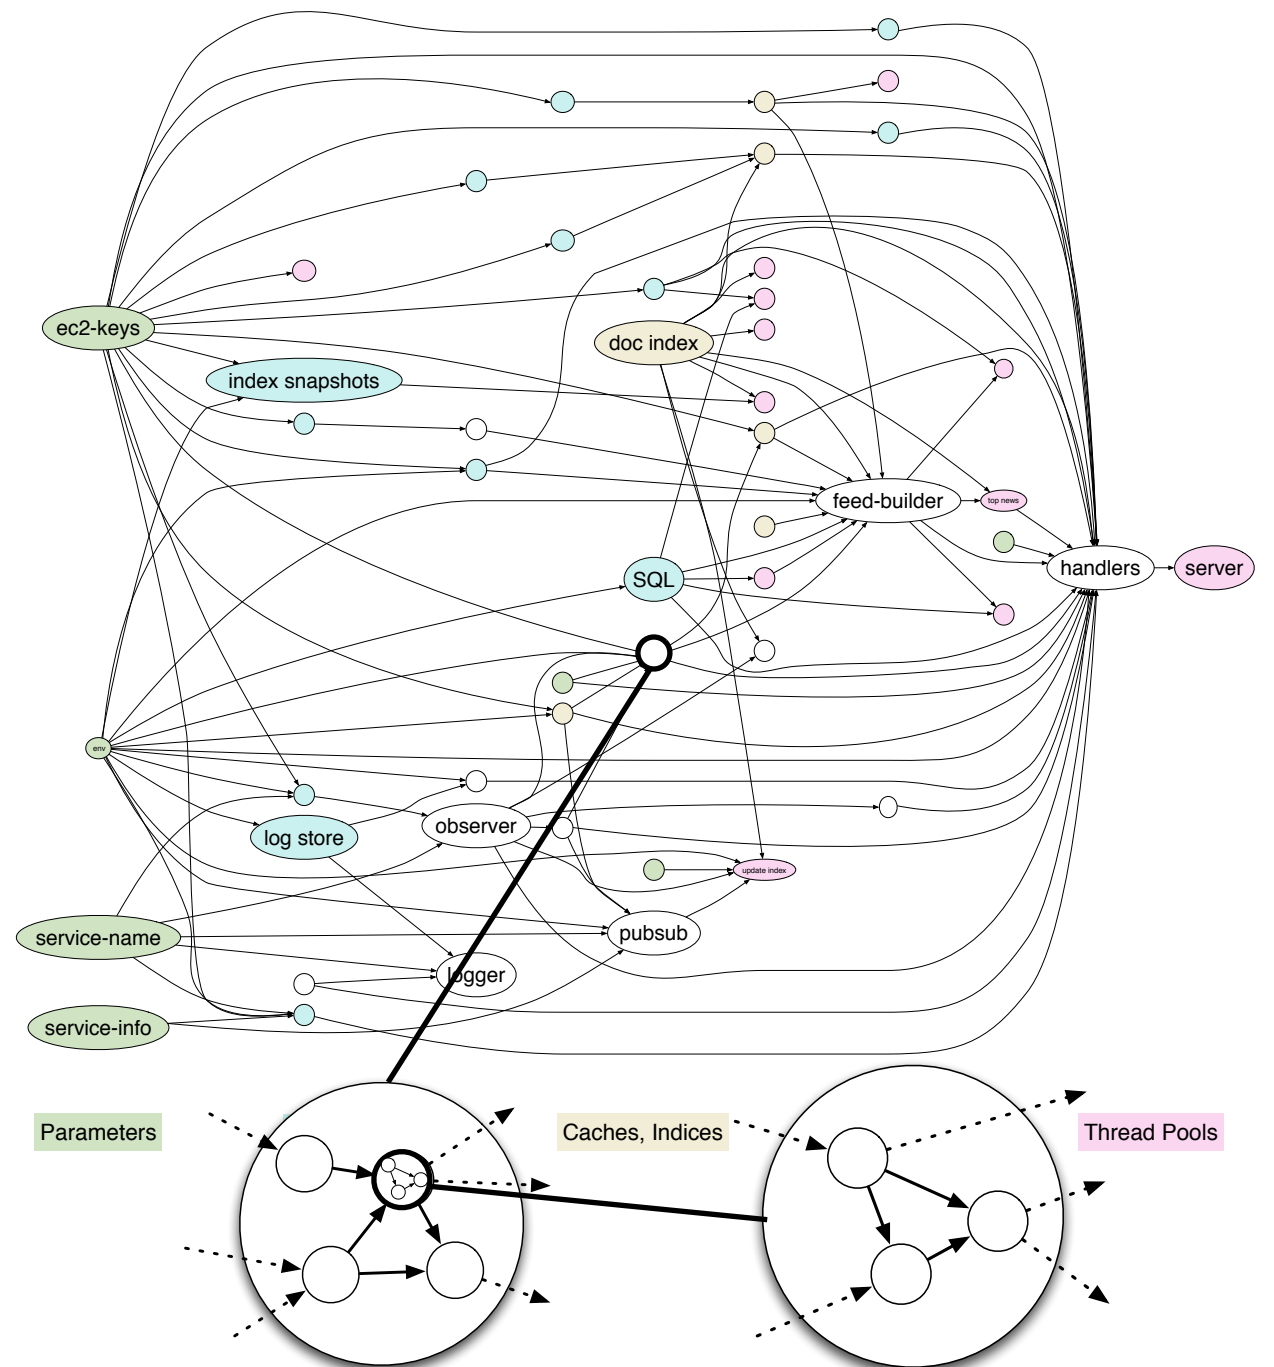
# Sub-Components



# Sub-Components



# Sub-Components





# Sub-Components

- Nodes can themselves be Graphs
  - just nested maps
- Package components as sub-graphs
- Sub-graphs are transparent
  - debugging
  - monitoring
  - imperfect abstractions

```
(def write-back-cache
  {:store
   (instance store ...)
  :write-queue
   (instance queue ...)
  :periodic-prune
   (instance task ...)})
```

# Easy system testing

- Old xxx-line Lets were impossible to test
- With graph, just merge in mock node fnks
- no elaborate mocks objects or redefs
- automatic, safe shutdown

```
(deftest home-feed-systest
  (test-service
    (assoc api-service
      :doc-index
      (fnk [] {:res fake-idx})
      :get-user
      (fnk []
        {:res (constantly me)})))
  (is (= (titles (slurp url))
        ["doc1" "doc2"]))))
```

# Summary

- Graph = way express complex compositions
  - declaratively
  - simply
- Widely applicable
- Simpler code, better tooling
- Hope to open source soon
  - (we're hiring!)

