

# Data Structures

the code that isn't there

“A data structure is just a  
stupid programming language.”

- Ralph William Gosper Jr.

“A data structure is just a  
stupid programming language  
tiny virtual machine.”

# data structures

Fundamentals:

linked list

array

hash table

binary tree

data structures

using list for a set?

ruby's require, circa 1.9.2

[bit.ly/krkvJP](http://bit.ly/krkvJP)

# FAIL



# data structures

"The cheapest, fastest, and most reliable components  
are **those that aren't there.**"

- Gordon Bell

data structures

...set the path of least resistance.

data structures

implementation details  
bubble up to the surface

data structures

git vs. hg

data structures

internal data formats

# data structures



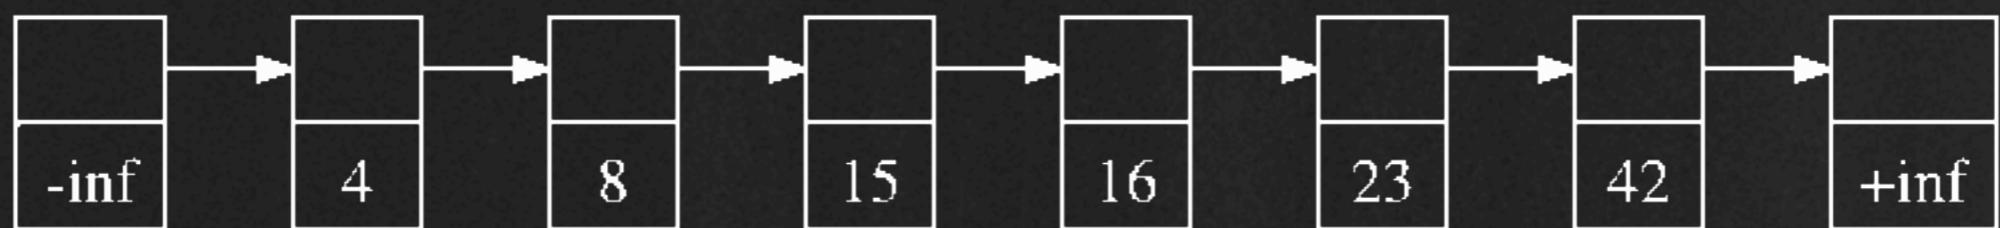
jarring segue

but, enough about the basics...

skip lists

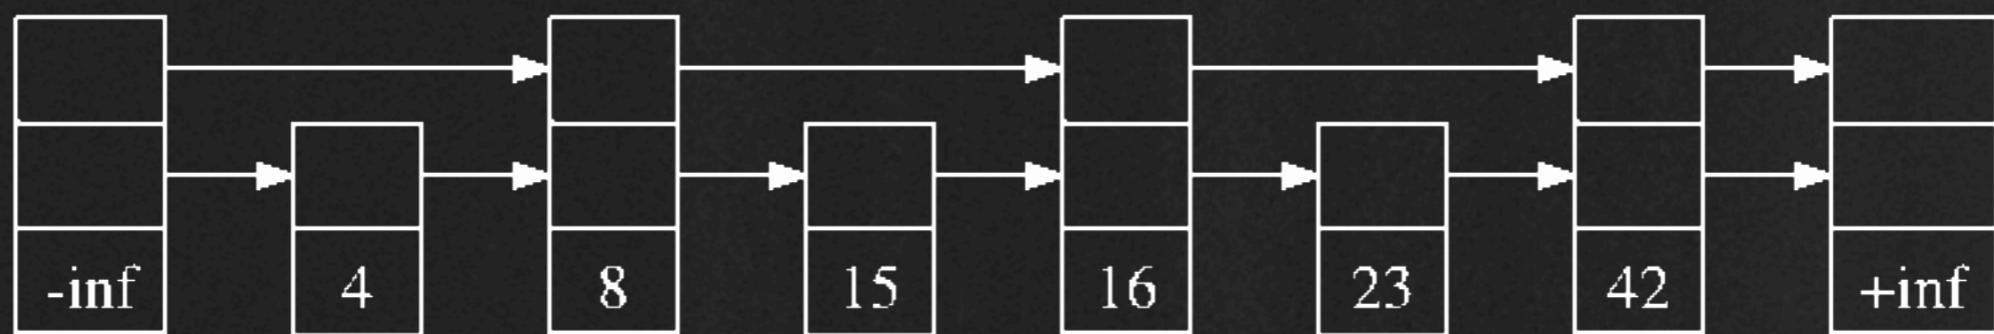
# skiplists

first, take an ordered linked list...



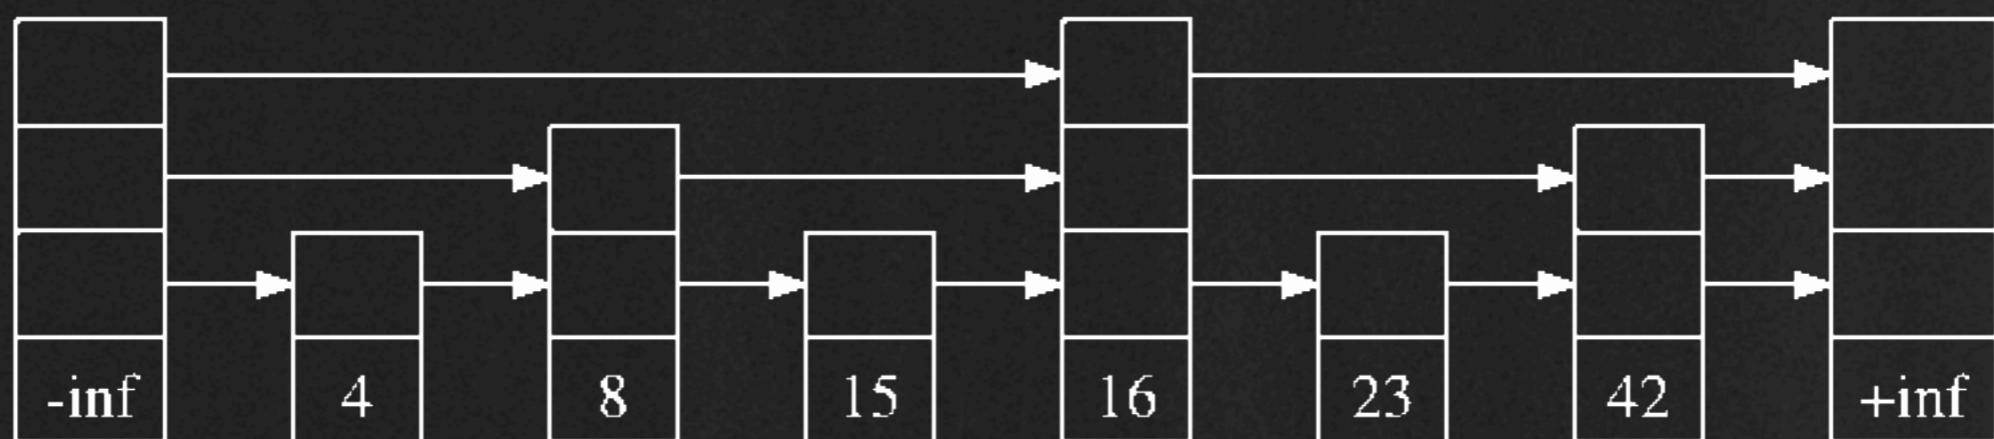
# skiplists

add an “express lane”...

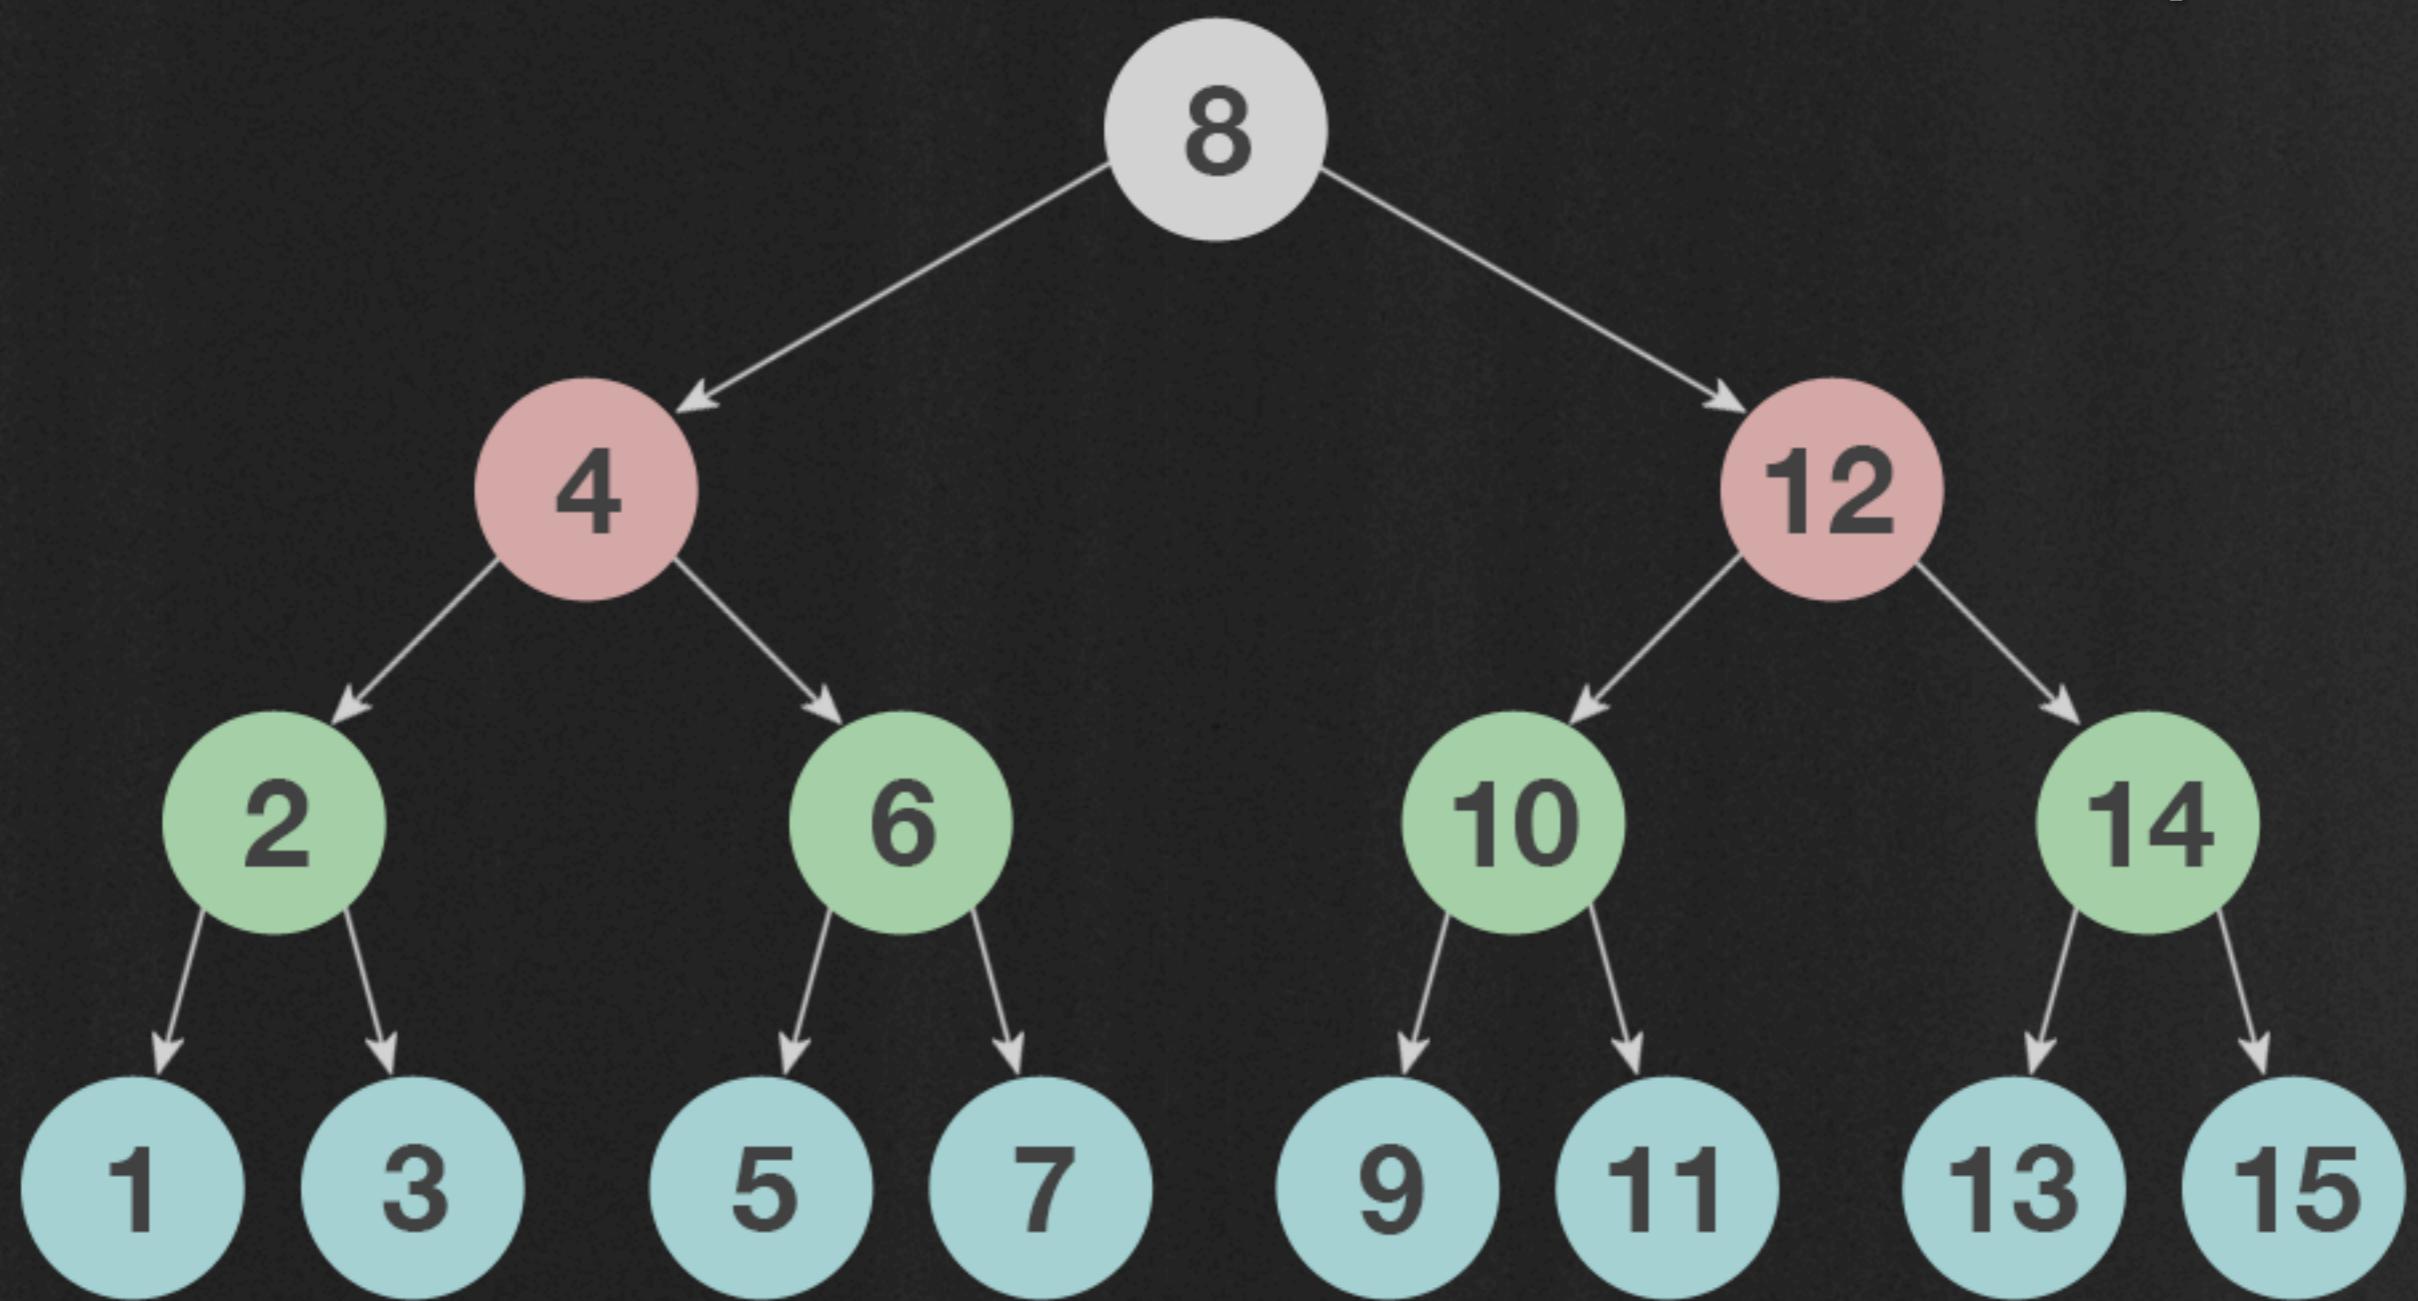


# skiplists

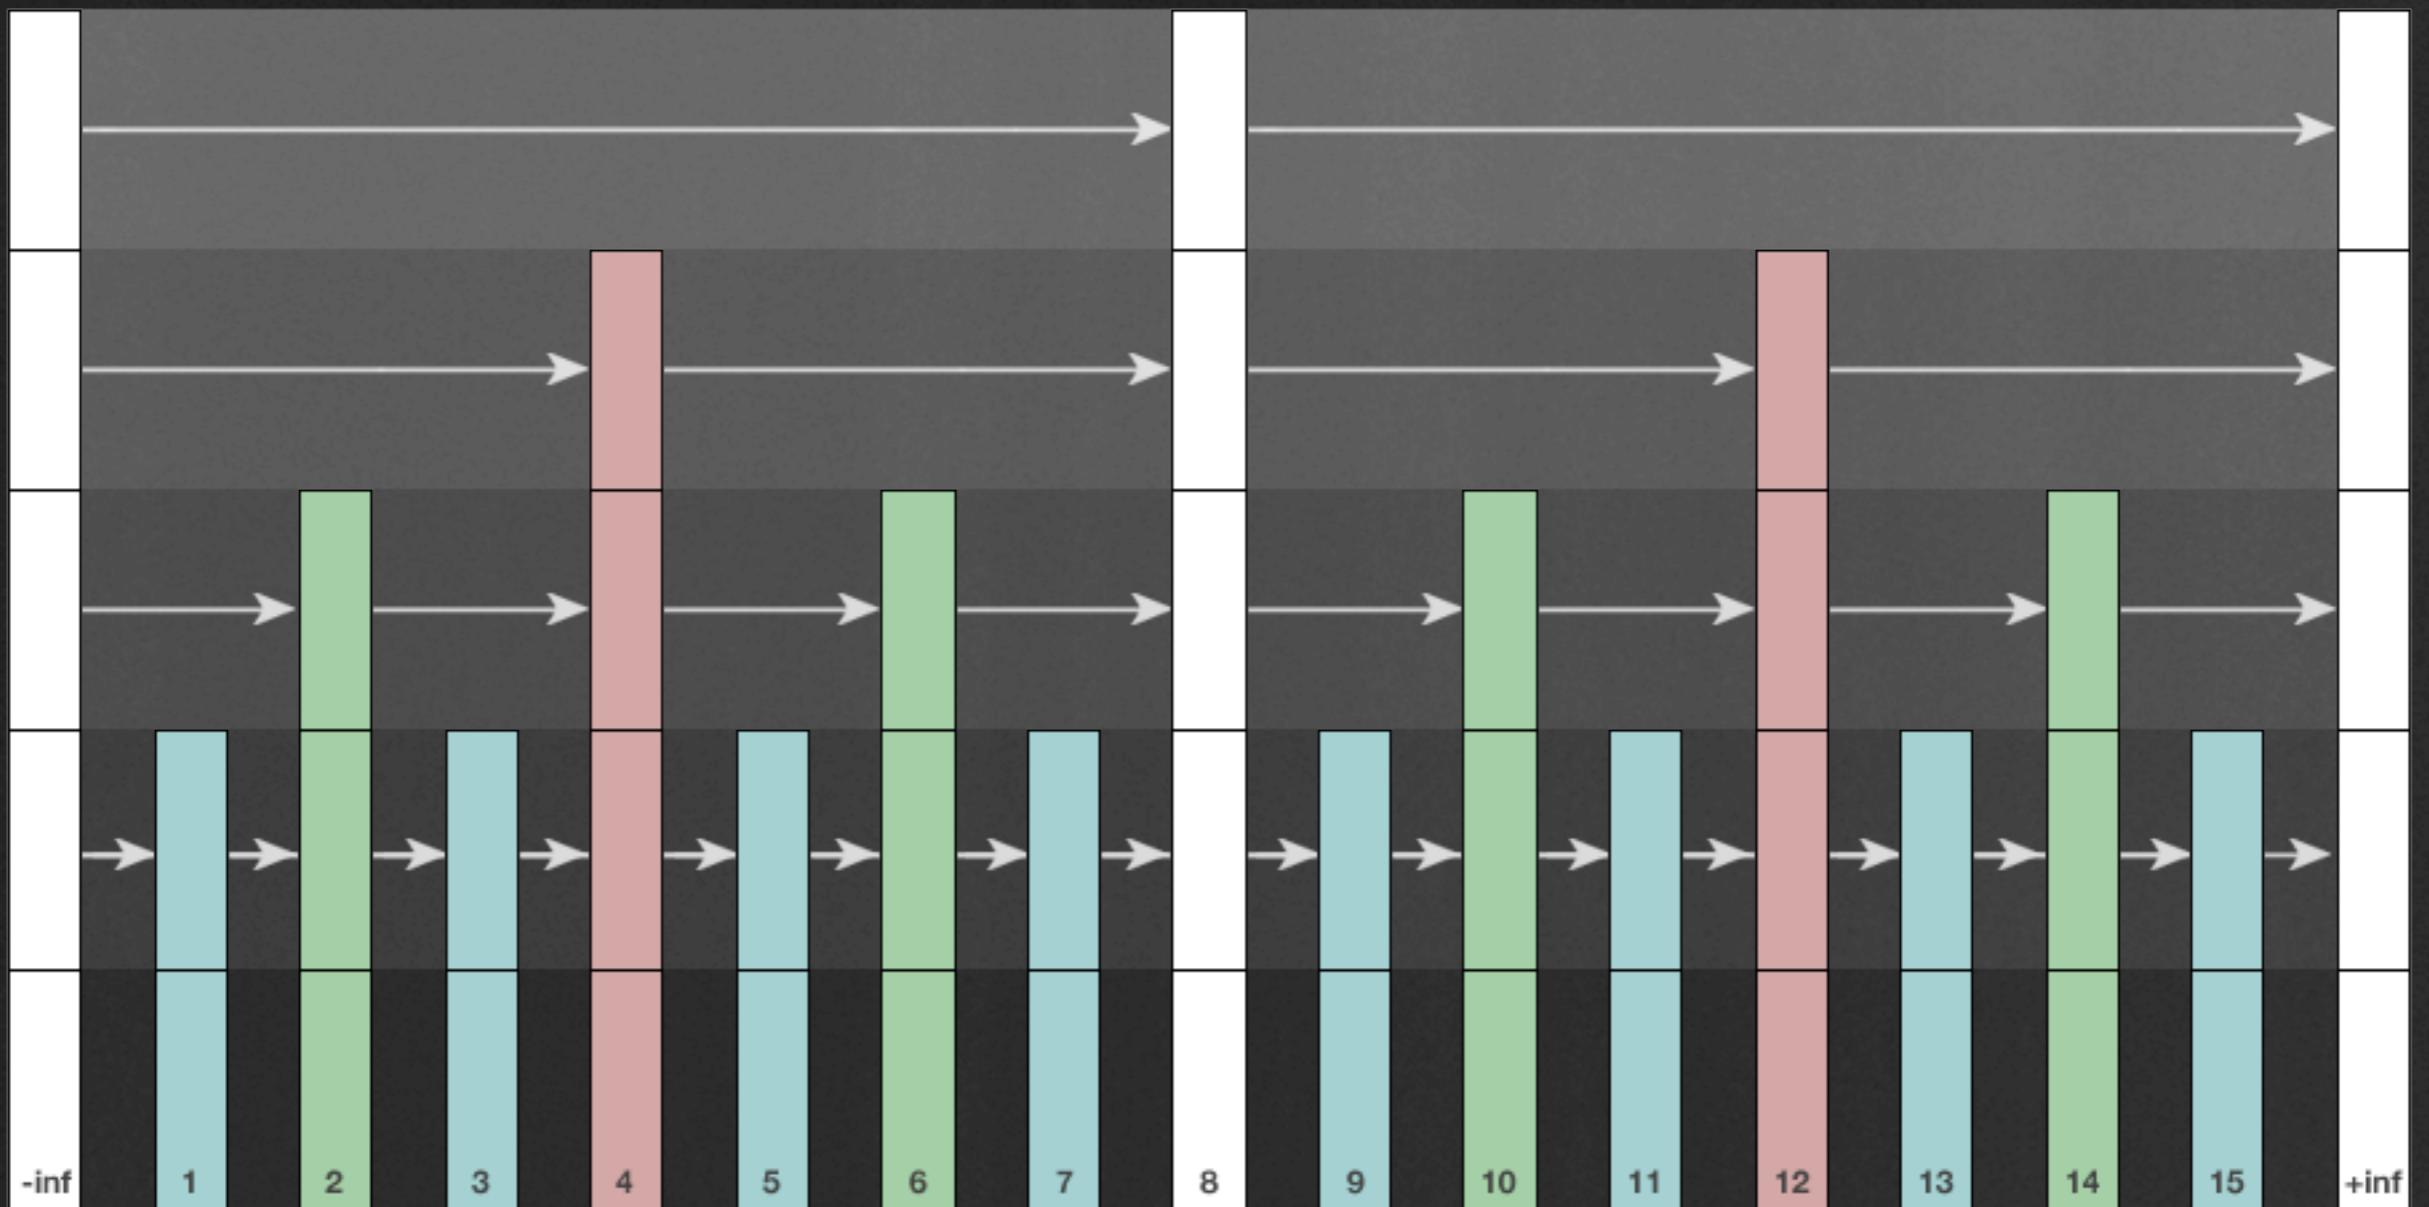
add another...



# skip lists



# skiplists



skiplists

but, how do we balance that?

real trees are not balanced,

skiplists



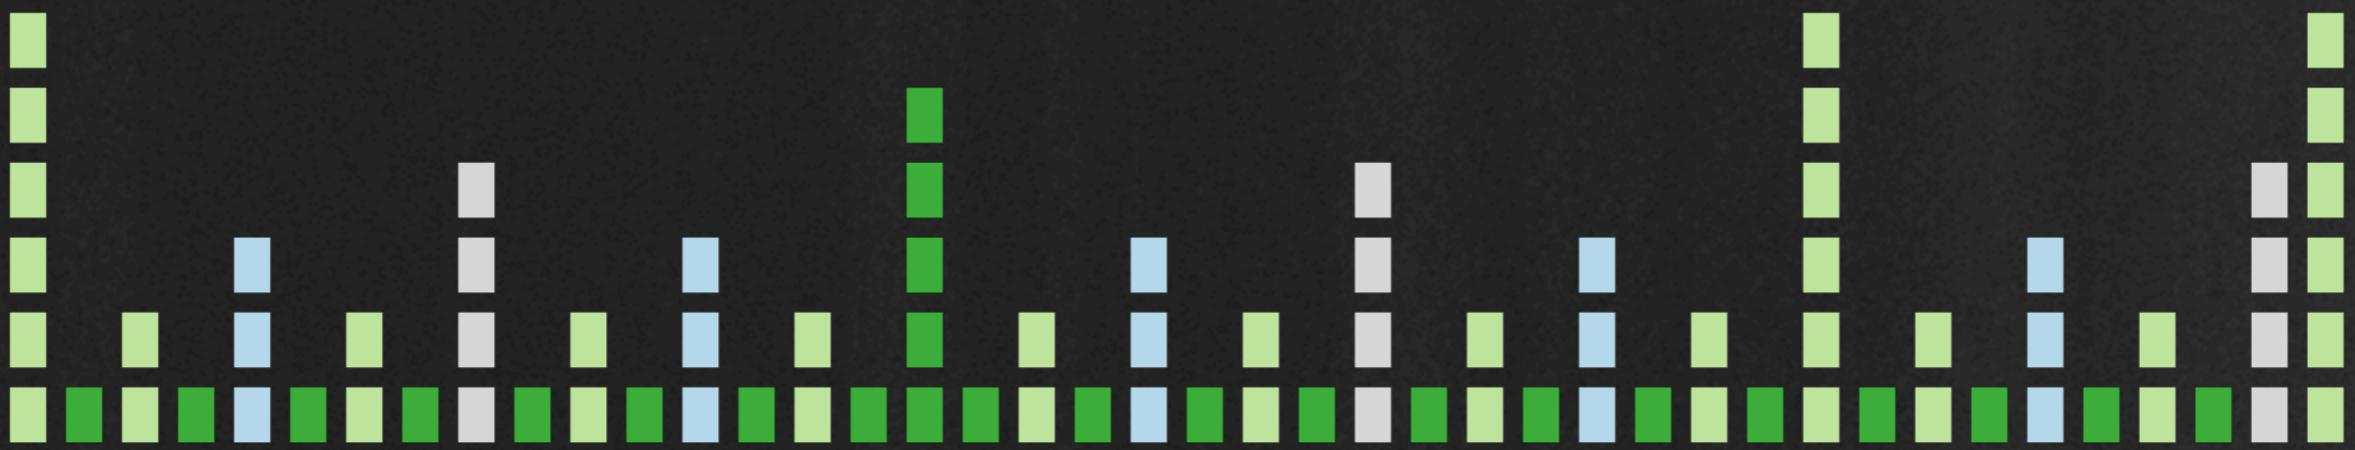
just balanced enough

skiplists

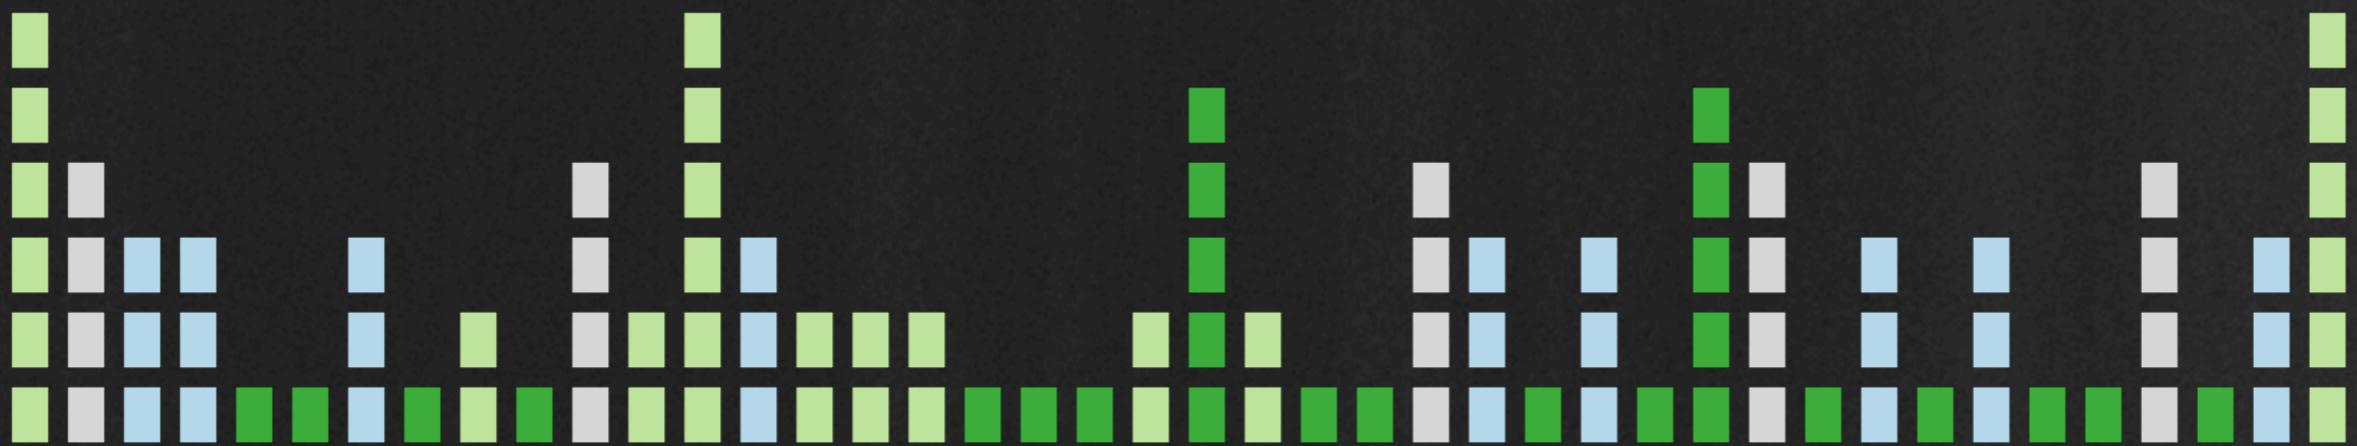
random levels

probability distribution

# skiplists



# skiplists



skiplists

controlled trade-offs

philosophy

“roundabouts, not traffic lights”



philosophy

traffic light

philosophy

roundabout



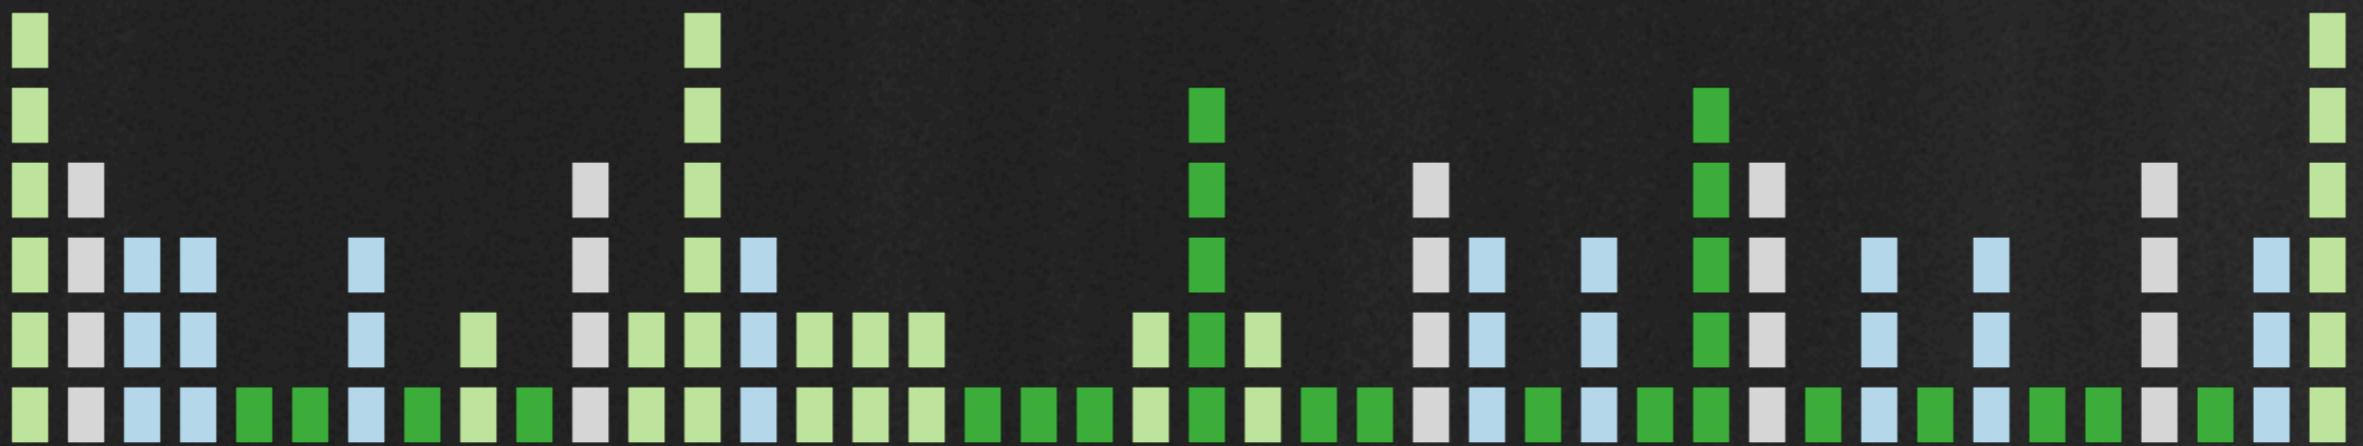
philosophy

global order emerges  
from local decisions

philosophy

only immediate neighbors are affected

most locking is local, little contention



skipping on

still: mutation?

skipping on

still: mutation?

isn't Rich Hickey in the building?

# difference lists

# difference lists

```
?- uses(prolog, Person).
```

no

difference lists

retroactive  
immutability!

# difference lists

## Unification

?- [X, Y, X] = [1, 2, Z].



X = 1

Y = 2

Z = 1

yes

difference lists

data structures with “holes” in them

difference lists

logic variables:

explicitly modeling partial information

# difference lists

## Unification

?- [X, Y, H] = [1, 2, z].

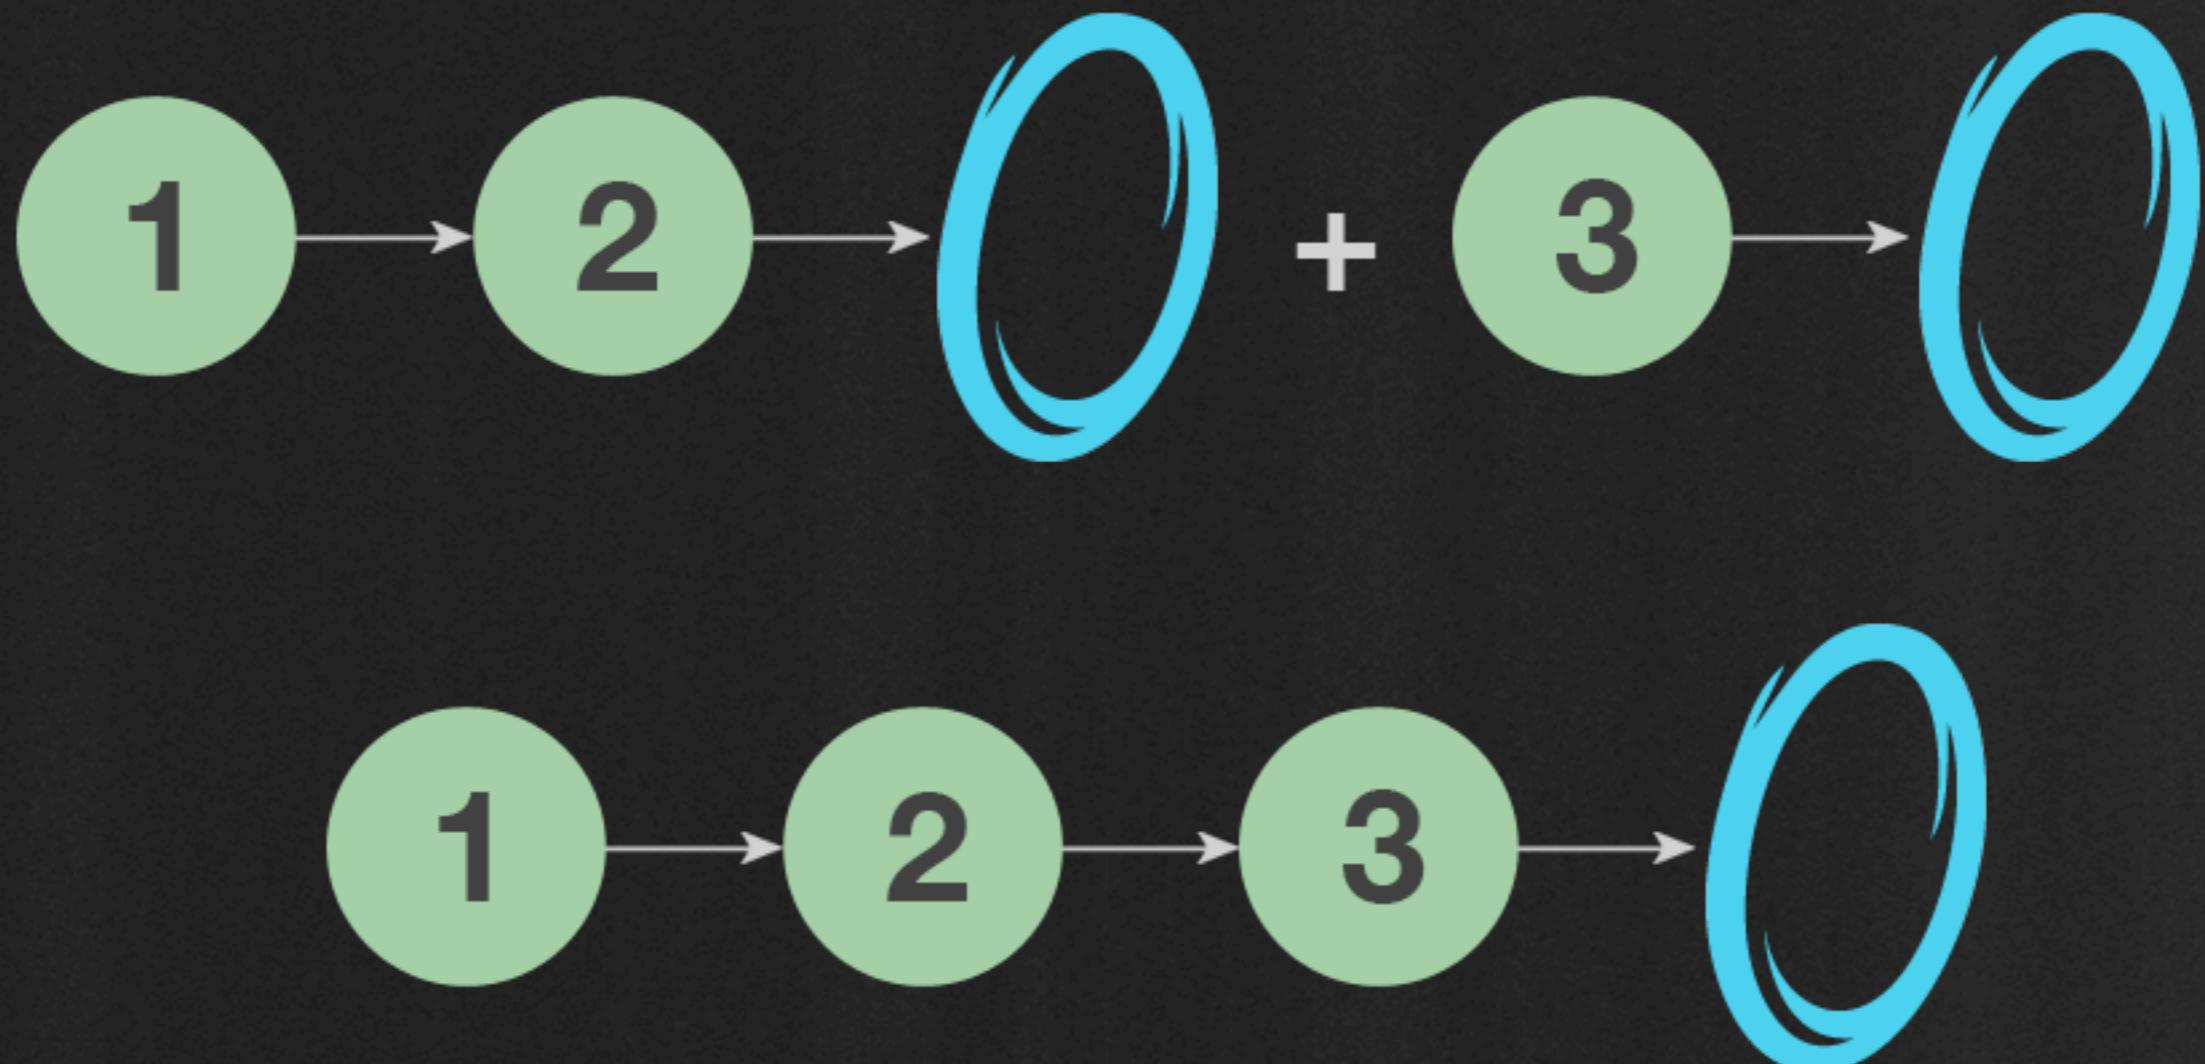
X = 1  
Y = 2  
Z = H ← ???

yes

difference lists

constant-time append  
to an immutable list

# difference lists



# difference lists

```
?- A = [1, 2 | B],  
       B = [3 | C]
```

```
A = [1, 2, 3 | C]
```

yes

# difference lists

```
?- A = [1, 2 | B],  
       B = [3 | C]
```

```
A = [1, 2, 3 | C]
```

yes



Yes!

difference lists

still immutable

...just unstuck in time

# difference lists

more fundamental than lazy evaluation  
or lazy streams

(a bit closer to futures/promises)

difference lists

...apply to more than just lists

difference trees?

difference dictionaries?

rolling along...

rolling hash

find matching/overlapping  
sequences in binary data

rolling hash

hash everything against everything?

**md5, sha1, etc. are too slow**

rolling hash

“Fast, Cheap, and Able to Roll”

## rolling hash

The_quick_brow	0xba5eba11
he_quick_brown	0x0b5e55ed
e_quick_brown_	0xdeadbeef
_quick_brown_f	0xf005ba11
quick_brown_fo	0xbea2ded1
uick_brown_fox	0xdefea7ed

rolling hash

fill then window buffer, then:

drop 1, take 1, new hash  
repeat

rolling hash

**rsync**

rolling hash

sync data across slow network

minimize passes & bandwidth

rolling hash

just send the hash for each block, eh?



# rolling hash

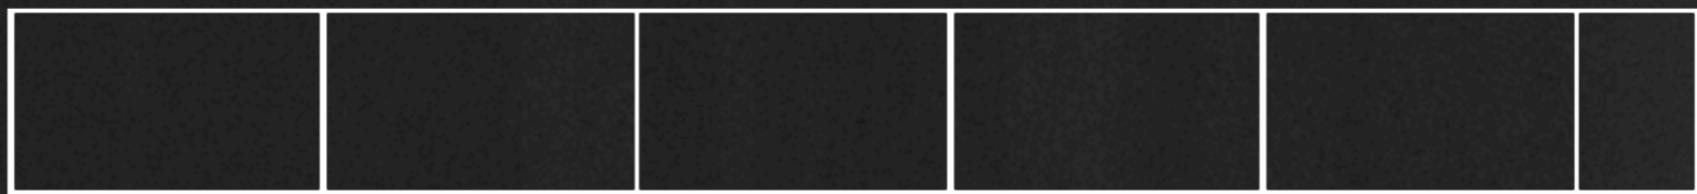
not so fast:  
inserts/deletes shift each block!



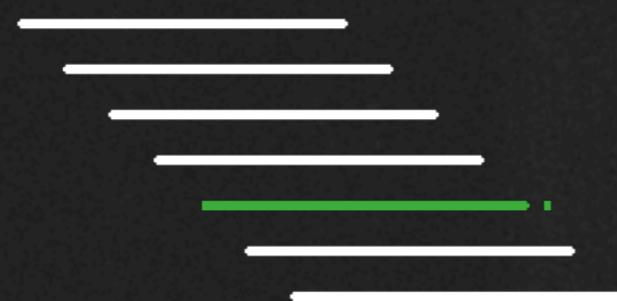
rolling hash

Match with a rolling hash, then sha1

S



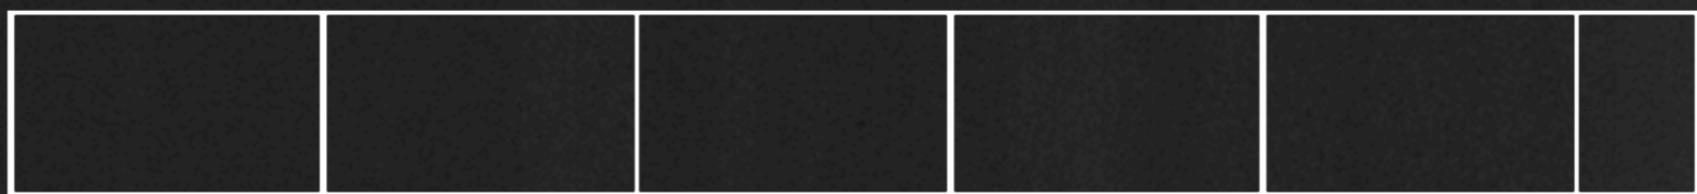
D



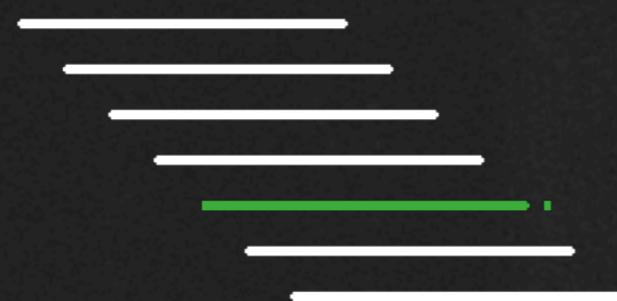
rolling hash

request unmatched regions

S



D



rolling hash

(for more details, check out  
Andrew Tridgell's thesis:

Efficient Algorithms for  
Sorting and Synchronization, pg. 64)

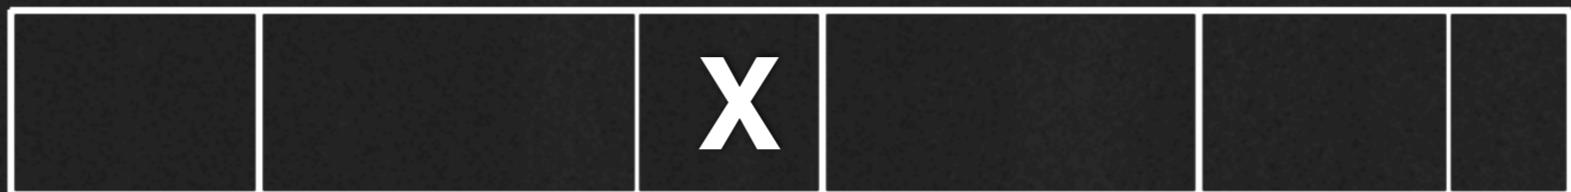
rolling hash

can also be used for chunking data



rolling hash

even with changes, falls back in sync



rolling hash

...finds deterministic breaks

...cheaply matches blocks

# jumprope



jumprope

storage for large binary strings

(where “string” can mean “file”)

jumprope

“content addressable storage”

no pointers, easy distribution

jumprope

persistent and immutable

can be freely cached

jumprope

de-duplication for free

jumprope

...kind of like a git repo

jumprope

three structural elements:

jumprope

leaf

raw\_data

jumprope

limb



# jumprope

trunk



jumprope

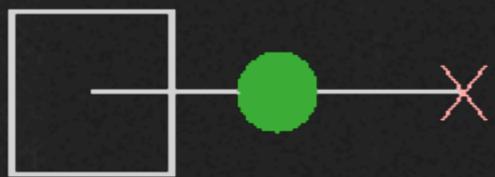
then, it just needs a key/value store

get(hash) => jr\_node

set(hash, jr\_node) => OK | ERROR

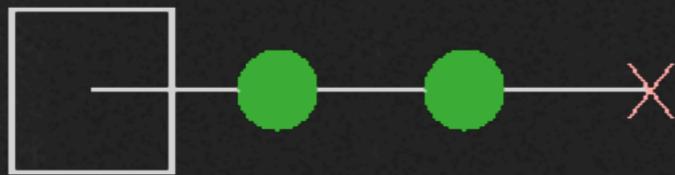
jumprope

building it up



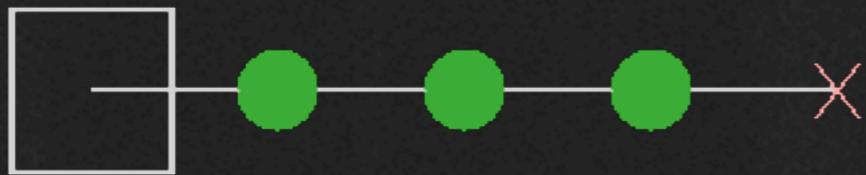
jumprope

add another



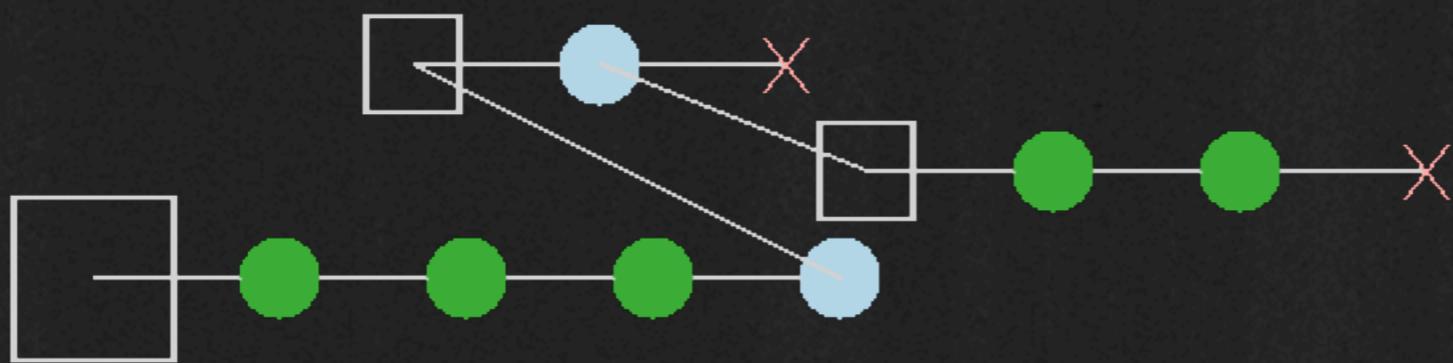
jumprope

and another



jumprope

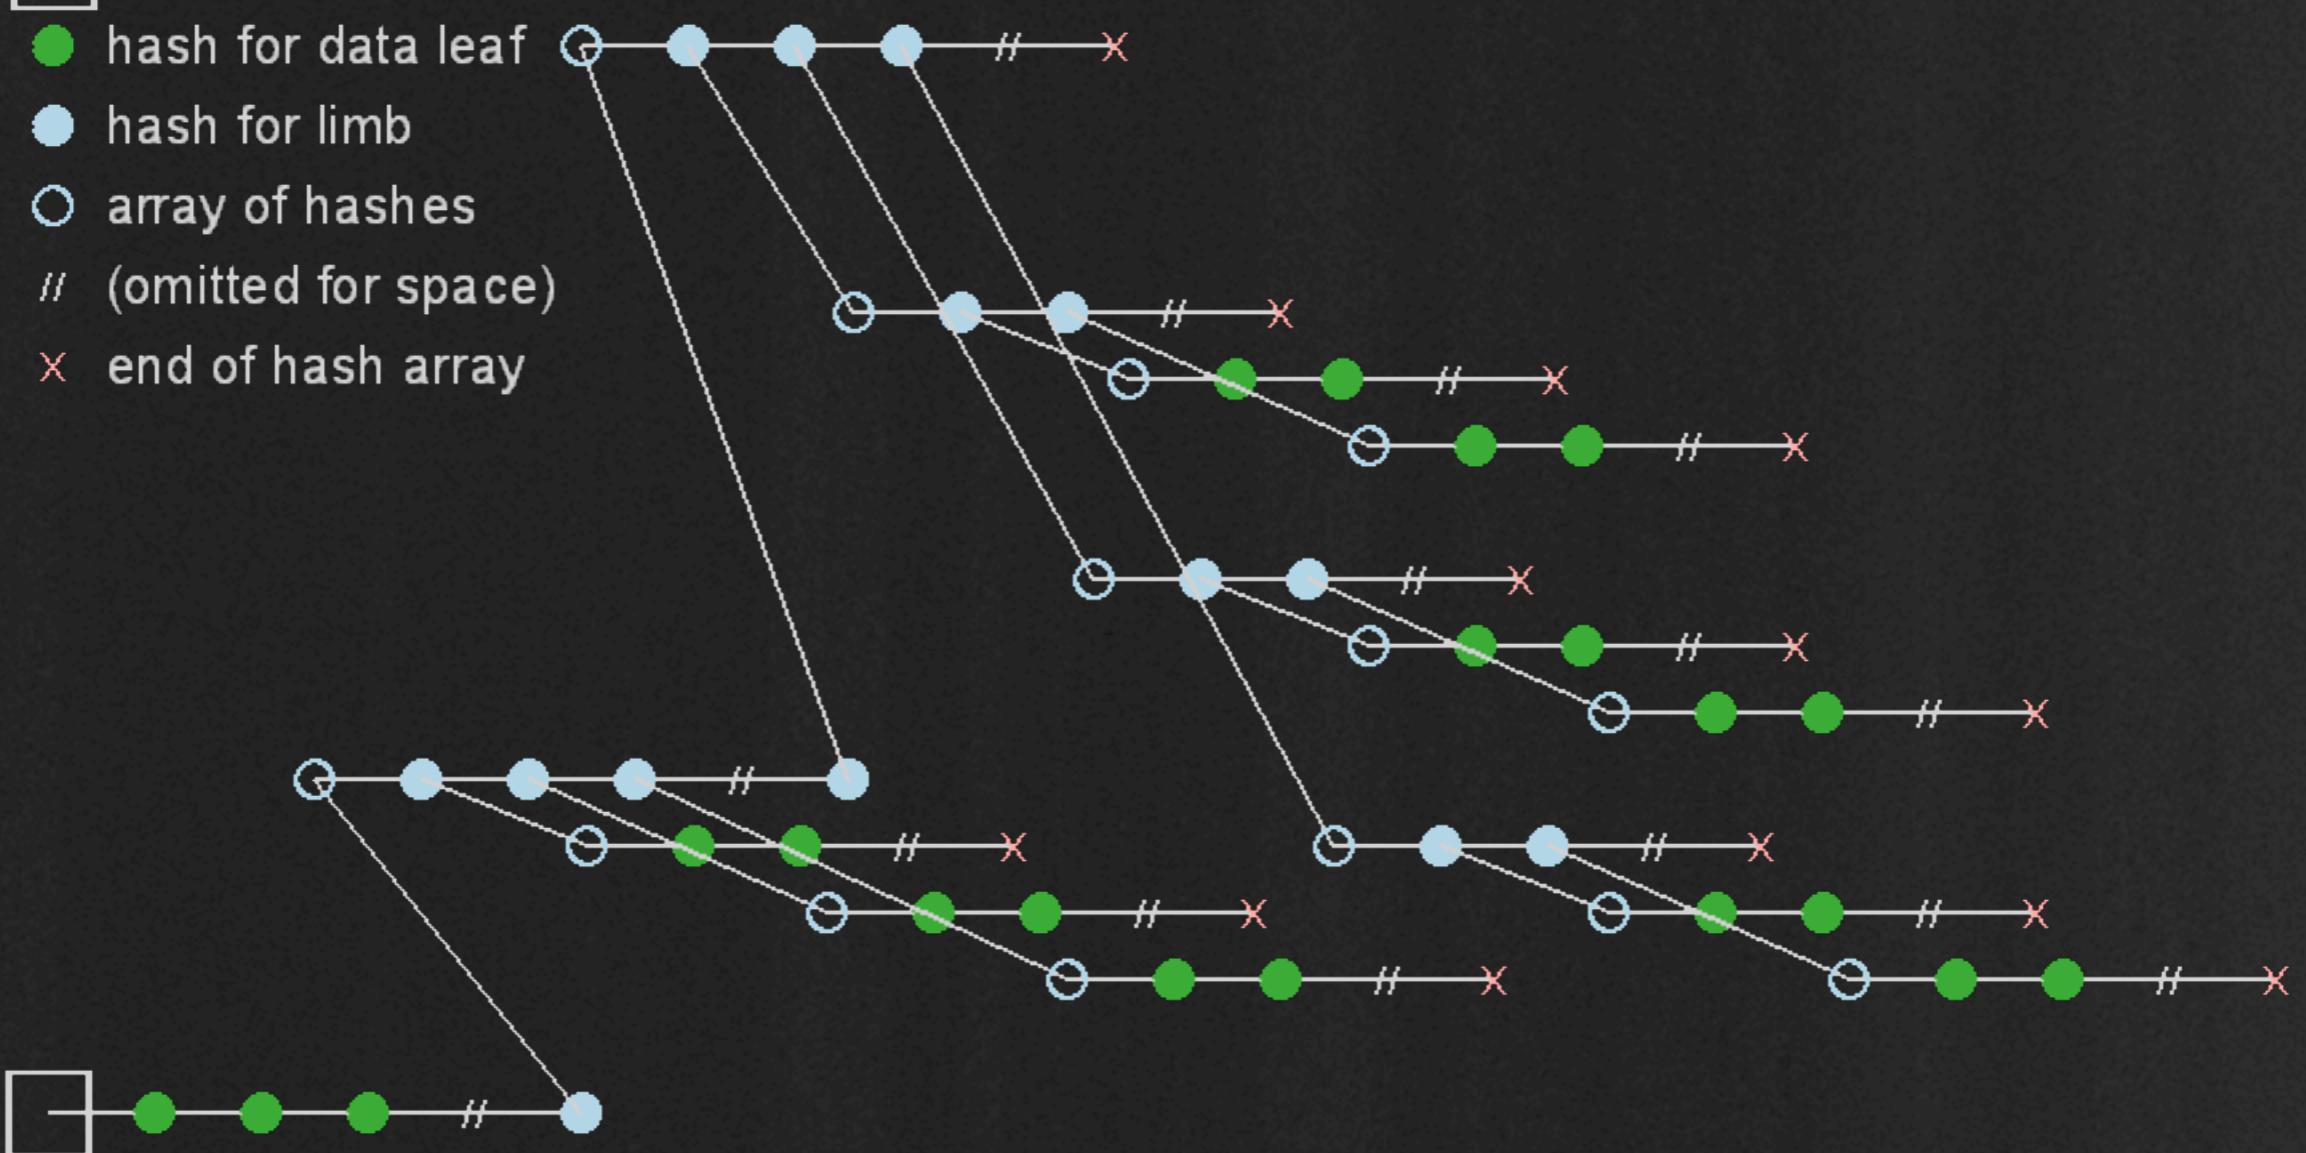
grow a level, then add a limb



# jumprope

## Legend

- ◻ head of jumprope
- hash for data leaf
- hash for limb
- array of hashes
- // (omitted for space)
- ✗ end of hash array



# jumprope

## seek into it



jumprope

duplicated nodes are shared

free de-duplication

jumprope

low metadata overhead

jumprope

fetch it (in parallel)

stream it

mirror it

jumprope

use case?

scatterbrain

properties

randomization

persistence & immutability

emergence from local behavior

lack of pointers (CAS)

We're  
Hiring!



cc attribution

the php hammer:

<http://www.flickr.com/photos/raindrift/sets/72157629492908038/>

jumprope:

<http://www.flickr.com/photos/adwriter/5826718940/>

traffic light:

<http://www.flickr.com/photos/bexross/2636921208/>

cc attribution

roundabout:

<http://www.flickr.com/photos/trendscout/4000299787/>

tree:

<http://www.flickr.com/photos/neilspicys/2348959159/>

leek:

<http://www.flickr.com/photos/richardnorth/7319091956/>

Questions?

Scott Vokes

@silentbicycle

[github.com/silentbicycle](https://github.com/silentbicycle)

